

109 Additional Python Problems

Ilkka Kokkarinen

Chang School of Continuing Education
Toronto Metropolitan University

Version of September 29, 2026

This document contains the specifications for 109 additional Python problems (plus fourteen bonus problems) created by [Ilkka Kokkarinen](#) starting in March 2023, to augment the original set of [109 Python Problems](#) designed for the course [CCPS 109 Computer Science I](#) for Chang School of Continuing Education, Toronto Metropolitan University, Canada (formerly Ryerson University).

After the first couple of dozen problems suitable for the first or second course in Python programming, the remaining problems suit more advanced programming and algorithm courses for computer science majors during the sophomore, junior, and senior years of their undergraduate degrees. These problems also make excellent preparation for a technical coding interview; the last fifth or so are even FAANG-level.

As in the original problem collection, these problems are listed roughly in order of increasing difficulty. The complexity of the solutions for these additional problems ranges from straightforward loops to convoluted branching recursive backtracking searches that require clever optimizations to prune the search tree enough to keep the total running time of the test suite within the twenty-second time limit.

The rules for solving and testing these problems are the same as they were for the original 109 problems. You must implement all your functions in a single source code file that must be named `labs109.py`, the same way as you wrote your solutions to the original 109 problems. The automated test suite to check that your solutions are correct, along with all associated files, is available in the GitHub repository [ikokkari/PythonProblems](#).

Python coders of all levels will hopefully find something interesting in this collection. The author believes there is room on Earth for all races to live, prosper, and get strong at coding while having fun solving these problems. Same as with the original 109 problems, if some problem doesn't spark joy for you, please don't get stuck with that problem, but just skip it and move on to the next problem that interests you. Life is too short to be wasted on solving stupid problems.

Table of Contents

1. The Fischer King	7
2. Multiplicative persistence	8
3. Top of the swops	9
4. Lowest common dominator	10
5. Discrete rounding	11
6. Translate	12
7. Ifs and butts	13
8. Word positions	14
9. Power prefix	15
10. Deterministic finite automata	16
11. The parking lot is full	17
12. Lychrel numbers	18
13. Powertrain	19
14. First fit bin packing	20
15. Count possible triangles	21
16. Arrow walk with me	22
17. Count Friday the Thirteenths	23
18. Median filter	24
19. The magic knight of Muhammad ibn Muhammad	25
20. Man of La Cancha	26
21. Nondeterministic finite automata	27
22. Count monochromatic rectangles	28
23. Count palindromic substrings	29
24. Longest mirrored substring	30
25. Square lamplighter	31
26. Pairwise lamps	32
27. Accumulating merge	33
28. And they walk in twos or threes or more	34
29. Ladies and gentlemen, Conway Bitty	35
30. Lowest fraction between two fractions	36
31. Reasonable filename comparison	37
32. List the Langford violations	38

33. Ten pins, not six, Dolores	39
34. Strict majority element	40
35. Add like an Egyptian	41
36. Sorting by pairwise swaps	42
37. Van der Corput sequence	43
38. Maximal repeats	44
39. It's a game, a reflection	45
40. Reverse a Fibonacci-like sequence	46
41. Recamán sequence	47
42. Mian–Chowla sequence	48
43. Stern–Brocot path	49
44. Carryless multiplication	50
45. Mū tōrere boom-de-ay	51
46. A place for everything and everything in its place	52
47. Scoring a tournament bridge hand	53
48. Manimix	54
49. Count distinct substrings	55
50. Replacement with perfect hindsight	56
51. Replacement with perfect foresight	57
52. When there's no item, there's no problem	58
53. Boxed away	59
54. 'Tis but a scratch	60
55. Do or die	61
56. Arithmetic skip	62
57. Carving Egyptian fractions	63
58. Largest square of ones	64
59. Infection affection	65
60. Prize strings	66
61. String shuffle	67
62. Weak Goodstein sequence	68
63. Markov distance	69
64. A very graphy caterpillar	70
65. St. Bitus' Dance	71
66. Digit string partition	72
67. Hofstadter's figure-figure sequences	73
68. [Be]t[Te]r [C][Al]l [Sm][Al]l	74
69. Spiral matrix	75
70. Baker–Norine dollar game	76
71. Total covered area	77
72. Balsam for the code	78
73. Flip those trips	79

74. Maximal disk placement	80
75. Nice sequence	81
76. Forbidden digit	82
77. Decoding the base $-1 + i$	83
78. Independent dominating set	84
79. Vertex cover	85
80. Shotgun sequence	86
81. Card row game	87
82. The remains	88
83. Count sublists with odd sums	89
84. Tailfins and hamburgers	90
85. Split the digits, maximize the product	91
86. Tower of cubes	92
87. Gijswijt sequence	93
88. Sum of consecutive squares	94
89. Balance of power	95
90. Longest zigzag subsequence	96
91. Kimberling's expulsion sequence	97
92. Kimberling's repetition-resistant sequence	98
93. Game with multiset	99
94. Tower of Babel	100
95. Make a list self-describing	101
96. Two pins, not three, Dolores	102
97. Out where the buses don't run	103
98. SMETANA interpreter	104
99. The sharpest axes	105
100. Vidrach Itky Leda	106
101. How's my coding? Call 1-800-3284778	107
102. Scatter her enemies	108
103. Sneaking	109
104. Inverse pair sums	110
105. Blocking pawns	111
106. Boggles the mind	112
107. The round number round	113
108. Complete a Costas array	114
109. Oppenhoppenheimereimer	115
Bonus problem 110: String stretching	116
Bonus problem 111: Casinos hate this Toronto man!	117
Bonus problem 112: Word bin packing	118
Bonus problem 113: Probabilistic tic-tac-toe	119
Bonus problem 114: Bandwidth minimization	120

Bonus problem 115: Set splitting	121
Bonus problem 116: Domino poppers	122
Bonus problem 117: Knight jam	123
Bonus problem 118: Cubes on the trailer	124
Bonus problem 119: Tom and Jerry	125
Bonus problem 120: Pinch to a pound	126
Bonus problem 121: Minimal Egyptian fractions	127
Bonus problem 122: Unity partition	128
Bonus problem 123: Ants on the rod	129

1. The Fischer King

```
def is_chess_960(row):
```

Whenever you think that you have found a good solution, you should still look around a bit more, since a great solution might be waiting just around the corner. Tired of seeing his beloved game of kings devolve into rote memorization of openings, the late chess grandmaster and generally colourful character Bobby Fischer generalized chess into a more exciting variation of “Chess 960” that plays otherwise the same except that the home rank pieces are randomly permuted before each match begins, which should render any memorization of openings moot.

So that neither player gains an unfair random advantage from the get-go, both black and white pieces are permuted the same way. Furthermore, to maintain the spirit of chess, the two bishops must be placed on squares of opposite colour (that is, different **parity** of their column numbers), and the king must be positioned between the two rooks to enable **castling**. Given the home rank as some permutation of 'KQrrbbkk' with these letters denoting the **K**ing, **Q**ueen, **r**ook, **b**ishop and **k**night, this function should determine whether that string constitutes a legal initial position in Chess 960 under this discipline. Your function can assume that the `row` is a string that contains the eight correct chess pieces in their right multiplicities.

row	Expected result	(reason)
'rbrkbkQK'	False	(king is not between rooks)
'rkbQKbkr'	True	
'rrkkKbQb'	False	(king is not between rooks, bishops same parity)
'rbbkKQkr'	True	
'brbQKkkr'	0	(bishops same parity)
'rKrQkbb'	True	
'kbQbkrKr'	0	(bishops same parity)
'bQrKkrkb'	True	

Discrete math enjoyers can convince themselves that exactly 960 out of the $8! / (2! \times 2! \times 2!) = 5040$ possible permutations of these chess pieces satisfy the required constraints, hence the name.

2. Multiplicative persistence

```
def multiplicative_persistence(n, ignore_zeros=False):
```

Multiply the digits of the given positive integer n to produce the next value of n , and repeat this operation until n becomes a single digit. For example, $n = 717867$ becomes 16464, then 576, then 210, finally 0, where it will stay. This function should return the **multiplicative persistence** of n , that is, the number of iteration rounds required to turn n into a single digit in this manner.

Since even one zero digit makes the entire product zero, most values of n become zero quickly. To make this process that much more interesting, if the keyword argument `ignore_zeros` is set to `True`, zero digits are ignored during multiplication.

n	ignore_zeros	Expected result
717867	False	4
717867	True	4
36982498883598862928	False	2
36982498883598862928	True	5
7899978679899687	False	3
7899978679899687	True	6

When zeros are not ignored, the number $n = 277777788888899$ holds the highest known multiplicative persistence of 11 rounds. Since multiplication doesn't depend on the order of digits, any permutation of those same digits has the same multiplicity, and listing the digits in ascending order gives the smallest number with that persistence. Persistences seem to go up to eleven but no higher, since no integers with a higher persistence than eleven are known, nor are believed to exist.

3. Top of the swops

```
def topswops(cards):
```

The late great John Horton Conway invented all sorts of whimsical mathematical puzzles and games that are simple enough even for wee children to understand, yet contain far deeper mathematical truths beneath a superficial veil. This problem deals with topswops, a patience puzzle played with `cards` numbered with positive integers from 1 to n .

These `cards` are initially laid out in a row in some permutation, given to your function as a tuple. Each move in topswops reverses the first k cards, where k is the card at the head of the row. For example, if the current permutation is `(3, 1, 5, 2, 4)`, reversing the first three cards gives the permutation `(5, 1, 3, 2, 4)`, from which reversing the first five cards gives the permutation `(4, 2, 3, 1, 5)`, from which reversing the first four cards gives `(1, 3, 2, 4, 5)`, where the game will remain stuck in this fixed state.

For every permutation of these n cards, the card that carries the number one must eventually end up as the first card in the row to terminate this process. (Discrete math enjoyers can again try proving this guaranteed termination as a side quest.) Your function should compute and return the number of moves needed to reach this fixed state from the given initial permutation of `cards`.

<code>cards</code>	Expected result
<code>(1, 2)</code>	0
<code>(2, 1)</code>	1
<code>(3, 2, 5, 1, 4)</code>	5
<code>(7, 1, 2, 4, 6, 3, 5)</code>	7
<code>(4, 1, 8, 5, 6, 9, 2, 7, 3)</code>	13
<code>(4, 10, 7, 6, 15, 8, 14, 19, 2, 17, 3, 16, 12, 13, 5, 1, 9, 11, 18)</code>	20
<code>(9, 1, 3, 17, 8, 7, 4, 16, 15, 11, 2, 19, 6, 14, 12, 10, 5, 18, 13)</code>	28

A still-open problem in combinatorics is to generate a permutation for the given n that requires the maximum number of moves to terminate. Exact solutions for this are known only up to $n = 17$. This problem illustrates the principle that even if a function f is straightforward to compute and understand, its inverse f^{-1} might not be equally simple, so finding even one x for a given y such that $y = f(x)$ can be exponentially hard.

4. Lowest common dominator

```
def lowest_common_dominator(beta, gamma):
```

The **prefix sum sequence** of the given integer sequence is defined so that each element equals the sum of the elements of the original sequence up to and including that position. For example, the prefix sum sequence of [3, -1, 0, 5, 2] would be [3, 2, 2, 7, 9].

An integer sequence α **dominates** another integer sequence β of the same length if every element of the prefix sequence of α is greater than or equal to the corresponding element of the prefix sequence of β . The concept of domination defined this way is a **partial order** that is **reflexive** (every sequence dominates itself) and **transitive** (if α dominates β and β dominates γ , then α dominates γ). However, note that it is not necessary for every element of α to be greater than or equal to the corresponding element of β for the sequence α to dominate β ; for a trivial example, the sequence [10, -2, -3] dominates the sequence [-10, 10, 4].

One can further prove that for any two sequences β and γ , there exists a unique **lowest common dominator** sequence α that dominates both β and γ , so that any other sequence that dominates β and γ will also dominate α . (If you want to get all technical about it, the domination relation defines a **lattice** where any two elements have a unique **supremum** and **infimum**.) This function should return the lowest common dominator of the sequences beta and gamma.

beta	gamma	Expected result
[5]	[7]	[7]
[1, -2]	[0, 0]	[1, 0]
[5, 5, -5]	[2, 5, 0]	[5, 10, 7]
[0, 3, -9, -12]	[-8, -4, 13, -13]	[0, 3, 1, -12]
[15, 12, -6, 5, 11]	[5, -1, -14, -17, 20]	[15, 27, 21, 26, 37]

5. Discrete rounding

```
def discrete_rounding(n):
```

Make the given positive integer n your current number. For each k counting down from $n - 1$ to 2, “round up” your current number by replacing it with the smallest positive integer that is exactly divisible by k and is greater than or equal to your current number. Return the final number acquired after these replacements.

For example, starting with $n = 5$ and counting down k from 4 down to 2, the smallest positive integer divisible by 4 and greater than or equal to 5 is 8. Next, the smallest positive integer divisible by 3 and greater than or equal to 8 is 9. Lastly, the smallest positive integer divisible by 2 and greater than or equal to 9 is 10, our final answer. The reader can similarly verify that, starting from $n = 12$, this iteration produces $12 \rightarrow 22 \rightarrow 30 \rightarrow 36 \rightarrow 40 \rightarrow 42 \rightarrow 42 \rightarrow 45 \rightarrow 48 \rightarrow 48 \rightarrow 48$.

n	Expected result
1	1
2	2
10	34
23	174
$10^{**}4$	31833630
$10^{**}6$	318310503562
$10^{**}7$	31830995532658

Of course, this problem makes for a simple exercise in Python loops and integer arithmetic for the first programming course. Still, it also has a curious mathematical property worth a closer look. If we denote the result of this process starting from n by $r(n)$, the expression $n^2 / r(n)$ approaches π in the limit as n approaches infinity! For example, when n equals one million, this expression already has the first four decimal places of π correct, and when n equals one billion, the first eight. Readers may enjoy pondering what connection integer arithmetic iterated in this manner could have to circles so that it produces the famous transcendental number as its limit using only integer arithmetic and divisibility.

6. Translate

```
def tr(text, ch_from, ch_to):
```

This simple text-processing function implements the behaviour of the Unix command-line tool `tr`, but, of course, operates on the given argument `text` as a Python function rather than reading or writing to standard input or output. This function should return a string that is otherwise the same as the original `text`, except that every occurrence of any character inside the `ch_from` string has been replaced by the corresponding character in the same position in the `ch_to` string. Both strings `ch_from` and `ch_to` are guaranteed to be the same length, and `ch_from` will never contain the same character twice.

text	ch_from	ch_to	Expected result
'X'	'Y'	'Z'	'X'
'jonny'	'jony'	'AAAA'	'AAAAA'
'abccba'	'bc'	'XY'	'aXYYXa'
'bee'	'fgabhed'	'jjidflb'	'dll'
'yeah!'	''	''	'yeah!'
'abcde'	'abcde'	'edcba'	'edcba'

7. Ifs and butts

`count_cigarettes(n, k):`

As we can see from expressions such as “to treat someone like a dog”, the past was not only a different country, but often a filthy and unhygienic time to be alive. This is illustrated by the classic puzzle from “the obvious immediate answer is actually wrong” genre, about a nicotine-addicted bum who can make himself a complete cigarette from the remains of five cigarette butts. Quickly now, if the bum has a pack of 25 cigarettes to begin with, how many cigarettes does he get to smoke? The answer is not the obvious 30 but actually 31, since the five cigarettes created from the butts of the original twenty-five will themselves produce five butts for the bum to make one more cigarette.

This function should compute the total number of cigarettes smoked by the bum who initially holds n complete cigarettes, assuming that k cigarette butts are needed to create a new cigarette. Your function should keep track of the current number of complete cigarettes and butts that you currently hold. After smoking all your complete cigarettes, create as many complete cigarettes from the butts as you can and repeat. Don't toss the leftover butts from each round, since they might come in handy in later rounds.

n	k	Expected result
25	5	31
9	3	13
219	3	328
1262911996619460878972361204638641351 1706819653640642644878	47	12903666052416230719934994 91696003119674392442871978 7050201

(Short terminology public service announcement; a **hobo** travels to work; a **tramp** travels to avoid work; a **bum** neither travels nor works.)

8. Word positions

```
def word_positions(sentence, word):
```

Here is a simple problem taken from the exercise set of the famous MIT Python course 6.0001. Your function is given a `sentence` made up of some words separated by single whitespace characters, and one particular `word` to look for. This function should return the list of all **word positions** in the `sentence` where that particular `word` appears. Note that this function does not ask for the character positions in the string, but operates at the granularity of individual words.

For simplicity, the sentence contains no punctuation, only words and the single spaces that separate them. Capitalization counts, so lowercase and uppercase versions of the same word count as separate words.

This problem is a straightforward exercise in Python loops and strings. Still, aspiring Pythonistas can take it as a challenge to implement this function as a one-liner using Python list comprehensions and carefully chosen string methods.

sentence	word	Expected result
'Stix nix hix pix'	'hix'	[2]
'Buffalo buffalo Buffalo buffalo buffalo buffalo Buffalo buffalo'	'buffalo'	[1, 3, 4, 5, 7]
'James while John had had had had had had had had had had had a better effect on the teacher'	'had'	[3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13]
'That that is is that that is not is not is that it it is'	'that'	[1, 4, 5, 11]
'That that is is that that is not is not is that it it is'	'is'	[2, 3, 6, 8, 10, 14]

9. Power prefix

```
def power_prefix(prefix):
```

Integer powers of two have an interesting property in that for any prefix of digits that you can think of, there exists a power of two whose base-ten representation starts with precisely that prefix. The non-constructive proof of this surprising fact based on the **pigeonhole principle** is presented, for example, in the free textbook “Mathematics for Computer Science” for all discrete math enjoyers. However, in this programming course, we will simply take this fact as given and apply it to practice writing some loops.

Given a `prefix` string that consists of digit characters 0 to 9 and possibly some asterisks that denote a “don't care” wildcard character that matches any single digit character to make this problem more interesting, this function should return the lowest positive integer k so that the power 2^k starts with that `prefix`.

prefix	Expected result
'*024'	10
'*22*3720'	63
'19**'	84
'299**41083164'	656

10. Deterministic finite automata

```
def dfa(rules, text):
```

Deterministic finite automata (DFAs) are an important tool for text processing in computer science. A DFA consists of a set of **states**, for simplicity numbered 0, 1, 2, ..., and transitions that determine how the automaton changes state in response to the characters it reads. Initially, the DFA starts its operation in state 0. It then processes the characters of the given `text` one at a time from left to right. At each character `c`, the automaton changes its internal state to the successor state `rules[(s, c)]`, where `s` is the state the automaton is currently in.

This function should return the final state the automaton with the given `rules` ends up in after processing the given `text`. You can assume that the rules dictionary is complete, meaning it provides the unique successor state for every possible pair of the current state and the text character.

rules	text	Expected result
{(0, 'a'): 1, (0, 'b'): 0, (1, 'a'): 1, (1, 'b'): 1}	'bb'	0
{(0, 'a'): 0, (0, 'b'): 1, (0, 'c'): 0, (1, 'a'): 0, (1, 'b'): 2, (1, 'c'): 1, (2, 'a'): 0, (2, 'b'): 2, (2, 'c'): 2}	'bca'	0
{(0, 'a'): 3, (0, 'b'): 0, (0, 'c'): 3, (0, 'd'): 2, (1, 'a'): 1, (1, 'b'): 3, (1, 'c'): 3, (1, 'd'): 1, (2, 'a'): 4, (2, 'b'): 4, (2, 'c'): 2, (2, 'd'): 3, (3, 'a'): 2, (3, 'b'): 4, (3, 'c'): 1, (3, 'd'): 3, (4, 'a'): 2, (4, 'b'): 0, (4, 'c'): 1, (4, 'd'): 1}	'aadbcbabba'	2

Deterministic finite automata are important in the theory of computer science because any **regular expression**, an important technique for advanced pattern matching inside text strings, can be converted to an equivalent DFA so that the original regular expression matches the given `text` if and only if the corresponding DFA ends up in one of its **accepting states** after processing that `text`. Once the regular expression is converted into an equivalent DFA, that automaton can match the regular expression with blazing speed in a single linear pass through the `text`, regardless of the original regular expression's complexity.

11. The parking lot is full

```
def parking_lot_permutation(preferred_spot):
```

A parking lot contains n parking spots 0 to $n - 1$, arranged in a circle. A total of n cars arrive at the lot one at a time, also numbered from 0 to $n - 1$ in order of their arrival. When car i arrives at the lot, it first drives to `preferred_spot[i]`. If that spot is still available, the car will park there. Otherwise, the car tries spots after its preferred spot in linear order until it finds an available spot, and parks in the first empty spot it finds. This function should return the list of parking spots that each car ends up in under this discipline.

preferred_spot	Expected result
[0, 0]	[0, 1]
[2, 0, 1]	[1, 2, 0]
[3, 2, 2, 2]	[2, 3, 1, 0]
[2, 4, 4, 0, 1, 4]	[3, 4, 0, 5, 1, 2]
[4, 7, 6, 7, 4, 6, 0, 6]	[3, 5, 6, 7, 0, 4, 2, 1]
[2, 5, 6, 0, 0, 3, 5, 6, 2, 5]	[3, 4, 0, 5, 8, 1, 2, 6, 7, 9]

Under this discipline, the result will always be some permutation of numbers from 0 to $n - 1$. Of course, if each car has its own unique parking spot that no other car prefers, the result will simply be the **inverse** of the `preferred_spot` permutation. Other variations of this combinatorial problem give rise to interesting calculations and theorems, such as the fact that exactly $(n + 1)^{n - 1}$ different `preferred_spot` lists of length n force no car to cycle back to the beginning of the lot.

12. Lychrel numbers

```
def lychrel(n, giveup):
```

Starting from the given positive integer n , repeatedly add the number to the number that you get from reversing its digits until the resulting number becomes a **palindrome**, that is, the same number when read forward and backward. This function should return the number of iterations needed to reach a palindrome number from the given initial value n . If the giveup iterations have reached no palindrome number, return `None`.

Even relatively small starting numbers display great diversity in how quickly this process converges to a palindromic number. For example, it is still an open problem whether this process will ever terminate starting from $n = 196$.

n	giveup	Expected result
101	100	0
168	100	3
196	10000	None
296	100	7
10911	100	55

Starting at $n = 196$, the process eventually reaches values so large they exceed Python's internal limit for integer-to-string conversion, which this author used to reverse the given integer when implementing the solution. Reversing the digits of an integer using only integer arithmetic without any string conversions is another interesting side exercise for aspiring Python maven.

13. Powertrain

```
def powertrain(n):
```

Another wacky creation of John H. Conway is the **power train**, an iterated operation on positive integers that quickly grinds even humongous numbers down to single digits. A single step of the powertrain turns every other digit into an exponent and computes the product of the resulting exponential terms. If n has an odd number of digits, the missing exponent for the last digit is considered zero. For example, the iteration step for $n = 12345$ would compute $1^2 3^4 5 = 1 \times 81 \times 1 = 81$. The second step would then similarly compute $8^1 = 8$. The power train iteration ends when the number contains only a single digit.

This function should return the number of iteration steps that turn the given n into a single-digit number. For this problem, we define $0^0 = 1$, the way Python's exponentiation operator `**` already does. Any higher power of zero already unambiguously equals zero.

n	Expected result
7	0
56	9
1335	3
37424714	2
965151392315953	10
44133127887521824384942738314	4
6484958651441123695313118445723169954837619982	2

Somewhat paradoxically, the larger n , the sooner running the train on it tends to grind it down into a single digit, since once the digit zero appears anywhere in the bases with a nonzero exponent, the entire resulting product becomes zero. Conway discovered that $n = 2592$ is a **fixed point** of the power train transformation and thus will never reach a single digit, and Neil Sloane later established the same property for $n = 24547284284866560000000000$.

14. First fit bin packing

```
def first_fit_bin_packing(items, capacity):
```

Bin packing is a famous problem in combinatorial optimization. The function takes a list of items, each a positive integer representing the item's size, and the capacity of an individual bin, which is the maximum total size of items placed in that bin. The task is to distribute these items efficiently into bins to minimize the total number of bins.

The general bin packing problem is quite difficult, and even the best algorithms will take exponential time to find and verify the optimal distribution of items to bins. However, many **approximation algorithms** have been devised for this problem. These algorithms are guaranteed to find a near-optimal solution quickly, but can sometimes produce a suboptimal solution that uses more bins than necessary. These approximation algorithms are typically **online algorithms**, meaning that the items are considered to arrive one at a time (no peeking ahead allowed!). The decision of which bin to place the current item must be made immediately and cannot be taken back later, even if some later item turns out to fit that bin better than the current item.

In this problem, you will implement the simple first-fit algorithm for online bin packing. This algorithm maintains a list of bins created so far, initially empty. It then loops through the items and places each in the first available bin with enough capacity left. If no bin is available, create a new empty bin at the end of the list and place the item there. This function should return the list of total bin sizes after the algorithm executes.

preferred_spot	capacity	Expected result
[3, 2, 3]	7	[5, 3]
[8, 6, 1, 9, 7]	9	[9, 6, 9, 7]
[9, 10, 6, 1, 4]	13	[10, 10, 10]
[11, 8, 3, 14, 10, 13, 14]	17	[14, 8, 14, 10, 13, 14]
[18, 12, 6, 10, 2, 3, 4, 17, 7]	53	[51, 28]

The first-fit algorithm can use up to 1.7 times as many bins as the optimal solution, for a particularly pathological sequence of items. Even with lookahead, no polynomial-time approximation algorithm can guarantee a better ratio than $3/2$ unless $P = NP$, in which case the problem could be solved exactly in polynomial time.

15. Count possible triangles

```
def count_triangles(sides):
```

The famous **triangle inequality** says that the sum of the two shorter sides of a triangle has to be at least equal to the longest side for that triangle to be geometrically possible. Given a list of possible lengths of triangle `sides`, this function should count the number of ways to choose three different sides to create a triangle. Since we care only about triangles that are proper and ignore triangles that degenerate into a line segment, we require the sum of the two shorter sides to be **strictly larger** than the longest chosen side length for the integer triple to count as a triangle.

Of course, you can solve this problem by iterating over all possible ways to choose three lengths from the sides list, either with three nested for loops or `combinations(items, 3)`. However, a cleverer design eliminates swathes of side-length triples that can't form a triangle, resulting in a significant speedup.

sides	Expected result
[1, 2, 3]	0
[1, 2, 4]	1
[3, 4, 7, 13, 28]	0
[5, 10, 12, 13, 14, 15]	19
[7, 14, 20, 25, 29, 33, 36, 39, 41, 42, 43, 44]	171

16. Arrow walk with me

```
def arrow_walk(board, position):
```

This cute little problem, taken from a recent Stack Overflow Code Golf post, has surprising depth in both theory and implementation. You are initially standing at the given position of a one-dimensional `board` encoded as a string of angle bracket characters. At each time step, you move one step left or right, depending on the direction of the angle bracket you are currently standing on. As you step away from a bracket, that bracket flips to point in the opposite direction. This function should return the number of steps needed to get off the board.

Since Python strings are immutable, you will definitely want to first convert the `board` to a list whose individual items can be mutated in constant time. Doing so will speed up your function by an order of magnitude compared to the naive implementation that keeps the board as a string throughout execution.

board	position	Expected result
'><>>'	1	5
'<<><<<'	1	2
'<>><>><>>'	3	18
'<<<><>><<>>'	0	1
'><>><<><<<<><<<<><<<<><>><>><<<<<'	22	400

Discrete math enthusiasts might enjoy proving that this process can never get stuck in an infinite loop, but must eventually terminate for any finite `board` and starting `position`.

17. Count Friday the Thirteenths

```
def count_friday_13s(start, end):
```

The problems in both this and the original 109 Python Problems collection are designed not to assume any knowledge about any special libraries and frameworks. However, dates and times are an important exception to this general principle. As you can see from the classic document “Falsehoods programmers believe about time”, a part of the larger “Awesome Falsehood” curated collection of such lists from various walks of life, anyone trying to roll their own date and time calculation functions is virtually guaranteed to mishandle many unexpected edge case situations. (For example, how many time zones do you think exist? Obviously 24, right? And all of them are an integer number of hours offset from each other?)

Instead of reinventing the wheel, delegate any calculations involving real-world dates and times to the appropriate standard library, such as Python's `datetime` module. To practice using the `date` and `timedelta` datatypes of this library, this function should calculate how many days from the given `datetime.date` object `start` to the similar `datetime.date` object `end`, inclusive, happen to be Fridays that fall on the thirteenth day of that particular month.

start	end	Expected result
<code>datetime.date(2024, 3, 4)</code>	<code>datetime.date(2024, 5, 21)</code>	0
<code>datetime.date(2024, 3, 4)</code>	<code>datetime.date(2024, 12, 13)</code>	2
<code>datetime.date(2022, 1, 28)</code>	<code>datetime.date(2026, 12, 14)</code>	9
<code>datetime.date(1986, 8, 29)</code>	<code>datetime.date(2035, 9, 12)</code>	84
<code>datetime.date(1987, 12, 3)</code>	<code>datetime.date(2064, 11, 10)</code>	130

18. Median filter

```
def median_filter(items, k):
```

A time series of measurement values that contains noise that occasionally produces random outliers is often run through some kind of **smoothing** process to create a series of values with the same general shape, but with the random outliers brought back closer to the actual trend. Of countless such filters, the k -median filter best suits this course because it can be implemented solely with element-order comparisons, without imprecise floating-point arithmetic.

Given a list of `items` that are such noisy measurements, create a result of the same length where each item has been replaced by the **median** value of that item and the $k - 1$ immediately preceding items. The median of the given sequence of values is defined to be the value that would end up at the center position if the list of those values were sorted. So that this definition is also meaningful for the first $k - 1$ items in the list, we pretend that the list has been implicitly prepended with $k - 1$ copies of its first item. To ensure the median element is unambiguous at each location, the parameter k is guaranteed to be odd.

items	k	Expected result
[2, 2, 3]	3	[2, 2, 2]
[33, -44, -9, -30, -3]	3	[33, 33, -9, -30, -9]
[-25, -28, -41, -15, 14, 46, 31, 34]	3	[-25, -25, -28, -28, -15, 14, 31, 34]
[62, -20, 49, 1, -59, 65, 102, -98, 41, -45, -40, -94, -63, -48, -61]	5	[62, 62, 62, 49, 1, 1, 49, 1, 41, 41, -40, -45, -45, -48, -61]

19. The magic knight of Muhammad ibn Muhammad

```
def magic_knight(n, items):
```

A magic square is an n -by- n grid of distinct integers where each row and column has an equal sum. Usually these numbers are integers from 1 to n^2 , making the magic square **normal**, but other constraints are not only possible; they also open doors to interesting problems about the existence of such squares. As documented, for example, in the book “The Zen of Magic Squares, Circles and Stars” by Clifford Pickover, nearly every weirdo throughout history has at some point dabbled in devising new methods to construct magic squares. (The number of different magic squares of order n is exponential with respect to n , more than enough for mathematicians of each new century to discover.)

In this problem, we look at the magic square construction method originally devised by Muhammad ibn Muhammad, based on the movement of a chess knight. If the x - and y -coordinates of an n -by- n board are denoted with integer coordinates from 0 to $n - 1$, the knight starts its route at the upper right corner $(n - 1, n - 1)$ and assigns the integers from 1 to n^2 to squares in the same order that it travels through these squares. In each move, the knight takes one step left and two steps down, wrapping around the edges whenever this move would take it off the board. If the square that the knight would step into has already been assigned a number, the knight continues from that square two steps directly to the left, again wrapping around the edge if necessary, into a square that by this construction is guaranteed not to have a number in it already.

This simple algorithm will always produce a magic square whenever the board size n is not divisible by 2 or 3. Furthermore, if n is not divisible by 5, the construction makes the sums of both diagonals equal to the same magic sum as the rows and columns. This function should fill in the n -by- n magic square in this manner. However, to keep the returned answers short, the function doesn't need to return the entire square but only the coordinates where the given `items` were placed, as proof of work that the entire route was correctly constructed along the way.

n	items	Expected result
5	[1, 10, 12, 14, 24, 25]	[(4, 4), (3, 1), (4, 2), (2, 3), (3, 3), (2, 1)]
19	[76, 86, 106, 259, 300, 361]	[(13, 1), (1, 0), (17, 17), (0, 15), (12, 9), (2, 1)]
295	[21024, 40895, 42272, 50760, 82217, 87025]	[(74, 138), (129, 221), (217, 122), (226, 256), (122, 177), (2, 1)]

20. Man of La Cancha

```
def jai_alai(n, results):
```

Back in the day, the hottest action in Miami took place at the jai alai frontons. The last time that this trad Chad sport showed up in mainstream culture was probably in the opening titles of *Miami Vice*. Even if we can't return to the golden days of tailfins, bowling shirts, and rock-hard pelotas, we can appreciate how this game lets multiple teams face each other simultaneously in an orderly fashion, without devolving into chaos.

The game of *jai alai* is played between n teams, one or two players per team. Before the game, these teams are randomly numbered from 0 to $n - 1$. Initially, team 0 holds the serve on the court, while teams 1 to $n - 1$ wait their turn in a queue, which is probably best implemented with the `Queue` standard library class. Each round, the first team in the queue enters the court to play against the team currently serving. The team that scores a point stays on the court, the other team goes back to the end of the queue, and the next team from the head of the queue enters the court to play the winner of the previous point.

Given the `results` of each point, this function should return the points that each team has scored so far. In these `results`, 'W' means the team holding the court scored the point, while 'L' means the team that came in scored the point.

n	results	Expected result
2	'WLWL'	[2, 2]
2	'LLLLLL'	[3, 3]
4	'WWWWWWWWW'	[9, 0, 0, 0]
4	'LLLLWWW'	[5, 1, 1, 1]
4	'LLLLLWLLL'	[2, 3, 2, 2]
5	'WWLLWWWWWLW'	[2, 2, 0, 2, 8]

Curiously, these rules make it maximally difficult for a colluding cabal of players to secretly choose the winner, in the style of tournament poker collusion teams or mob-controlled “point shaving” schemes in basketball. Combinatorics appreciators may enjoy the challenge of analyzing how many of the n teams need to secretly collude to throw (heh) the points to make one of this cabal the winner with high probability, over all possible random permutations of teams.

21. Nondeterministic finite automata

`nfa(rules, text):`

Deterministic finite automata (DFAs) from the earlier problem are a powerful tool for regular expression pattern matching. Still, their major downside is the potentially exponentially large number of states relative to the length of the regular expression. To keep the number of states reasonably small while still allowing efficient pattern matching algorithms, we can generalize these automata to **nondeterministic finite automata (NFAs)** so that for each possible pair of current state `s` and character `c`, the dictionary item `rules[(s, c)]` is not merely a single state, but a list of possible successor states that the nondeterministic automaton can move to after reading the character `c` while in the current state `s`.

This function should return the list of all possible states that the given NFA can end up in after processing the characters in the given `text`, and these possible final states should be listed in sorted ascending order. Instead of just keeping track of the current state with an integer scalar variable, your function now needs to maintain an entire set of possible states that the automaton could be in after having processed the current prefix of the `text`, and for each character of the `text`, compute the set of possible successor states from the set of possible current states.

Note also that it is quite possible for some item of `rules[(s, c)]` to be the empty list, which makes it possible for some NFA and `text` inputs that the automaton ends up having an empty set of possible final states.

rules	text	Expected result
<code>{(0, 'a'): [1], (0, 'b'): [1], (1, 'a'): [0], (1, 'b'): [1]}</code>	<code>'aab'</code>	<code>[1]</code>
<code>{(0, 'a'): [], (0, 'b'): [], (1, 'a'): [0], (1, 'b'): [1]}</code>	<code>'ab'</code>	<code>[]</code>
<code>{(0, 'a'): [1, 0], (0, 'b'): [], (0, 'c'): [0, 3, 2], (0, 'd'): [2, 0], (1, 'a'): [0], (1, 'b'): [2], (1, 'c'): [0, 3], (1, 'd'): [0], (2, 'a'): [0], (2, 'b'): [], (2, 'c'): [3, 1], (2, 'd'): [3, 0], (3, 'a'): [3, 2], (3, 'b'): [3], (3, 'c'): [], (3, 'd'): [0]}</code>	<code>'cacabaa'</code>	<code>[0, 1, 2, 3]</code>

22. Count monochromatic rectangles

```
def count_unicolour_rectangles(grid):
```

The body of a Python for-loop can consist of any statements, including for-loops nested for arbitrary depths to create wheels spinning within wheels. This problem lets you play with a whopping **four** levels of nested for-loops: two for the coordinates of the top-left corner of the rectangle, with either one or two more levels of nested loops than inside those outer two.

You are given a two-dimensional `grid` whose rows are strings made up from the alphabet "abcd" to represent four different possible colours in each cell. This function should count the number of axis-aligned rectangles in the `grid`, where all four corners of each rectangle have the same colour. That is, count the number of coordinates (x, y) and width-height pairs (w, h) with $w > 0$ and $h > 0$ so that all four cells (x, y) , $(x + w, y)$, $(x, y + h)$ and $(x + w, y + h)$ have the same colour.

grid	Expected result
<code>['aba', 'ccc', 'ada']</code>	1
<code>['bbbbcba', 'cdacaabc', 'acdcdcdc', 'cddacdab', 'cbaaaabd']</code>	5 (how quickly can you find them all?)

This problem connects to an interesting combinatorial question: filling an n -by- n rectangle with four colours without creating any monochromatic rectangles. For small values of n , you can solve the problem by filling the grid randomly, and whenever you find a monochromatic rectangle, choose one of its corners randomly and flip it to a different random colour. This approach doesn't work for larger values of n , since the solutions get too sparse among the exponentially many possible colourings for this headless chicken to have a snowball's chance of accidentally hitting them, so more finesse is needed. As discussed in Brian Hayes's post "The 17×17 Challenge", when $n > 18$, colouring the grid with four colours without creating at least one monochromatic rectangle is known to be impossible. Colourings without monochromatic rectangles have been constructed for all n up to 17, but the case of $n = 18$ remains tantalizingly open.

23. Count palindromic substrings

```
def count_palindromes(text):
```

Given a `text` string that consists of letters `a` and `b` only, count how many substrings of the `text` are **palindromes** whose length is at least three characters. A palindrome is a string that reads the same forward and backwards, such as `'babab'` or `'abba'`. If the `text` contains several equal palindromic substrings, each of these substrings should be included separately in the count.

To count the palindromes efficiently without having to loop through a quadratic number of substrings and then test for each such substring whether it is a palindrome, we get a more efficient solution by realizing that every palindrome has a **center** that is either a single character for a palindrome of an odd length, or two equal characters for a palindrome of an even length. Loop through all positions `i` of the `text` as potential centers of palindromic substrings in the `text`. To count all the palindromes with the center `i`, use a nested while loop to look at characters `k` positions left and right from the center `i`, with `k` initially set to 1. As long as the two characters that surround the center `k` steps away in both directions are equal to each other, increment both the total count of palindromes and the value of `k` for the next round of the inner loop. Once you have looked at all possible centers inside the `text` and added up all the palindromes emanating from them, return the total count accumulated throughout the function.

text	Expected result
'aa'	0
'ababa'	4
'bbbbba'	6
'aababaabab'	10
'bbbabbbbbbbbbbbbbbbbbbbbbbbbbbbbbbbbbbaab'	747
'a' * 1000	498501

24. Longest mirrored substring

```
def str_rts(text):
```

Right on the heels of the previous problem, here is a similar problem about substrings, taken from Jeff Erickson's collection of old homework and exam questions. Your function should find the longest string that occurs both **as it is** and **as its reversed version** somewhere inside the `text` of lowercase English characters so that these two occurrences do not overlap with each other. Since this substring is not necessarily unique, this function should return the length of the longest possible substring with this property. Note that, as the last example string shows, the first substring may well occur after the halfway point of the `text`.

For this lab, a simple brute-force solution that loops through all possible pairs of starting points for the two strings is not sufficient to pass the automated tests, so a faster solution is needed. One possible way to do this is to build a dictionary `pos` so that `pos[c]` is the list of the positions where that character `c` occurs in the `text`. For example, if this dictionary is constructed for the text "helloworld", then `pos['o'] == [4, 6]`. Once you have constructed this dictionary in a single loop over the `text`, iterating through all possible starting points of the longest substring and its mirror image should be a breeze. Perhaps the reader can think up even faster techniques to solve this problem.

text	Expected result
' '	0
'abcd'	0
' aa '	1
'u ft tu ft '	3
'ha hahex o xehah '	5
'g gotjpr j togz gotjpr'	4
'xcz xqwcduc xa axzcud cwqx'	10
'djettezzez zettezexeexezc zettezexe ev tveexezettez '	11

25. Square lamplighter

```
def square_lamps(n, flips):
```

Lamps arranged in an n -by- n square grid are all initially turned off. A series of `flips` is then performed so that each flip changes the state of every lamp in either a single row (that is, every lamp in that row that is on is turned off, and every lamp in that row that is off is turned on) or a single column. The rows and columns of the grid have been numbered from 1 to n . Each flip f with $f > 0$ applies to the row f , whereas a flip with $f < 0$ applies to the column $-f$. This function should compute how many lamps on the grid are on after completing the given sequence of `flips`.

We could solve this by tracking each lamp's state separately and tallying the lit lamps after all the `flips`. However, this brute-force solution would be horrendously inefficient for large n , both in time and in memory. To solve this problem in linear time, we first need to realize that the order in which the `flips` are performed is irrelevant for the end state of each lamp, and that two flips performed in the same row or same column cancel each other out as if those two flips had never been performed at all. Finally, permuting the row or column order after the flips does not affect the final tally of lit lamps.

With these insights, the invisible light bulb hovering over your head should suddenly turn on and illuminate your way to compute the final tally with integer arithmetic by tracking only the **parity** of flips for each row and column.

n	flips	Expected result
1	[1, -1]	0
2	[1, 2]	4
3	[-3, 2, 1, 3, 1, 1]	6
4	[-2, 3, -2, 2, -1, 2, -3, 4, -1, 1, 1]	8
7	[2, 2, 3, -4, -3, 6, 1, 6, -4, 5, 7, -7, 4, -6, 4, 2, -4, 5, 4, -7, 4]	25
10000	list(range(-10000, 0)) + list(range(1, 10001))	0

26. Pairwise lamps

```
def lamp_pairs(lamps):
```

A one-dimensional row of `lamps` is represented as a string where 0 denotes a lamp that is turned off, and 1 denotes a lamp that is turned on. In a single move, you can choose any two adjacent lamps that are in the same state (that is, either both lamps are on, or both lamps are off) and flip them together to the opposite state. What is the smallest number of moves after which the entire row of `lamps` is simultaneously turned on?

Since each flip preserves the parity of the number of lamps that are off, a solution can exist only if the initial number of lamps that are off is even. However, as you can see from the test case `'00010'`, this necessary condition is not sufficient: no sequence of flips can turn on all lamps simultaneously, so the function should return `None`.

To solve this problem, find the leftmost lamp that is off. If its immediate neighbour to the right is also off, flip both these lamps on, and continue with the rest of the row. Otherwise, keep going right until you find the first pair of lamps in the same state (either both on or both off), and use them to make the right neighbour of that leftmost lamp off, so you can flip both lamps on together. Implemented properly, this function should be a single linear loop that traverses the `lamps` in a single pass from left to right.

lamps	Expected result
'100'	1
'0110'	3
'00010'	None
'011000'	4
'00110111011111001110010010001011'	45
'0001000000000000100100110101011'	None

This problem is “Paired furnaces” in the excellent Kattis problem collection.

27. Accumulating merge

```
def accumulating_merge(items1, items2):
```

The `itertools` standard library module contains many useful functions for common transformations of sequences and pairs of sequences. However, it does not contain a function for the following artificial variation: **merging** the elements of two sequences, `items1` and `items2`, into a result that contains each element of both sequences exactly once.

This function should start by creating an initially empty result list, and then iterate through both sequences one element at a time, but not necessarily in lockstep. This function should keep track of the position it is at in each list (both initially at the beginning of the list), and the **accumulated sum** of elements for each list it has processed so far (both sums initially zero). The function should also remember which list is currently being processed, initially always the first list `items1`.

Each round, the function appends the next element of the list currently being processed to the result and adds it to the accumulated sum of that list. If processing of the current list reaches the end of that list or the accumulated sum of the current list becomes strictly larger than the accumulated sum of the other list, the function switches to process the elements of the other list, provided that the function has not yet reached the end of that list. Processing the elements one at a time continues in this same manner until the result list contains all elements from both lists.

items1	items2	Expected result
[1, 2, 3]	[4, 5]	[1, 4, 2, 3, 5]
[3, 6, 2, 14]	[5, 9, 11, 3]	[3, 5, 6, 9, 2, 14, 11, 3]
[15]	[1, 17, 14]	[15, 1, 17, 14]
[22, 17, 4, 16]	[1]	[22, 1, 17, 4, 16]
[11, 2, 8, 4, 16, 5, 19]	[9, 19, 18, 12]	[11, 9, 19, 2, 8, 4, 16, 18, 5, 19, 12]
[16, 8, 20, 17, 11, 16, 10]	[14, 11, 4, 5]	[16, 14, 11, 8, 20, 4, 5, 17, 11, 16, 10]

28. And they walk in twos or threes or more

```
def twos_and_threes(n):
```

The fact that any positive integer can be represented as a sum of distinct powers of two in exactly one way lets us represent integers as **binary numbers**. However, a more curious representation uses distinct terms of the form $2^a 3^b$ where a and b are natural numbers: binary numbers are a trivial special case of this with $b = 0$ in every term. For example, the integer 21 can be expressed as the sum $21 = 12 + 9 = 2^2 3^1 + 2^0 3^2$. Unlike with powers of two, this representation allows the terms $2^a 3^b$ to be chosen so that no chosen term is divisible by any other chosen term! The **constructive** proof Paul Erdős discovered gives us a recursive algorithm to compute these terms efficiently for any given n .

If $n = 2m$ is an even number, simply construct this sum for m , and then just multiply every term by two. Since the original terms were not divisible by each other, multiplying each term by two does not change the set of terms. On the other hand, if n is an odd number, find the largest b so that 3^b is at most equal to n . Then construct the sum for $(n - 3^b) / 2$, half of the remainder, and multiply the terms of this latter sum by two. These terms for the remainder plus the original 3^b are then the desired sum whose terms add up to n . Since 3^b clearly does not contain the prime factor two, whereas all the remaining terms do and their exponent of three must be lower than b , it is easy to verify that none of these terms can divide each other.

This function should compute the twos-and-threes representation of the given positive integer n and return the list of terms in descending order. To make the answers more concise and useful, each term $2^a 3^b$ should be given as a two-tuple (a, b) of its exponents of two and three, instead of the actual integer of multiplying these powers together.

n	Expected result
3	[(0, 1)]
4	[(2, 0)]
19	[(0, 2), (1, 1), (2, 0)]
2901	[(0, 7), (1, 5), (2, 3), (4, 1), (3, 2)]
771084	[(2, 11), (11, 2), (5, 6), (8, 4)]
2578298957581455	[(0, 32), (1, 30), (2, 29), (4, 25), (5, 24), (13, 19), (28, 2), (15, 17), (26, 3), (25, 5), (24, 6), (17, 15), (23, 9), (20, 12)]

29. Ladies and gentlemen, Conway Bitty

```
def conway_coin_race(a, b):
```

Two players each choose a bit sequence, expressed here as two strings `a` and `b`, and the player whose chosen bit sequence appears first in the series of repeated flips of a fair random coin (0 means heads and 1 means tails) wins the race. If both players choose a sequence of equal length, one might expect that both have the same 1/2 probability of winning, but surprisingly, this isn't so! For example, the sequence '100' beats the sequence '000' with probability 7/8, since the second sequence can win only if it appears at the very beginning of the coin flip series, after which it has no chance of winning, regardless of how the rest of the flips come up. (DUCY?)

An ingenious algorithm to determine the probability that the sequence `a` beats John Horton Conway devised the sequence `b`, because of course it was. This algorithm uses a **leading function** *lead*(`x`, `y`) defined between two bit strings `x` and `y`, not necessarily of the same length, such that *lead*(`x`, `y`) is a binary string the same length as `x`. For each position of the result, the bit in that position is on if and only if the **suffix** of the string `x` starting at that position is also a **prefix** of the string `y`. For example, the reader can quickly verify that *lead*(01101, 11010) = 01001. (To help you work with binary numbers, the Python `int` function to convert a string into an integer takes an optional base parameter, so that, for example, `int('10010', 2)` equals 18.)

Odds of sequence `a` winning over sequence `b` can now be expressed in **odds form**

$$\text{lead}(b, b) - \text{lead}(b, a) : \text{lead}(a, a) - \text{lead}(a, b)$$

which is then converted to a `Fraction` of the probability that the sequence `a` wins the race. For example, the odds of 67:29 would become the probability $67 / (67 + 29) = 67 / 96$.

a	b	Expected result (as <code>Fraction</code>)
'1'	'0'	1/2
'00'	'10'	1/4
'100'	'110'	1/3
'1101111111'	'10001101'	60/317
'01010010111'	'1110'	7/1024

As explained in “The Colossal Book of Mathematics” by Martin Gardner, the ordering between sequences of the same length is **intransitive**. Any sequence of three or more bits is beaten by the sequence that first places the complement of the second-to-last bit of the original sequence in front of it, then discards the last bit.

30. Lowest fraction between two fractions

```
def lowest_fraction_between(first, second):
```

Given two `Fraction` objects `first < second`, this function should compute a third `Fraction` object `result` strictly between the other two so that `first < result < second`, and that among all possible `result` fractions with this property, the returned `result` has the lowest possible **denominator**. (If multiple fractions with the same denominator happen to fit in the wide gap between `first` and `second`, return the one with the lowest numerator.)

To make the brute-force search for the result fractions more efficient, let us denote these three fractions by $a/b < e/f < c/d$. Multiplying each term by the so far unknown denominator f of the result (note that f is already known to be nonzero), this is equivalent to $(a f)/b < e < (c f)/d$. It therefore suffices to try out increasing values of f starting from one until you get to the value of f that is large enough to make the whole number parts of the two fractions $(a f)/b$ and $(c f)/d$ different. The numerator of the `result` then equals the whole number part of the first fraction $(a f)/b$ plus one, while the denominator equals the f that you ended up at.

first	second	Expected result (as Fraction)
<code>Fraction(7, 3)</code>	<code>Fraction(10, 3)</code>	3
<code>Fraction(49, 100)</code>	<code>Fraction(51, 100)</code>	1/2
<code>Fraction(11, 8)</code>	<code>Fraction(13, 8)</code>	3/2
<code>Fraction(11, 9)</code>	<code>Fraction(196, 153)</code>	5/4
<code>Fraction(127, 48)</code>	<code>Fraction(53345, 20112)</code>	45/17
<code>Fraction(5, 3)</code>	<code>Fraction(97, 57)</code>	17/18

This problem comes from the collection “Mathematical Diamonds” by the late great Ross Honsberger.

31. Reasonable filename comparison

```
def tog_comparison(first, second):
```

In his list of Ten Most Wanted Design Bugs in the “Bughouse” collection of usability blunders that have plagued computer users longer than most students reading this have been walking on this Earth, the usability expert Bruce “Tog” Tognazzini includes “Stupid sort” of listing file names in alphabetical order in a mechanistic manner. Since it is a fine summer day up here in 2023, I saw a screenshot of the Eclipse Java IDE listing the results of running JUnit test methods in that stupid sort order. The time has come to put a stop to this idiocy once and for all.

If the order comparison of filenames is performed on a per-character basis, the files `foo1.txt`, `foo10.txt` and `foo2.txt` end up being listed in this order, since the stupid string comparison used in the sorting algorithm does not realize that integers that consist of consecutive digit characters from '0' to '9' ought to be compared according to the arithmetic order between integers, instead of the **lexicographic** dictionary order of their characters. Comparisons between non-digit characters are done according to their Unicode code points, as the Python string comparison operators `<` et al. already do for individual characters.

This function should compare the two given text strings in this reasonable manner, returning `-1` if the `first` string is lower than the `second`, `+1` if the `first` string is higher than the `second`, and `0` if both strings are equal. Also, to keep this problem simple, the comparison function does not need to treat any leading minus sign as part of the number. What sort of perverts would number their files with negative numbers anyway?

first	second	Expected result
'foo10.txt'	'foo2.txt'	1
'foo1.txt'	'foo2.txt'	-1
'bar42qux99joe.txt'	'bar42qux100moe.txt'	-1
'qPX938.doc'	'qPX7105.doc'	-1
'60359Z30571.dat'	'60359Z30571.dat'	0

32. List the Langford violations

```
def langford_violations(items):
```

A list of $2n$ elements where each number $1, \dots, n$ occurs exactly twice is called a Langford pairing if, for every k from 1 to n , there are exactly k elements between the two occurrences of k in the list. For example, $[2, 3, 1, 2, 1, 3]$ is a Langford pairing for $n = 3$, as the reader can quickly verify. Langford pairings exist whenever n is a multiple of four or one less than a multiple of four. All Langford pairings for the given n can, in principle, be generated with a recursive backtracking algorithm, or a special algorithm can be used to rapidly produce just one such pairing for any given admissible value of n . However, the exact count of Langford pairings for a given n is a famous open problem in combinatorics; these exponentially growing counts remain unknown for $n = 31$ onwards.

Instead of generating Langford pairings, this function simply receives a list of `items` guaranteed to contain each number $1, \dots, n$ exactly twice, and should return a sorted list of all numbers that violate the Langford requirement in that list of `items`. You should implement this function to operate in a single pass through the `items`, instead of using two nested loops in a sluggish “Shlemiel” fashion. Keep track of which numbers you have already seen through the pass, and whenever you see some number the second time, append that number to the result list if the distance between its two occurrences is incorrect. After processing all the `items` in a single pass, return the sorted result list.

items	Expected result
[2, 3, 4, 2, 1, 3, 1, 4]	[]
[5, 8, 4, 1, 7, 1, 5, 4, 6, 3, 8, 2, 7, 3, 2, 6]	[]
[5, 8, 1, 1, 10, 2, 11, 6, 2, 9, 8, 4, 6, 7, 4, 10, 7, 3, 11, 9, 5, 3]	[1, 4, 5, 6, 7]
[14, 4, 1, 1, 2, 13, 15, 2, 11, 10, 12, 5, 10, 6, 9, 14, 8, 5, 7, 13, 6, 11, 4, 12, 9, 8, 7, 3, 15, 3]	[1, 3, 4, 10, 11, 15]
[3, 4, 1, 15, 3, 16, 4, 2, 16, 13, 12, 20, 5, 8, 14, 7, 10, 2, 5, 9, 18, 6, 1, 13, 19, 17, 11, 10, 6, 14, 12, 17, 19, 20, 8, 15, 7, 9, 11, 18]	[1, 2, 7, 8, 9, 12, 15, 16, 17, 19, 20]

Interested readers can check out the page “Langford’s Problem” at [Dialectrix](#), or the page “Langford’s Pairing” by [Susam Pal](#), for more information about this classic combinatorial problem.

33. Ten pins, not six, Dolores

```
def bowling_score(frames):
```

Even if we post-postmodern people will end up bowling alone as our society becomes more alienated and atomized, calculating the score in ten-pin bowling should still make for a fine exercise in Python programming. This function should compute the score for the list of `frames` encoded as strings. Strikes are encoded as 'X', misses as '-', and spares as '/' in each frame. For simplicity, we assume that no fouls occur during the game.

For readers unfamiliar with the notation, rules for scoring a bowling game are available on several online sites. Note also the special three-character form of the last frame when it contains either a strike or a spare.

frames	Expected result
['X', '4-', '54', '-2', '72', 'X', 'X', '5/', 'X', '--']	113
['2/', '34', '9/', '52', '4-', '11', '81', '8/', 'X', '6-']	99
['X', '7/', '-/', '31', '18', '11', '4/', '33', '43', '72']	93
['X', '2/', '9/', 'X', '8/', '6/', '25', '8/', '1/', '71']	150
['X', 'X', 'X', '6/', '2/', 'X', 'X', '2/', '4/', 'XXX']	214
['8/', '7/', '8/', '9/', 'X', 'X', '-9', '3/', 'X', 'XX7']	199
['42', '8/', '8/', '-8', 'X', 'X', '6/', '31', '5/', 'X51']	141
['21', '11', '8/', 'X', '61', '7/', '4/', 'X', 'X', '4/X']	147
['3/', '45', '2/', '7/', '2/', '8/', '11', '4/', '3/', '61']	119

34. Strict majority element

```
def has_majority(items):
```

This function should determine whether the given list of `items` contains a strict majority element, that is, an element that occurs more times than all other elements combined. Unlike the sloppy definition “at least half”, this definition makes the majority element unique if it exists.

This problem can be solved efficiently in two different ways. The first way uses a **counting dictionary** to tally the occurrence counts of each element in `items`, then iterates through the keys of the counter to check how many times each key occurred in `items`, keeping track of the element with the highest count so far. Nothing wrong with this technique, and Python makes it easy to implement for a good enough solution for this problem.

However, a cleverer way to crack this dusty old algorithmic chestnut is to notice that if you ignore any even-length prefix of `items` that contains k copies of some element along with k other elements, the majority element of the original list will necessarily also be the majority element of the rest of the list after ignoring these $2k$ elements. This second way solves this problem with two sequential for-loops that iterate over the items and operate **in place**, that is, use only a constant amount of bookkeeping memory in addition to the items themselves. There is a non-trivial chance that you can use this solution to dazzle your future job interviewer.

items	Expected result
[]	False
[42]	True
[1, 2, 1, 2]	False
[7, 5, 5, 5, 1, 5, 5, 1, 5]	True
[16, 2, 2, 2, 2, 4, 2, 2, 2, 2, 2, 2, 2, 2, 4, 4]	True
[3, 9, 20, 33, 33, 4, 20, 33, 3, 20, 20, 33, 4, 33, 3, 33, 33, 33, 3, 33, 33, 5, 3, 6, 3, 3, 3, 3, 3, 3, 3, 3, 6, 3, 3]	False

35. Add like an Egyptian

```
def greedy_egyptian(f):
```

As explained on the Wikipedia page “Egyptian Fractions”, ancient Egyptians expressed integer fractions as sums of **unit fractions** (that is, fractions of the form $1/n$ with $n > 0$) so that no denominator was ever repeated. For example, the ordinary fractions $2/3$ and $43/48$ can be expressed as the Egyptian fraction sums $1/2 + 1/6$ and $1/2 + 1/3 + 1/16$, respectively. (The simpler representation of $2/3$ as $1/3 + 1/3$ is not allowed here, as it repeats a denominator.) Every positive integer fraction can be expressed this way as a sum of unit fractions. In fact, any fraction can be expressed as a sum of distinct unit fractions in infinitely many ways, since any unit fraction $1/n$ can be further broken down into two unit fractions: $1/(n+1) + 1/(n(n+1))$.

Several algorithms for constructing a unit-fraction representation of a given fraction are known. This function has you implement the greedy algorithm first described and proven to always terminate by Fibonacci himself. Even though this algorithm works for any positive integer fraction, we will restrict our attention to fractions of the form a/b where $a < b$. If $a = 1$, the fraction is already a unit fraction representation, and the function can add n to the result and terminate. Otherwise, compute the smallest positive integer n for which $a/b > 1/n$ (you can do this with integer arithmetic without any loops), add n to the result list, and continue converting the remaining $a/b - 1/n$.

This function should return the list of integer denominators of the greedy unit fraction representation of the given `Fraction` object `f`, listed in ascending order.

f	Expected result
<code>Fraction(1, 2)</code>	[2]
<code>Fraction(2, 3)</code>	[2, 6]
<code>Fraction(43, 48)</code>	[2, 3, 16]
<code>Fraction(9, 19)</code>	[3, 8, 66, 5016]
<code>Fraction(4, 25)</code>	[7, 59, 5163, 53307975]
<code>Fraction(53, 71)</code>	[2, 5, 22, 977, 1271729, 2425940702433, 11770376583439808963536545]

As you can see from the last example, the greedy breakdown into unit fractions can sometimes end up quite unbalanced, even for relatively simple fractions. Other algorithms can produce more balanced breakdowns, and possibly produce sums with fewer terms.

36. Sorting by pairwise swaps

```
def pair_swaps(perm):
```

The parameter `perm` is a **permutation** of order n , that is, a list with each integer from 0 to $n - 1$ exactly once somewhere in it. (Being computer scientists, we start counting from zero, not one.) A single move consists of swapping any two elements of the permutation. What is the minimum number of **pairwise element swaps** necessary to sort this permutation to ascending order so that `perm[i] == i` for every position i in the list?

Any permutation can be sorted in at most $n - 1$ pairwise swaps by first swapping element 0 to its correct location if it isn't already there, then swapping element 1 to its correct location if it isn't already there, and so on. However, a smaller total number of swaps can be realized for some permutations by preferring swaps that move two elements into their final resting positions in a single move, or, if no such fortunate double swaps are available, choosing a swap that creates such a double swap for the next move. How can you organize and implement this computation most simply and efficiently?

perm	Expected result
[0, 2, 1]	1
[2, 0, 1]	2
[0, 1, 2, 3, 4]	0
[2, 4, 1, 0, 5, 3]	5
[2, 5, 0, 3, 1, 4]	3
[26, 13, 7, 4, 5, 14, 0, 2, 19, 16, 10, 17, 15, 27, 20, 8, 22, 21, 24, 9, 1, 11, 12, 25, 3, 23, 6, 18]	21

This problem is from the classic collection “The Moscow Puzzles: 359 Mathematical Recreations”.

37. Van der Corput sequence

```
def van_der_corput(base, n):
```

The Van der Corput sequence is an infinite low-discrepancy sequence of rational numbers that fills in the open unit interval $(0, 1)$. Any finite prefix of this sequence can be used as positions of sample points over the unit interval, possibly aided with some small random **jittering**, so that all regions of the interval get sampled fairly in the sense that the sampling does not concentrate on some parts of the interval and leave other parts without attention. The sequence is usually defined in base 2, but can, in principle, be defined in any base. The sequence's position numbering starts from 1, not 0. This function should compute the value at position n in the van der Corput sequence of the given base.

To compute the value in position n , start by expressing n in the given base. For example, in the binary van der Corput sequence, position $n = 6$ is 110_2 in binary. Reading the individual digits in this representation from lowest to highest, use each digit in position i inside the number as the coefficient of the corresponding negative power $(-1 - i)$ of the base. For base two, these negative powers would be $1/2, 1/4, 1/8, 1/16$, and so on. The value in position n is the sum of these negative powers multiplied by their respective coefficients. For example, the value at position 6 of the van der Corput sequence is $1/4 + 1/8 = 3/8$.

base	n	Expected result
2	1	Fraction(1, 2)
3	7	Fraction(5, 9)
4	9	Fraction(3, 8)
5	17	Fraction(13, 25)
42	$10^{**}10$	Fraction(23346056329, 230539333248)

Interested readers can check out the blog post “When Random Numbers Are Too Random: Low Discrepancy Sequences” at “The blog at the bottom of the sea” to learn more about the need and application of this and more advanced low discrepancy sequences in fair sampling over one- and higher-dimensional intervals.

38. Maximal repeats

```
def maximal_repeats(text, k):
```

Define a **maximal repeat** as a substring that occurs inside the `text` at least in two different locations (these occurrences are allowed to overlap, as can be seen in the second example `text` in the table below), and that some two of these occurrences are followed by different characters. (We can conceptually pretend that every string is followed by an invisible “end-of-string” character that cannot occur inside the string itself. Every substring that occurs at the end of the `text` therefore automatically qualifies as such a continuation mismatch.)

This function should return the list of all maximal repeats whose length is at least `k` characters that can be found inside the given `text`. These repeats should be listed in descending order of length; the repeats of equal length should be listed in ascending alphabetical order.

text	k	Expected result
'babbababba'	2	['babba', 'abba', 'bab', 'bba', 'ab', 'ba']
'adfadababab'	2	['abab', 'bab', 'ab', 'ad']
'bcfefefefbcfefefef'	3	['bcfefefef', 'cfefefef', 'fefefef', 'efefef', 'fefef', 'efef', 'fef']
'gagagaagagdagdabdag'	5	[]

This problem is from Rosalind, an excellent problem collection that emphasizes the algorithms and techniques used in bioinformatics. Their suggested solution uses an advanced data structure called a suffix tree, which would make this function fast even for text strings with millions of characters. Here, you can pass the automated tester within the time limit with a reasonable approach that uses only Python's bread-and-butter sets and dictionaries.

39. It's a game, a reflection

```
def abacaba(n):
```

As illustrated on the Wikipedia page ABACABA, this is one of the simplest fractal patterns to appear in all sorts of familiar (or semi-familiar) mathematical sequences and patterns. Defined over an alphabet starting from A, the ABACABA string S_k of order k can be defined recursively so that the string S_0 equals just A, and the string S_k is the concatenation of two copies of the string S_{k-1} with the k :th character of the alphabet stuck between those copies to separate them. For example, S_1 equals ABA, S_2 equals ABACABA, S_3 equals ABACABADABACABA, and so on.

Since each S_{k+1} contains the previous S_k as a prefix, it makes sense to define the infinite ABACABA string as the limit of this process as k approaches infinity. This function should return the character at position n in the infinite ABACABA string, counting from zero. It is easy to prove by induction that the string S_k contains exactly $2^k - 1$ characters, so your function should not try to construct this string explicitly, but compute the result with integer computations.

To return a meaningful result for arbitrarily large alphabets (not just the 26 letters of English), this function should return an integer: 0 for A, 1 for B, 2 for C, and so on.

n	Expected result
0	0
1	1
2	0
3	2
$2^{42} - 1$	42
$2^{42} - 1 + 2^{41}$	41
$2^{41} - 1 - 2^{41}$	41
10^{100}	0

40. Reverse a Fibonacci-like sequence

```
def reversenacci(i, n):
```

A sequence of integers is **Fibonacci-like** if, after the given first two elements, each element equals the sum of the previous two elements. Of course, the original Fibonacci sequence starting with (1, 1) is the most famous and important of such a series. Still, other choices of the first two elements give us an infinity of sequences, such as the Lucas sequence starting with (2, 1).

At this stage, it would be anticlimactic to ask the reader to generate the Fibonacci-like sequence from the given two elements. A more interesting question is whether a Fibonacci-like sequence exists with the value n at position i , such that all previous elements are positive and non-decreasing. (Position counting starts from zero.)

The reader may enjoy thinking of ways to speed up this search compared to brute-forcing all possible starting values and generating the sequence from each start value up to position i .

i	n	Expected result
5	8	True
4	12	False
4	13	True
4	14	True
17	28048	False
17	28047	True

41. Recamán sequence

`def recaman(n):`

The Recamán sequence is a famous integer sequence in which iterating a simple rule produces chaotic and unpredictable behaviour. The first element of this sequence at the position $n = 0$ equals zero. After that, the element a_n at position n equals $a_{n-1} - n$, provided this difference is positive and $a_{n-1} - n$ has not already appeared in the sequence so far. If this is not the case, the element a_n equals $a_{n-1} + n$, even if this value has already appeared in the sequence.

This function should compute and return the element a_n located at the position n of this sequence. The automated tester will call this function with values from 1 up to one million, by far the largest number of test cases given to any function in this problem collection! Since these test cases are given to your function in strictly increasing order, your function needs to remember only the most recent element a_{n-1} (initially zero), and the set of values that have appeared in the sequence so far.

Since storing a Python set of up to one million elements seems excessive, you can use a useful programming technique in this situation. When maintaining a **monotonic** set of positive integers where no element ever gets removed from the set (this is a common scenario in problems where the set is used to remember integers that the function has seen and done), in addition to the set itself, you should also maintain a scalar integer variable x that remembers the smallest positive integer that is not a member of the set. The set itself then only needs to remember the elements that are larger than x , since all values less than x are automatically members of the set. When you add the threshold element x to the set, you can keep incrementing x as long as it is a member of the set, and remove the elements you skip to free up those memory bytes for larger elements later.

n	Expected result
1	1
10	11
100	164
1000	3686
10000	199508
100000	2057164

The Numberphile video “The Slightly Spooky Recamán Sequence” discusses the curious properties of this seemingly simple sequence.

42. Mian–Chowla sequence

```
def mian_chowla(n):
```

The Mian–Chowla sequence is another interesting sequence of increasing positive integers defined by a simple rule from which a chaotic level of complexity emerges. This sequence starts with 1 in the first position ($n = 0$). After that, the Mian–Chowla sequence can be defined in two equivalent ways; here we use the one that makes implementation easier: **all differences between two distinct elements in the sequence must be pairwise distinct**. That is, if the sequence contains two distinct elements a and b , it cannot contain any other two distinct elements c and d so that $b - a = d - c$. The sequence begins with 1, 2, 4, 8, giving the first impression of simply being powers of two, but then continues with 13, 21, 31, 45, 66, 81, 97, 123, 148, ...

This function should return the element in position n in the sequence, using zero-based indexing. The automated tester is guaranteed to give your function the values of n in strictly increasing order, so your function should store its progress in generating this sequence in some global variables outside the function, so that your progress doesn't always get erased between the function calls and force you to again and again generate the entire sequence from the beginning like some modern-day Sisyphus. Same as in the previous Recamán sequence, use a Python `set` to remember all the differences you have already encountered in the sequence, aided by a variable that keeps track of the smallest positive integer that you haven't seen as a difference of two elements so far.

n	Expected result
0	1
1	2
19	475
46	4297
99	27219
300	524307
400	1145152
500	2107005

43. Stern–Brocot path

```
def stern_brocot(x):
```

The Stern–Brocot tree is an infinite binary tree that contains every positive rational number exactly once, always in simplest reduced form, as a node. The tree's construction starts with the lower and upper limits that do not appear in the tree. The lower limit is initially set to $0/1$, which represents zero, and the upper limit is set to $1/0$, which represents positive infinity.

The root node of the Stern–Brocot tree contains the median of the two initial limits. Since computing this median gives the numerator $0 + 1 = 1$ and the denominator $1 + 0 = 1$, the root node contains the rational number $1/1$, the unity from which this entire panoply emerges. To compute the value stored in the left child of the given node that holds the value x , replace the current upper limit with x and keep the lower limit as is. Symmetrically for the right child, replace the current lower limit with x and keep the upper limit as is. The infinite subtree constructed to hang from that node then contains all rational numbers between that lower and upper limit.

Instead of constructing the entire tree up to some particular depth, this function should return the list of values that appear in nodes on the linear path from the root node to the node that contains the rational number x . The **binary search algorithm** for achieving this is explained in the section “Mediants and binary search” on the Wikipedia page for the Stern–Brocot tree.

x	Expected result
Fraction(1, 1)	[Fraction(1, 1)]
Fraction(4, 5)	[Fraction(1, 1), Fraction(1, 2), Fraction(2, 3), Fraction(3, 4), Fraction(4, 5)]
Fraction(4, 3)	[Fraction(1, 1), Fraction(2, 1), Fraction(3, 2), Fraction(4, 3)]
Fraction(35, 32)	[Fraction(1, 1), Fraction(2, 1), Fraction(3, 2), Fraction(4, 3), Fraction(5, 4), Fraction(6, 5), Fraction(7, 6), Fraction(8, 7), Fraction(9, 8), Fraction(10, 9), Fraction(11, 10), Fraction(12, 11), Fraction(23, 21), Fraction(35, 32)]

44. Carryless multiplication

```
def carryless_multiplication(a, b):
```

The blog post “I Could Carry Less” by Brian Hayes links to the article “Carryless Arithmetic Mod 10” that defines interesting addition and multiplication operations for integers. These operations work otherwise the same as the familiar addition and multiplication, except that addition and multiplication of individual digits are performed **reductively** modulo 10, meaning that any carries over to the next column are simply ignored.

This weird but mathematically perfectly consistent operation gives birth to interesting integer sequences that the curious reader might want to check out after completing this problem. (Especially the section explaining how to properly define **primality** in this algebra should be highly educational.) However, in this problem, you are only asked to implement the carryless multiplication operation between the given nonnegative integers a and b .

It is probably useful to first implement a similar function `carryless_addition` to add two given integers in the carryless fashion, if only to get the gist of the general behaviour of this wacky arithmetic. This function will also simplify the logic of the `carryless_multiplication` proper.

a	b	Expected result
12	0	0
2	5	0
7	666	222
643	59	417
123	456	43878
59999888866666	66666228887466	48264684815151480266668626

The reader may also compare and contrast this operation to Problem 87 in the original 109 Python Problems collection, “Lunatic multiplication”, where instead of ignoring the carry digits, there is never any carry since addition and multiplication were defined as taking the maximum and minimum of the digits being added and multiplied, respectively.

45. Mū tōrere boom-de-ay

```
def mu_torere_moves(board, player):
```

Mū tōrere is a traditional Māori board game with surprisingly complex strategy veiled behind deceptively simple rules. As with the Oware move problem in the original collection, the minimax algorithm for this game will have to wait until the course CCPS 721 Artificial Intelligence I, but we can already implement the move generator to list the possible moves for the given `player` on the current `board`. Being computer scientists, we of course also immediately generalize this game from eight points to $2k$ points around the shared center for arbitrary $k > 0$.

The reader should consult the linked Wikipedia page for the setup and rules of this game, or the page “Mu Torere” at Mayhematics for a deeper analysis of game positions. Pieces never capture each other or get removed from the board, but on his turn, a player must move any one of his pieces to a neighbouring space. The stalemated player who can't make any legal moves loses. To prevent the game from ending after the first move, a piece can move into the empty center space only if at least one opponent's piece is next to the piece being moved.

The `board` is given as a string of $2k + 1$ characters, each either 'B' or 'W' for black and white pieces, or '-' denoting a space. The first $2k$ characters list the pieces around the center, and the last character gives the contents of the center. This function should return a list of possible boards the `player` can reach with a single move from the current `board`. To ensure that the result is unique, this list should be returned alphabetically sorted.

board	player	Expected result
'-BW'	'W'	['WB-']
'-BW'	'B'	['B-W']
'BBWW-'	'B'	['-BWWB', 'B-WWB']
'WBWWBB-'	'W'	['-BWWBBW', 'WB-WBBW', 'WBW-BBW']
'BBBBBW-WW'	'B'	[]
'BBBBBW-WW'	'W'	['BBBBW-WWW', 'BBBBWWW-W', 'BBBBWWWW-']
'BWWBWBWB-'	'B'	['-WWBWBWBB', 'BWW-WBWBB', 'BWWBW-WBB', 'BWWBWBW-B']

46. A place for everything and everything in its place

```
def insertion_sort_swaps(items):
```

Of all the simple sorting algorithms, **insertion sort** is not only the shortest but also provably the most efficient. Starting from a singleton prefix that contains only one element and is therefore sorted by default, insertion sort maintains and grows a **prefix of sorted elements**. The outer loop of insertion sort iterates through the `items` from left to right. The inner loop grabs the **element** just past the current sorted prefix and repeatedly swaps it with the preceding larger elements until it either ends up at the beginning or is preceded by an element at most equal to it.

For example, the list `[5, 2, 7, 4, 1]` first becomes `[2, 5, 7, 4, 1]`, which then becomes `[2, 5, 7, 4, 1]`, which becomes `[2, 4, 5, 7, 1]`, which becomes `[1, 2, 4, 5, 7]`. Since your function could just pass the buck of sorting the `items` to its native `sort` method, and our automated tester doesn't have X-ray glasses to see how exactly your function does the job, your function must return the total number of adjacent element swaps during the insertion sort as **proof of work** of actually having gone through the rigmarole of the insertion sort algorithm.

items	Expected result
<code>[1, 2, 3, 4]</code>	0
<code>[2, 1, 0, -1]</code>	6
<code>[4, 3, 1, 1, 0, -1, 2, -3]</code>	23
<code>[3, 2, 21, 7, 23, 5, -25, 7, -18, -21, 8, 11, 23, -19]</code>	43
<code>list(range(10000, 0, -1))</code>	49995000

For the given list of `items`, the number of required element swaps is an important combinatorial quantity known as the **inversion count**: the number of element pairs in unsorted order. A sorted list has zero inversions, and swapping two adjacent elements can decrease the inversion count by at most one. Therefore, any sorting algorithm that can only swap adjacent elements must necessarily take at least linear time in the inversion count of the list. Insertion sort requires exactly that much, plus the linear time needed to take a butcher's at each element even before any comparisons and swaps, and is thus optimal among sorting algorithms limited to compare and swap adjacent elements only.

47. Scoring a tournament bridge hand

```
def bridge_score(strain, level, vul, doubled, made):
```

Once the dust has settled after both bidding and play of a bridge hand are over, it's time to add up the score while muttering under your breath about the mistakes that your partner again did, probably just to annoy you. To keep this problem simple, we will focus only on successful contracts. Given the `strain`, `level`, `vulnerability` and the `doubled` status of the contract (as one of the possibilities `'- '`, `'X'` and `'XX'`), along with the level of tricks made, this function should compute the score for the declarer.

The rules for scoring the successful bridge contract can be found in many places on the Internet, for example, in the Wikipedia page “Bridge Scoring”. Being gentlemen instead of some kind of animals that wallow in their own filth, we score according to the rules of tournament bridge, not rubber bridge, which is better suited for kitchen bridge players.

strain	level	vul	doubled	made	Expected result
'notrump'	1	True	'- '	2	120
'spades'	3	False	'X'	4	650
'diamonds'	3	True	'- '	5	150
'hearts'	4	True	'- '	4	620
'clubs'	6	False	'- '	7	940
'notrump'	7	True	'XX'	7	2980

48. Manimix

```
def manimix(expression):
```

The given `expression` string consists of digit characters placed inside properly nested and balanced pairs of parentheses and square brackets. Define a string of digit characters inside square brackets as evaluating to the maximum of those digits, and a string of digit characters inside parentheses as evaluating to the minimum of those digits. For example, the string `'[283]'` evaluates to 8, and the string `'(283)'` evaluates to 2. Nested expressions are evaluated inside out, so that `'([28](56))'` evaluates as if it were `'(85)'`, which then evaluates to 5. This function should evaluate the given `expression` under these rules. There is no limit to how long an `expression` can be and how deep its nesting can reach.

Probably the simplest way to solve this problem is to use a list as an initially empty **stack**, and loop through the characters of the `expression` one at a time from left to right. Any character other than a closing parenthesis or a closing square bracket is simply appended to the stack. Whenever the loop encounters a closing parenthesis or square bracket, it should use an inner loop to pop characters from the stack up to and including the corresponding opening parenthesis or square bracket, and push the minimum (if parentheses) or the maximum (if square brackets) of the popped digits into the stack. After processing the entire expression, the stack will contain the final answer as its only element.

expression	Expected result
'(123456)'	1
'[(45)(27)]'	4
'(88(0(41)))'	0
'[5[0[5743]](7(816)8942)8]'	8
'[[(872)(824)30(6565)][5379][[11][89]][(92([9170](0295)44))(1[00742])(35391))[6[2507]746][[[1701]9(1557)]5]90]4]'	9

49. Count distinct substrings

```
def count_distinct_substrings(text):
```

A nonempty substring can be sliced from the given `text` of length n in $n(n - 1)/2$ different ways, for all possible ways to choose the start and end of that slice. There is exactly one substring of length n , namely the `text` itself, two substrings of length $n - 1$, three substrings of length $n - 2$, and so on, down to the n substrings of length one. If all characters in the `text` are distinct, all these substrings will also end up being distinct, but if the `text` contains repeated characters, some of these substrings will be equal.

This function should compute how many different nonempty substrings the given `text` contains. This problem can be solved by looping through all $n(n - 1)/2$ possible slices of text, adding them to a result set one by one, and returning the size of that set as the final answer. However, this approach is horrendously inefficient when the `text` contains many repeated characters and substrings, making most of them equal.

A better solution springs forth from the realization that any substring of length $k - 1$ in `text` must necessarily be either a prefix or a suffix of some substring of length k . (DUCY?) The function should therefore maintain the set of all substrings of length k , initially containing only the `text` itself, for the starting state $k = n$. We can easily compute the set of all substrings of length $k - 1$ by looping through each substring `s` in the set of known substrings of length k , and adding both strings `s[1:]` and `s[:-1]` to the set of substrings of length $k - 1$ currently being constructed. After completing the set of all substrings of length $k - 1$, add the size of this set to the current tally, and go to the next round to compute the set of all substrings of length $k - 2$.

text	Expected result
'ab'	3
'bbbb'	4
'fegg'	9
'lfvffffffv'	32
'eokoooooooo'	33
'tboobooobbbb'	61
'ab' * 10000	39999

50. Replacement with perfect hindsight

```
def mmu_lru(n, pages):
```

The topic of page replacement algorithms in computer **virtual memory management**, whose intricacies we explore further in the later course CCPS 530 Operating Systems, can be illustrated with an analogy that gives us an interesting coding problem in this first programming course. You are carrying a binder that has n **slots** available, each slot initially empty and able to carry one **page**. You are given a list of `pages` to process in this exact order. A nonnegative integer identifies each page, and the same page can appear multiple times. Before being processed, the page must first be in a binder. If that page is not already in the binder, you need to place it in a slot, costing you one page fault.

If the binder contains unused slots to hold the current page, there is no problem. However, if the binder is full and there are no unused slots, one of the pages in the binder must be discarded to create an empty slot. An interesting problem in memory management is choosing which page to evict to minimize the expected number of future page faults. Following the general principle that the best predictor of a system's future behaviour is its past behaviour, a straightforward discipline is to remove the **least-recently used page** as the best guess of the page that won't be used in the future, in the style and spirit of Marie Kondo. This function should compute the total number of page faults for the given sequence of pages under the least-recently-used page discipline and return that count.

n	pages	Expected result
2	[3, 1, 4, 4, 3, 3]	4
2	[4, 1, 2, 2, 4, 1, 2]	6
1	[3, 1, 1, 3, 1, 3, 3, 3]	5
2	[4, 0, 8, 4, 0, 8, 8, 4, 8]	7
5	[8, 6, 3, 8, 3, 6, 8, 3, 3, 6]	3
4	[5, 8, 10, 5, 10, 8, 1, 10, 5, 10]	4
3	[1, 12, 4, 4, 4, 11, 4, 1, 1, 8, 1, 4]	6

51. Replacement with perfect foresight

```
def mmu_lru(n, pages):
```

After solving the previous problem in the theme of page replacement algorithms in virtual memory management, we can see that making an educated guess of which page to ditch to make room for the current page can be difficult for anybody who is not a descendant of Madame de Thèbes and blessed with the gift of perfect foresight of future events. However, for this problem, let us pretend we have that perfect foresight, so we already know all the future pages that will be processed.

Instead of throwing out the least-recently used page to make room for the newcomer, we throw out the page whose next use is the farthest in the future. Of course, if a page in the binder will never be needed again, it's automatically the one that gets tossed. This function should compute and return the total number of page faults over the given pages under this discipline.

n	pages	Expected result
2	[3, 1, 4, 4, 3, 3]	3
2	[1, 5, 2, 1, 5, 5, 1, 1]	4
2	[4, 0, 8, 4, 0, 8, 8, 4, 8]	5
3	[1, 12, 4, 4, 4, 11, 4, 1, 1, 8, 1, 4]	5
4	[12, 12, 5, 15, 16, 6, 6, 12, 12, 16, 6, 12, 6, 16, 12]	5

Discrete math enjoyers can later prove that this prescient approach is guaranteed to minimize the total number of page faults over the processing of the given list of `pages` in all possible cases. (Take any optimal solution, and show that replacing some other choice with the choice done here can never make that solution worse; in other words, this problem has the **greedy choice property**.)

52. When there's no item, there's no problem

```
def stalin_sort(items):
```

Unlike insertion sort and other such non-Lysenkovian **comparison sorting** algorithms devised by the CIA to weaken the spirit of the hardy people of Brutopia, **Stalin sort** is a famous joke sort algorithm that wastes no time with bourgeois considerations of “fairness” or “equality”. Inspired by and named after the original anti-capitalist social justice warrior, Stalin sort operates in one brutally effective loop through the list of `items`, giving every subversive element, that is, those smaller than their predecessor, a one-way train ticket to gulag. After this loop, the elements that remain in the list will be in sorted order! For example, given the list `[3, 2, 4, 1]`, Stalin sort would remove the wreckers 1 and 2 (note that after removing element 1, the next element 2 is compared to its surviving predecessor 3), leaving the sorted items `[3, 4]`.

Stalin sort can be turned into an inefficient but working sorting algorithm by allowing the subversive elements to return after having learned their lesson in the gulag. After each purge, the subversive elements are returned to the list in the order they were sent away. For example, the previous list becomes `[2, 1, 3, 4]` after the first round of purges. The same algorithm is then applied to this new list as many times as required to achieve and verify the utopia, here reached once comrade 1 has returned from his second re-education journey.

Instead of returning the sorted list, which you could do simply by calling the list function `sort`, the automated tester expects this function not to be alienated from the toils of production but to prove having done this work by returning the number of rounds needed to arrange the `items` to sorted order and verify that order. However, you should be aware of that wretched wrecker, Shlemiel, who wants to sabotage this glorious effort by making each purge run in quadratic time by hiding the inner loop inside a seemingly innocent list operation, such as `remove`. Proper praxis of the glorious revolution will steamroll through the items in linear time with a single for-loop.

items	Expected result
<code>[1, 2, 3, 4]</code>	1
<code>[2, 1, 0, -1]</code>	4
<code>[4, 3, 1, 1, 0, -1, 2, -3]</code>	6
<code>list(range(10000, 0, -1))</code>	10000

It should be clear and rational to everyone outside the oppressor class that this algorithm is guaranteed to terminate after a finite number of rounds. Don't be afraid to place your finger on every item and ask why, comrade!

53. Boxed away

```
def fewest_boxes(items, weight_limit):
```

Each item in the given list has a weight that is a unique positive integer, listed in ascending order. You are given an unlimited supply of identical boxes, each of which can hold up to two items, provided that the weight of these items is at most equal to the given `weight_limit` of an individual box. No box can hold three or more items, even if those items together happen to fit within the weight limit. Your function should calculate the minimum number of boxes needed to pack away all the `items`.

The bin packing problem, the general version of this problem in which each box can hold not just two but any number of items within the weight limit, is known to be NP-hard, so no efficient polynomial-time algorithms are believed to exist for it. Even the special case of the partition problem that asks to determine whether two boxes, each with a weight limit equal to one half of the total weight of the `items`, suffice to pack away all the given items exactly is already NP-hard, so any algorithm aiming to solve this problem in the general case is doomed to juggle these `items` between the two boxes through an exponential number of partitions. However, hardcoding the capacity of each box to at most 2 items should eliminate all this massive branching and allow the optimal solution to be extracted in far less time.

items	weight_limit	Expected result
[41]	42	1
[4, 5, 8]	11	2
[3, 6, 12, 14, 21]	24	3
[12, 23, 29, 42, 53, 54, 56]	58	5
[15, 18, 31, 37, 38, 50, 57]	74	4
[19, 23, 44, 69, 72, 77, 89, 118, 141]	162	5
[18, 20, 51, 73, 118, 159, 172, 213, 233, 235, 282]	308	6

54. 'Tis but a scratch

```
def knight_survival(n, x, y, k):
```

A lone chess knight starts at the position (x, y) on the generalized n -by- n chessboard, position indexing again zero-based. The knight must make k random moves, each chosen with equal probability $(1/8)$ from the eight possible knight moves. Any move that ends outside the board makes the knight fall to his doom. This function should compute the exact probability that the knight survives on the board for the duration of the entire series of k moves, and return that exact probability as an exact `Fraction` object.

The survival probability $S(n, x, y, k)$ can be computed recursively. The two base cases are when the coordinates (x, y) are outside the board, so the survival probability is zero regardless of k , and when the coordinates are inside the board with $k = 0$, so the survival probability is one. Otherwise, the survival probability is the sum of the eight probabilities $S(n, x', y', k - 1)$, multiplied by the transition probability $1/8$ to **normalize** these probabilities to add up to one, where x' and y' go through all the possible squares reachable by a single knight's move from the square (x, y) .

This recursion can be expressed as a recursion aided by some handy `@lru_cache` magic. Alternatively, you can solve this with **dynamic programming** by creating an n -by- n table of probabilities, all initialized to the same value 1. From this table, which contains the survival probabilities for the current k (initially 0), you can compute the next-highest table of survival probabilities for $k + 1$.

n	x	y	k	Expected result
4	3	0	3	5/128
9	8	2	1	1/2
9	6	5	1	1
10	1	6	3	119/256
11	3	5	11	124974217/536870912
13	1	12	13	25928405477/549755813888

This problem is “Knight Probability in Chessboard” on LeetCode. Of course, very few companies would probably interview candidates using problems that involve probabilities and recursion. But if you are lucky enough to be hired by such a company, aim to make yourself indispensable there.

55. Do or die

```
def accumulate_dice(d, goal):
```

Right on the heels of the previous problem about recursively computed probabilities, a fair d -sided die is repeatedly rolled, and the results are added together until the total is at least the given `goal`. The final total can therefore be anything from `goal` to `goal + (d-1)`, inclusive. This function should compute the exact probabilities of the final total ending up at each of those values, and return the answer as a list of `Fraction` objects with exactly d elements.

Calculating the exact probabilities correctly requires a bit more finesse. To do this, we define $P(s, k)$ as the probability that the total sum of dice equals exactly s after k rolls. The base cases of this recursion are $P(0, 0) = 1$ and $P(s, 0) = 0$ for $s > 0$. The general probability $P(s, k)$ can be computed from the probabilities $P(s', k-1)$ where s' goes through all sums that a single roll could turn into the total of s . Note that the sums greater than or equal to the `goal` are **absorbing states** where further rolls do not change the achieved total. (As you can see in the third row in the table below, it is also possible for d to be greater than `goal`, and the problem is still perfectly well-defined with a definite answer.)

d	goal	Expected result
2	4	[Fraction(11, 16), Fraction(5, 16)]
4	5	[Fraction(369, 1024), Fraction(305, 1024), Fraction(225, 1024), Fraction(125, 1024)]
3	8	[Fraction(3289, 6561), Fraction(2200, 6561), Fraction(1072, 6561)]
6	11	[Fraction(106442161, 362797056), Fraction(87771985, 362797056), Fraction(65990113, 362797056), Fraction(50655625, 362797056), Fraction(34445005, 362797056), Fraction(17492167, 362797056)]
8	19	An eight-element list whose first element equals Fraction(31945752545707729, 144115188075855872)

This problem was inspired by a problem titled “Rolling a Die” in the collection “Mathematical Morsels” by Ross Honsberger, asking the reader to prove that the `goal` is always the most probable total to achieve with these rolls, regardless of d . It's always nice to have numerical confirmation of intuitively acceptable things you have only proven, not actually tried out.

56. Arithmetic skip

```
def arithmetic_skip(n):
```

A person is standing at the origin (0) of the sequence of integers, extending infinitely in both positive and negative directions. As his first move, he can take either one step to the left or one step to the right. As his second move, he can take either two steps to the left or two steps to the right. In general, as his k :th move, he can take either k steps to the left or k steps to the right. It can be proven that you can reach any target number n from the origin with these moves. This function should compute the smallest number of moves necessary to reach the number n .

This problem can be solved with a **breadth-first** search, but that approach would be far too slow to pass the test-case torrent from the automated tester. Instead, you should first define a recursive helper function `can_reach(n, k)` that returns `True` if it is possible to reach the goal n from the origin in precisely k moves, and returns `False` otherwise. The base case is $k = 0$, in which case the answer is the same as testing whether $n = 0$. For efficiency, also add a base case for when n is too large to reach from the origin in k moves, using the well-known formula for the **arithmetic sum** of the first k positive integers. Otherwise, n can be reached from the origin in exactly k moves if and only if either one of the predecessor states $n - k$ or $n + k$ can be reached from the origin in $k - 1$ moves.

After writing this utility function, the actual function should keep calling `can_reach(n, k)` for ever-increasing values of k until it receives the answer `True`, and return the k that achieves this. Since the same recursive subproblems will be reached many times during the entire run of automated test cases, you should use the `@lru_cache` decorator on your `can_reach` function to store these subproblem solutions for future use. To cut the size of this cache in half, note that $-n$ can be reached from the origin in k moves if and only if $+n$ can be so reached.

n	Expected result
0	0
2	3
-15	5
38	11
215	21
-999	45

57. Carving Egyptian fractions

```
def carving_egyptian(f):
```

The earlier problem, “Add like an Egyptian”, asked you to express a positive integer fraction as a sum of distinct unit fractions using the greedy algorithm to achieve this breakdown. This problem continues with the same task, but with a different method based on the observation that if any fraction a/b that is not a unit fraction can always be trivially expressed as $(a - 1)/b + 1/b$, immediately giving us one unit fraction $1/b$ and leaving a fraction with a smaller numerator to be converted.

The difficulty is keeping the generated unit fractions distinct. This can be achieved using the following technique, which maintains the set of unit fractions accepted into the `result` set so far, along with the `buffer` list of unit fractions waiting to be added. Each time that you carve out the new unit fraction $1/b$ in the outer loop, append b to the `buffer`. Then, as long as the `buffer` is nonempty, pop out any one element n from the `buffer`. (The result will not be affected by the choice of element to pop, so you might as well pop out the last element of the list.)

If n is not already in the `result` set, add it there. If n is already in the `result` set, remove it and continue based on the parity of n . If n is even, append $n/2$ to the `buffer`. If n is odd, append two values $(n + 1)/2$ and $n(n + 1)/2$ to the `buffer`. Either way, process the `buffer` in this manner until it becomes empty, and continue carving the remaining fraction $(a - 1)/b$.

f	Expected result
Fraction(1, 2)	[2]
Fraction(2, 3)	[2, 6]
Fraction(43, 48)	[2, 4, 8, 48]
Fraction(9, 19)	[3, 15, 19, 48, 4560]
Fraction(4, 25)	[7, 91, 163, 52975]
Fraction(53, 71)	[3, 5, 15, 18, 23, 45, 71, 160, 320, 1035, 1278, 102240, 204480]

As the reader may note by comparing the above example results with those from the same arguments in the greedy Egyptian fractions problem, the carving method tends to produce more terms. Still, these terms tend to be more balanced so the subdivision doesn't contain ugly slivers.

58. Largest square of ones

```
def largest_ones_square(board):
```

Here is another classic job interview chestnut that has a cute solution: either recursion aided by the `lru_cache` memoization magic, or dynamic programming to fill in the memoization table directly. Your function receives a `board` of n -by- n elements, where each row is given as a string of characters '0' and '1'. This function should find the largest square made of ones on the `board` and return its side length.

Once again, some “Shlemiel” would solve this with four levels of nested loops to iterate through all possible top-left corners and areas of that square. We can solve this problem more efficiently by defining a recursive function $S(row, col)$ that returns the side length of the largest square whose **bottom-right corner** is at the given *row* and *column*. The base case is $S(row, col) = 0$ whenever that position contains a '0' or is out of bounds of the `board`. For positions that contain a '1', the value of $S(row, col)$ is computed as $1 + \min(S(row - 1, col), S(row - 1, col - 1), S(row, col - 1))$, which depends on the solutions of its west, northwest and north neighbours. DUCY?

board	Expected result
<pre>['0000', '1111', '0111', '1111']</pre>	3
<pre>['00111', '11010', '11001', '00110', '10110']</pre>	2
<pre>['0010000', '1101111', '1111111', '0111111', '0111111', '0111111', '0111111', '1011110']</pre>	4

59. Infection affection

```
def infected_cells(infected):
```

On the infinite two-dimensional grid of integers, some cells (x, y) have become infected with a virus. Once infected, there is no cure, and an infected cell remains infected forever. The rule for the spread of the infection is that as soon as two or more of a cell's four neighbours (in the main compass directions) become infected, that cell also becomes infected. This function should compute the total number of infected cells at the end,, once the infection has reached all the cells it will spread to.

To efficiently determine all the cells the infection will reach, you should maintain a set of all cells infected so far and a queue of cells waiting to be examined. Initialize both of these to contain the infected cells. While the queue is not empty, pop any one cell from there, and look at its four neighbours. For each uninfected neighbour, count how many of its neighbours are infected. If this count is at least 2, that neighbour becomes infected, so you add it to both the queue and the infected set. Once the queue has become empty, the infection can spread no further, and you can return the size of the infected set as your final answer.

infected	Expected result
<code>[(0, 0), (1, 1), (1, 3), (3, 2)]</code>	16
<code>[(-1, 1), (1, -1)]</code>	2
<code>[(0, 1), (-1, 0), (0, -2)]</code>	8
<code>[(0, 1), (2, -1), (5, -1), (4, -2), (0, -1)]</code>	24
<code>[(8, 1), (5, 4), (7, 0), (6, -1), (5, 6), (5, 3), (3, 2), (4, 1)]</code>	48
<code>[(x, x) for x in range(1000)]</code>	1000000

This particular problem was inspired by a famous discrete math exercise of proving claims with **invariants**. Can you prove rigorously that for all the cells of an n -by- n box to become infected in this manner, at least n cells inside the box must be initially infected? This problem is tricky, since any hand-waving “But surely at least one cell in each row must initially be infected, *maaan*” is not correct. However, once you get the “Aha!” moment of the lightbulb turning on above your head to see that the total **perimeter** of all infected cells is an invariant of the infection process that cannot increase in any infection step, it immediately follows that to completely infect the n -by- n box that has a perimeter of $4n$, at least n cells of perimeter 4 per cell must be initially infected.

60. Prize strings

```
def prize_strings(n, late_limit, absent_limit):
```

The problems offered on Project Euler are generally beyond the reach of undergraduate students. Even if the problem itself were simple, adding up the results for gazillion possible inputs tends to require a mathematical knack beyond what you can expect at this level of study. However, some problems are simple enough to be included here and generalized beyond the original problem setting. A case in point is problem 191, “Prize Strings”.

The attendance record of a student over n days can be expressed as an n -character string consisting of characters 'O' for “on time”, 'A' for “absent,” and 'L' for “late”. For example, the attendance record 'OOALAO' means that the student was on time the first two days, absent the third, late the fourth, absent again on the fifth, and on time the sixth day. The student gets a plaster reward if they are late fewer than `late_limit` days in total, and are never absent for any `absent_limit` consecutive days. This function should calculate how many different attendance record strings allow the student to receive the reward.

This problem can be solved with a nested recursive function, aided by `@lru_magic` to keep the running time from blowing up exponentially. This recursive function needs three parameters to describe the state reached today: the number of days remaining, the the total number of late days accrued so far, and the days of absenteeism that immediately precede the current day. The function should then compute the number of ways to extend the sequence for each of the three possible ways to spend the current day (on time, late, or absent), and add up the results for these three ways for the answer returned to the previous level of recursion.

n	late_limit	absent_limit	Expected result
4	2	3	43
30	2	3	1918080160
5	4	3	211
20	2	6	10444797
31	1	4	13752712373
59	1	5	221333601789363969

61. String shuffle

```
def is_string_shuffle(first, second, result):
```

Analogous to shuffling two halves of a deck of cards, characters from two strings can be shuffled together to create a result string that contains all characters from both original strings, while keeping the order of characters from each original string unchanged. This problem is not quite as simple as it seems when the same characters can appear in both strings. For example, two strings "hello" and "world" can be shuffled together in exponentially many ways to create many result strings such as "worldhello", "hewolrldlo" or "wohrellldo".

This function should determine whether the `first` and `second` string can be shuffled together to produce the given `result` string, and return that answer.

first	second	result	Expected result
'knickerbocker'	'kamala'	'kkanacmilareebockkr'	False
'chandler'	'mxuzptlk'	'chmandlxuzerptlk'	True
'laurie'	'grayson'	'grlaysaurione'	True
'joy'	'travis'	'jtoaryvis'	False
'ouagadougou'	'ellie'	'ouagadoeulglouie'	True
'turk'	'urban'	'urbanturk'	True

62. Weak Goodstein sequence

```
def goodstein(n, k):
```

Another nice Project Euler problem suitable for this collection is problem 396, “Weak Goodstein sequence”. As you may know, other than the fact that we have ten fingers, there is nothing magical about the base ten that we use to represent integers; analogous to how every species in the universe calls itself “human” in its own language, every base is “base 10” from its own point of view! The Python built-in function `int` lets the caller supply the base as the second argument, which defaults to 10, but this argument has a maximum value that is far too small for this problem. You will therefore have to write the base conversion routines back and forth all by yourself using repeated remainder and integer division operations.

The **weak Goodstein sequence** starts from the given value $g_1 = n$ (for simplicity, we index its position from 1), and continues with the rule that involves base conversions. To compute the next element g_{i+1} from the current element g_i , first compute the list of digits of g_i in base i . Then interpret these very same digits as a number in the next higher base $i + 1$, and subtract one. For example, as shown on the Project Euler problem page above, the sequence for $n = 6$ starts as 6, 11, 17, 25, 35, 39, 43, 47, 51, 55, 59, 62, 65...

One might expect these values to keep increasing forever, since, of course, the same digits interpreted in base $i + 1$ represent a larger number than they do in base i . However, subtracting one at each step is enough to rein in this growth, so for any starting value n , the weak Goodstein sequence must eventually fall to zero and stay there! This function should return the value at position $k + 1$ in the weak Goodstein sequence starting from n .

n	k	Expected result
1	1	1
3	3	3
5	3	15
16	187	13328639
33	642	513618635782
41	504	32844064606486

The rule that defines this wacky sequence is still pretty weak sauce compared to the real deal Goodstein sequence, whose termination time can and will reach some pretty dizzying heights.

63. Markov distance

```
def markov_distance(triple1, triple2):
```

As explained on the Wikipedia page “Markov Numbers”, the positive integer solutions to the polynomial equation $x^2 + y^2 + z^2 = 3xyz$ have a curious property that if (x, y, z) is one solution triple that satisfies that equation, then $(x, y, 3xy - z)$ is another solution for that equation. Since this equation is symmetric, any of the three values of x , y , and z can be permuted into the role of the last term, giving us up to three more solution triples from any known solution triple (x, y, z) . For example, the solution triple $(1, 5, 13)$ gives us three other solutions $(1, 2, 5)$, $(1, 13, 34)$ and $(5, 13, 194)$. Furthermore, these solutions form an infinite tree depicted on the Wikipedia page, so that any solution triple can be reached from any other solution triple by traversing up and down this tree.

Your function should compute how many steps apart the two given solution triples, `triple1` and `triple2`, are in this tree. Since the solution space forms a tree, we don't need the full power of breadth-first search to find the shortest path in arbitrary state spaces; we can solve this problem by climbing up the tree. Start by initializing a step counter to zero. While the two solution triples are not equal to each other, replace the solution triple with the largest maximum value with its parent triple in this tree, and increment the step counter by one. Once the two solution triples become equal, stop and return the final value of the step counter.

triple1	triple2	Expected result
(1, 2, 5)	(1, 2, 5)	0
(13, 34, 1325)	(2, 29, 169)	5
(13, 1325, 51641)	(2, 29, 169)	6
(14701, 56399710225, 2487396418774357)	(5, 13, 194)	8
(985, 4360711162037, 12885900008113189)	(4400489, 33809973988413085, 446341255878891162918658)	14

64. A very graphy caterpillar

```
def is_caterpillar(edges):
```

An **undirected graph** consists of n nodes numbered from 0 to $n - 1$. Each node is connected to some other nodes by undirected edges, so if node u is a neighbour of node v , then node v is also a neighbour of node u . These edges are given as a list to your function so that each element `edges[u]` lists the neighbours of node u .

This function should determine whether the given graph is a caterpillar tree, a special type of tree consisting of a linear path of nodes with all other nodes one step away. To determine whether the graph is a caterpillar tree, first check whether it is connected, which is easiest to do by running a breadth-first search from any node and verifying that it reaches all nodes in the graph. After checking for connectedness, you can choose any one of the equivalent definitions from the Wikipedia page and implement that check as Python code. The easiest one is probably the one that removes **leaf nodes** (that is, nodes connected to only one other node) from the original graph, and then checks whether the remaining graph is a path graph. The latter check is trivial: look at the number of neighbours of each node.

edges	Expected result
[[1, 2, 3, 4], [0], [0], [0, 5], [0], [3]]	True
[[2, 5], [2], [0, 1], [5], [], [0, 3, 6], [5]]	False
[[3], [2], [1, 3, 4, 5], [0, 2, 6, 7], [2], [2], [3], [3]]	True
[[2, 6, 7], [7], [0, 4, 5], [5], [2], [2, 3, 7], [0], [0, 1, 5]]	False

Most problems on graphs become utterly trivial when restricted to caterpillar trees. One interesting exception is the bandwidth minimization problem that stubbornly remains NP-complete even when restricted to caterpillars. (The seemingly obvious algorithm of numbering the spine nodes from left to right and numbering each leaf hanging from a node with consecutive values does not always produce the optimal numbering.)

65. St. Bitus' Dance

```
def super_tiny_rng(seed, n, bits):
```

Because we handle individual integers stored in computer memory as their smallest elementary particles at the hardware level, introductory Python courses rarely cover low-level bitwise arithmetic and its operators. However, students familiar with the concept and bitwise operators in general might enjoy implementing a function described in the post “A Super-Tiny Random Number Generator” from “The Blog at the Bottom of the Sea.”

Every **pseudorandom number generator** maintains an internal **state**, initialized to the given `seed` value. Whenever a new random integer is requested, the generator performs some operation on its current state to produce the new state, and returns some part of the new state as that random quantity. This hardcoded operation applied to the state has been carefully chosen so that, even though it is perfectly deterministic, so that the same `seed` will always produce the same random numbers, it is highly **chaotic**, so that its results behave as if they were produced by a genuine random process, shaking and twisting like some poor soul caught in a *danse macabre*.

The super-tiny generator in the blog post above uses an internal state of a 32-bit integer. It updates its internal state `x` each time using the assignment `x += (x*x) | 5`, where the `|` operator is the **bitwise inclusive OR**, to make the bits scramble like in mad chaotic dance. The bitwise operation to chop the resulting `x` to 32 bits in a Procrustean manner is left as an exercise for the reader. Since the lowest bits of the state are highly non-random, each update uses only the highest-order bit as the random bit; this update must be performed `bits` times to produce a result integer with the given number of `bits`. Use the first random bit produced this way as the highest-order bit of the generated random number, and the later bits to fill in the lower-order bits.

Starting from the given `seed` value, this function should produce a list of `n` random unsigned integers, each one consisting of the given number of `bits`.

seed	n	bits	Expected result
831769172	3	8	[39, 170, 1]
2376066489	3	12	[2138, 3320, 2731]
2528054762	5	31	[1425110951, 1066116851, 1565670465, 1613085173, 1921709156]
1419340007	7	64	[2372072775304310267, 14757952967543449990, 2275344065262535408, 668817182494655816, 14695084848187285985, 6287545658816032331, 13919832216633589446]

66. Digit string partition

```
def digit_partition(digits, n):
```

This function should determine whether it is possible to partition the given `digit` string into pieces so that the sum of these pieces, when added as integers, equals the goal value `n`. For example, the `digit` string `"123456789"` can be broken into pieces `123 + 4 + 56 + 789` to sum to 972, and into pieces `12 + 34567 + 89` to sum to 34668.

This problem can be solved recursively by looping through all possible first numbers, and for each one, recursively checking whether the remaining digits can be partitioned to reach the original goal minus that first number. Alternatively, you can loop through all possible ways to choose the last number in the partition, and try to partition the preceding digits to add up to the remaining goal. However, unlike most other recursions in this collection, subproblems don't tend to repeat through the recursive search, so using `@lru_cache` would waste both time and memory. You should also ensure that your function handles the zero digits correctly, regardless of whether these zeros are at the beginning, middle or the end of the digit string.

digits	n	Expected result
'00010203000'	6	True
'1000055'	100010	True
'26389'	26389	True
'0453962'	1337	False
'482951'	4880	True
'0024644418'	24644417	False
'625950539'	1196	True
'601372126390244791992'	5940825159181	False
'2413331523846696659'	523846937993	True

67. Hofstadter's figure-figure sequences

```
def hofstadter_figure_figure(n):
```

In his now-classic tome “Gödel, Escher, Bach”, Douglas Hofstadter defined a bunch of wacky, chaotic and self-referential sequences that make for interesting coding problems. This problem looks at Hofstadter's figure-figure sequences, two complementary sequences R and S of increasing positive integers. Counting the position numbers from zero, as us computer scientists made of sterner stuff, the sequences are initialized with $R(0) = 1$ and $S(0) = 2$. After that, $R(n) = R(n - 1) + S(n - 1)$, and $S(n)$ equals the smallest positive integer that has not appeared in either sequence so far. The reader can easily verify that these two sequences start as

R : 1, 3, 7, 12, 18, 26, 35, 45, 56, 69, 83, 98, 114, 131, 150, 170, 191, 213, 236, 260, ...

S : 2, 4, 5, 6, 8, 9, 10, 11, 13, 14, 15, 16, 17, 19, 20, 21, 22, 23, 24, 25, ...

The R -sequence grows approximately quadratically, while the S -sequence patiently crawls behind to collect all the numbers that the R -sequence skipped over on its merry way.

Given the zero-indexed position n , this function should return a tuple of two values, $R(n)$ and $S(n)$. The automated tester will provide your function values of n from 0 up to 100,000 in order. To pass the test within the time limit, you should cache the sequences generated so far, along with any other information you choose to store, so you can quickly determine the smallest positive integer that has not yet made its debut.

n	Expected result
0	(1, 2)
6	(35, 10)
30	(602, 38)
49	(1509, 59)
100000	(5028615179, 100432)

68. [Be]t[Te]r [C][Al]l [Sm][Al]l

```
def break_bad(word, symbols):
```

The title card of the television series *Breaking Bad* famously stylized the show title by boxing chemical element symbols within the title. In this problem, you will write a function that performs a similar transformation for the given word made of lowercase English letters, boxing the chemical element `symbols` between square brackets. Your task is to break the word into chemical symbols to minimize the cost function, where each unboxed letter costs two points, each single-letter chemical symbol costs one point, and each two-letter chemical symbol costs zero points. If two different breakdowns of the given word have the same cost, return the one that uses a two-letter symbol earlier than the other solution.

This problem is best solved with **dynamic programming**. Given the n -letter word, start by creating two lists, `cost` and `move`, so that `cost[i]` gives the count of unbroken letters in the best solution for breaking the substring `word[i:]` into chemical elements, and `move[i]` says what symbol the optimal solution for `word[i:]` starts with. You then fill these two lists from right to left, deciding for each `move[i]` and its ensuing optimal `cost[i]` based on the latter parts of these lists filled in the previous rounds of this loop. After both lists, `cost` and `move`, have been filled, a simple left-to-right iteration can construct the optimal solution.

word	Expected result (using the real chemical elements)
'no'	'[No]'
'overgrows'	'[O][V][Er]gr[O][W][S]'
'multinode'	'm[U]l[Ti][No]de'
'reprised'	'[Re][Pr][I][Se]d'
'oxysalicylic'	'[O]x[Y][S][Al][I][C][Y][Li][C]'
'felina'	'[Fe][Li][Na]'

69. Spiral matrix

```
def spiral_matrix(n, row, column):
```

A **spiral matrix** of order n is an n -by- n numerical grid that contains each integer from 0 to $n^2 - 1$ exactly once, arranged in a clockwise spiral starting from the top-left corner. For example, the spiral matrix for $n = 5$ would look like this:

0	1	2	3	4
15	16	17	18	5
14	23	24	19	6
13	22	21	20	7
12	11	10	9	8

This function should return the element at the given row and column of the spiral matrix of order n , with both rows and columns numbered from 0 at the top-left corner. As always, your function can assume that row and column numbers fall properly inside the matrix.

Writing out the successive values to form the spiral matrix until you reach the goal position is very inefficient for large n and will cause the automated tester to reject your function for being too slow to solve the torrent of test cases. You need to think of shortcuts to speed up this evaluation. You can start by noticing that snug as a bug in a rug, inside a spiral matrix of order n there is a smaller spiral matrix of order $n - 2$ whose starting number is offset by $4(n - 1)$, the number of border elements surrounding this smaller spiral matrix.

n	row	column	Expected result
1	0	0	0
2	1	1	2
4	2	2	14
5	2	2	24
25	18	13	481
10001	5000	5000	100020000
10000001	5000000	5000000	100000020000000

70. Baker–Norine dollar game

```
def baker_norine_dollar_game(edges, balance):
```

An **undirected graph** consists of n nodes numbered from 0 to $n - 1$. Each node is connected to other nodes by undirected edges, so if node u is a neighbour of node v , then node v is also a neighbour of node u . These edges are given as a list to your function so that each element `edges[u]` lists the neighbours of node u .

In the **Baker–Norine variant** of the classic chip-firing game, each node has a coin balance that can initially be positive, negative or zero. A node is **ready to fire** if its balance is at least equal to its **degree**, that is, its number of neighbours. A single move consists of choosing any node that is ready to fire, and moving one coin from that node to each of its neighbours, maintaining the overall total of coins throughout the graph. This function should determine whether there exists some sequence of moves starting from the initial balance to make the balance of every node positive.

This problem can be solved with a simple graph search in a **state-space graph** whose nodes are the coin balances of the original nodes. Your function should maintain the list of states that constitute the current **search frontier**, initially containing only the initial balance. Repeatedly pop one state out of this frontier (doesn't really matter which one), and loop through all possible nodes that can fire in that state. For each node that ready to fire, compute its resulting successor state. If the balance of every node is positive in this successor state, your function can return `True`. Otherwise, if that successor state has not been previously encountered during the search (use a separate set to keep track of this, storing each state in that set as an immutable **tuple**), append that successor state to the search frontier. Should the search frontier become empty once all reachable state configurations have been exhausted without reaching a solution, your function can return `False`.

edges	balance	Expected result
[[1, 2], [2, 0, 3], [1, 0], [1]]	[7, -1, -1, -1]	False
[[1, 2], [4, 0], [0, 3], [4, 2], [3, 1]]	[6, -1, -1, 3, -2]	True
[[5, 1, 7, 3, 6], [0, 3, 2, 5], [5, 4, 1, 7, 6], [4, 1, 0, 6, 5], [3, 2, 7], [1, 0, 3, 2], [2, 0, 3], [4, 0, 2]]	[7, 9, -1, -1, 10, -1, -1, -1]	True

There is some pretty deep graph-theoretic math in this problem, for anyone interested in such stuff.

71. Total covered area

```
def squares_total_area(points):
```

You are given a list of `points` located in the nonnegative quadrant of the infinite two-dimensional integer grid. Your task is to process these points in order so that, using each point on its turn as the **top right corner**, you construct the **largest possible axis-aligned square** that does not reach across the x - or y -axis into the negative coordinate values and does not share a nonzero area with any of the previously created squares. This function should add up and return the total area covered by the squares constructed in this manner.

For example, a reader equipped with a pen and some grid paper can quickly verify that processing the points `[(6, 2), (5, 3), (9, 4)]` would first create a 2-by-2 square whose top-right corner is at `(6, 2)`. Then, add a 1-by-1 square with its top-right corner at `(5, 3)`. Finally, add a 3-by-3 square with its top-right corner at `(9, 4)`. Adding the areas of these squares gives the total area $4 + 1 + 9 = 14$. Note how the order in which the `points` are given can sometimes make a big difference; had these points been `[(9, 4), (6, 2), (5, 3)]`, the resulting squares would have added up to the total area of $16 + 0 + 9 = 25$ instead.

The automated tester will give your function enough test cases with the point x - and y -coordinates large enough so that any brute force solution that loops through the integer grid points inside each square to diligently add up these points one by one is guaranteed to run out of time before reaching the finish line of the test. The correct, efficient logic uses only integer arithmetic and comparisons to quickly reduce the squares to their final sizes, and it would not care much much even if these point coordinates were in the octillions.

points	Expected result
<code>[(5, 3)]</code>	9
<code>[(3, 3), (5, 5), (7, 3)]</code>	17
<code>[(4, 5), (3, 4), (9, 6)]</code>	41
<code>[(2, 2), (5, 5), (6, 2)]</code>	17
<code>[(0, 8), (6, 0), (0, 10), (2, 9)]</code>	4
<code>[(10, 2), (9, 8), (5, 11), (2, 16)]</code>	53
<code>[(3, 20), (26, 6), (11, 24), (14, 33), (21, 2)]</code>	190
<code>[(11, 81), (80, 99), (76, 32), (97, 144), (159, 14), (145, 28), (124, 30), (24, 157)]</code>	8467

72. Balsam for the code

```
def measure_balsam(flasks, goal):
```

You are a medieval balsam merchant equipped with a tuple of `flasks` of their capacities in ounces, listed in descending order. The largest flask is initially full of balsam, and the smaller flasks are empty. Your task is to measure exactly `goal` ounces of balsam into any one of these flasks in the smallest possible number of moves. Since these flasks are unmarked except for their capacities, you can't measure the desired pour amount just by eye and feel; each move must keep pouring balsam from some nonempty flask into some other flask until either the flask that you pour from becomes empty, or the flask that you pour into becomes full. This function should return the smallest number of moves to measure exactly `goal` ounces into some flask, or `None` if this measurement is not possible to achieve in any number of moves with the given `flasks`.

This problem is perhaps easiest solved using **breadth-first search**. Initially, the list of reachable states contains only the initial starting state. As long as none of your reachable states contains a flask with exactly `goal` ounces of balsam in it, loop through these reachable states to produce the list of **successor states** that can be reached from the previous row of states in one pouring move. You should also maintain the set of all states that you have already seen at any time during the search, and ignore such previously seen states should they happen to be generated again, to prevent your search from running around in a circle like some mangy dog chasing its own tail. Eventually, one of two things must happen: either a desired goal state appears in the reachable states, or the list of current states becomes empty, which lets you conclude that the desired `goal` was impossible to measure using the given `flasks` to begin with.

flasks	goal	Expected result
(8, 5, 1)	7	1
(10, 7, 2)	9	7
(5, 3, 3)	1	None
(12, 8, 3, 1)	10	3
(15, 11, 8, 6, 5, 4)	9	1
(27, 22, 21, 20, 20, 10, 9, 9)	23	5

The original version of this problem asks three medieval balsam thieves to divide the stolen flask of balsam equally, using three given empty flasks of different sizes, so each thief can scurry off in a different direction with exactly one third of the balsam.

73. Flip those trips

```
def trip_flip(switches):
```

In front of you lies a finite row of `switches`, each with four possible values: 0, 1, 2, or 3. In a single move, you can choose any three consecutive switches whose values are all different, and flip all three switches together to the fourth possible value. For example, one possible **move** on the row `[0, 2, 1, 3, 2, 3]` turns it into `[0, 2, 0, 0, 0, 3]`.

These flips can be performed in any order, not necessarily strictly left to right. However, no matter how the initial `switches` are positioned or how you choose your moves, you will inevitably reach a dead end after a finite number of flips, so no further moves are possible.

Your function should compute the maximum number of moves that can be performed for the given `switches`. Note that two flips on disjoint triples of switches are **commutative**: they can be performed in either order without affecting the outcome. Therefore, if your previous flip in the flip sequence constructed so far occurred at position i , the search for the next flip can start at position $i - 2$ instead of the beginning of the list.

switches	Expected result
[3, 0, 2, 2, 1]	1
[3, 3, 0, 2, 0, 3, 2, 0]	4
[3, 0, 1, 0, 1, 2, 0, 1, 0, 2]	6
[3, 1, 2, 2, 3, 1, 3, 1, 0, 2, 0, 3, 2, 3, 3, 2]	13
[2, 3, 0, 1, 3, 3, 0, 3, 2, 1, 2, 3, 3, 2, 1, 1, 2]	7
[1, 2, 0, 2, 2, 3, 0, 3, 0, 0, 2, 0, 3, 0, 2, 1, 0]	12

This sparkling gem of a problem comes from a recent tweet by Joel David Hamkins, whose writings on logic are highly recommended not just for computer science undergraduates aiming for graduate studies in theoretical computer science, but also for anyone looking for serious yet accessible material on these topics. The tomes “Lectures on Philosophy of Mathematics” and “Proof and the Art of Mathematics” are basically top-notch undergraduate courses on their own. The tweet asked for a proof that this process eventually terminates. This is easiest to see by noticing that the leftmost position that ever gets flipped in the sequence will be flipped only once and never again, and the result follows by induction on the remaining positions and the flip sequence.

74. Maximal disk placement

```
def place_disks(points, r):
```

This function should place as many non-overlapping disks of radius r on the two-dimensional plane as possible, when the list of `points` gives you the possible centers that these disks can be placed on, and return the number of disks placed. We count two **kissing** disks as overlapping, so two disks centred on (x_1, y_1) and (x_2, y_2) overlap if $(x_2 - x_1)^2 + (y_2 - y_1)^2 \leq 4 * r * r$.

To speed up the recursive search for the optimal placement, you should first preprocess the `points` by creating a dictionary that associates each center point with the list of other points where a disk would overlap the disk placed at that center point. Your recursive backtracking search should maintain the set of points where a disk can be placed without overlapping any previously placed disks. Each time you place a disk at a point, eliminate those overlapping disks from future consideration. Be careful not to bring back a point that has been eliminated at multiple levels of recursion until the recursive search has backtracked all the way up to the level where that point was first eliminated.

points	r	Expected result
[(4, 10), (0, 5), (6, 11)]	2	2
[(10, 4), (-4, 0), (4, 13), (0, 16)]	3	3
[(10, 4), (-4, 0), (4, 13), (0, 16)]	3	3
[(41, 17), (35, 11), (27, 17), (46, 12), (57, 3), (69, 15), (61, 24), (10, 18)]	6	6
[(x, y) for x in range(10) for y in range(10)]	3	4

75. Nice sequence

```
def nice_sequence(items, start):
```

♪♪ *Many n, came here as numbers, many n, will pass this way; Many n, will count the factors, as they fill the longest chain; Many n, are smooth and clique-y, many n, are here to stay; Many n, won't see the base case, when it ends the longest chain...* ♪♪ – Paul Anka, *The Longest Chain*

This problem idea is taken from “Problems with a Point” by William Gasarch and Clyde Kruskal. Using the positive integer `items` given in sorted order so that each item can be used at most once, create the longest possible sequence so that after the given `start` value, each element is always some integer multiple or a factor of the immediate previous element. Since this sequence might not be unique, the function needs to return its length.

This problem can be solved with a recursive function that takes the sequence constructed so far and tries all possible ways to extend it from the current last element. It might be useful to preprocess the `items` into a dictionary that maps each item to the list of other items that are its multiples or factors, to speed up looping through the local possibilities during the recursion. The reader may also think of other ways to speed up the algorithm so that the tests complete within the time limit.

items	start	Expected result
[2, 8, 12, 26, 108, 132, 264]	132	6
[3, 9, 55, 95, 106, 190]	106	1
[4, 5, 11, 17, 22, 24, 52, 88, 93, 120, 264]	5	9
[3, 6, 9, 27, 28, 45, 60, 84, 90, 168, 360]	28	11

Alternatively, your task is to find the longest non-repeating path through a graph whose nodes are the given integer `items`, and two nodes are connected by an edge if and only if one number is an integer multiple of the other. Under this definition, the edges are undirected, so each edge can be traversed in either direction. The general problem of finding the longest non-repeating path in an undirected graph is **NP-complete**, so no algorithms exist guaranteed to solve it efficiently in every case. Perhaps this restricted form of nodes may (or may not) allow helpful optimizations that are not available for general graphs. Who knows?

76. Forbidden digit

```
def forbidden_digit(n, d):
```

Suppose that all nonnegative integers that do not contain the digit d are listed in increasing order. For example, if d equals 3, this list would start as 0, 1, 2, 4, 5, 6, 7, 8, 9, 10, 11, 12, 14, 15, ...

This function should return the integer located at position n in this list, with counting starting from zero. Of course, the automated tester will again give your function large enough values of n so that constructing this list explicitly is a futile effort from the start. Instead, you need to use a bit of combinatorial thinking to find the result sooner. After the necessary flash of insight, the solution does not require any higher mathematics.

n	d	Expected result
3	3	4
5	0	6
21	3	24
7085	4	10752
29030	7	43835
10**100	5	676816616466104301364981403306033134233466210907249090703330489389102173000617274221681128634706638779211

77. Decoding the base $-1 + i$

```
def complex_base_decode(bits):
```

This exercise is suitable for mathematically minded students interested in Gaussian integers, that is, complex numbers with integer real and imaginary parts, normally represented as tuples of signed integers. However, as explained in the YouTube video “Counting in a Complex Base”, all Gaussian integers can also be encoded as bit strings whose digits are either 0 or 1; instead of the familiar binary base two, we use the Gaussian integer $-1 + i$ as the base. The k :th bit from the end then indicates whether the term $(-1 + i)^k$ is included in the sum, instead of the term 2^k in the ordinary binary numbers. These powers of $-1 + i$ mesh together so perfectly that every Gaussian integer (regardless of the signs of its real and imaginary parts) has a unique representation as a sum of distinct integer powers of $-1 + i$, analogous to how every integer, positive or negative, has a unique representation as a sum of distinct powers of minus two.

This function should return the Gaussian integer from the given `bits` in this representation. Yes, Python has a complex number type. Still, since its real and imaginary parts are floating-point numbers, their range isn't sufficient to represent the numbers decoded from bit strings thousands of bits long that the automated tester will be merrily throwing at this function. You therefore need to represent Gaussian integers as tuples of integer real and imaginary components, and calculate each power of $-1 + i$ yourself with the help of formula $(a + bi)(-1 + i) = -(a + b) + (a - b)i$, where $(a + bi)$ is the current power of $-1 + i$.

bits	Expected result
'0'	(0, 0)
'1'	(1, 0)
'1011'	(2, 3)
'10110'	(-5, -1)
'111011'	(2, -1)
'1101111110010101100110001010001101000001000000101001110101111100000101101000110001111010100110110000100110001011011100010100111111111011001100010111101100011110'	(-26781393650563464134561123, -25093735552551739092060943)

78. Independent dominating set

```
def independent_dominating_set(edges):
```

An **undirected graph** consists of n nodes numbered from 0 to $n - 1$. Each node is connected to other nodes by undirected edges, so if node u is a neighbour of node v , then node v is also a neighbour of node u . These edges are given as a list to your function so that each element `edges[u]` lists the neighbours of node u .

An independent set is a subset of nodes in a graph such that no two nodes in the subset are neighbours. On the other hand, a dominating set of a graph is a subset of nodes such that every node in the graph either is in this subset or is a neighbour of at least one node in this subset. On their own, both problems are NP-complete, so any algorithm to find the largest possible independent set or the smallest possible dominating set in a given graph will take exponential time to find and verify the solution for some particularly tricky graph.

However, combining these two problems restricts the possibilities enough that, at least for the graphs generated by our trusty automated tester, this function should be able to find the smallest **independent dominating set** quickly using recursive backtracking. For each node u , you have a simple two-way choice. You can choose to take the node u in the set, in which case you are not allowed to take any of the neighbours of u in the set. Or, you can choose to leave the node u out, in which case you must eventually take at least one of the neighbours of u in the set.

It is probably a good idea to first sort the nodes by their number of neighbours before the recursion searches through them. Since the smallest independent dominating set of the given graph is not necessarily unique, this function should only return the size of that set.

edges	Expected result
[[4, 3, 2], [2, 3], [1, 0], [1, 0], [0]]	2
[[5], [2, 5, 3], [1, 3, 5], [1, 2, 5], [], [1, 0, 3, 2]]	2
[[2, 5, 6], [6, 2, 4], [5, 4, 1, 0, 3, 6], [4, 2], [3, 2, 5, 1], [4, 0, 2], [2, 1, 0]]	1
[[2, 1, 7, 3, 9], [0, 8, 4], [4, 0, 9], [4, 0, 6], [3, 2, 1, 7], [7, 6], [5, 7, 3, 9], [4, 0, 9, 6, 5, 8], [1, 7], [2, 7, 0, 6]]	3

79. Vertex cover

`def vertex cover(edges):`

An **undirected graph** again consists of n nodes numbered from 0 to $n - 1$, given the same way as the list of `edges` as in the previous problem of finding the minimal independent dominating set. However, this time your function needs to find the minimum vertex cover; that is, a subset of nodes such that for every edge (u, v) of the graph, at least one of u and v belongs to the chosen subset.

The backtracking recursion for this problem is similar to the recursion for finding the minimum independent dominating set. It is straightforward to see that for every node u , the backtracking recursion gives a simple two-way choice: either you include u in the vertex cover, or you include **every** neighbour of u in the vertex cover. (In the dominating set problem, you were forced to take in at least one neighbour of u . It follows that the minimal vertex cover must necessarily contain at least as many nodes as the minimal dominating set, independent or not.)

It is again useful to sort the nodes in descending order of their number of neighbours before the backtracking recursion traverses them. Another good thing to notice is that if u is a node with only one neighbour v , we might as well nail in this neighbour v in the vertex cover before starting the recursive search, since v covers not only the edge between u and v , but also all the edges emanating from v to the rest of the graph. This heuristic does not apply to the construction of the optimal independent dominating set, although it does apply to unrestricted dominating sets (DUCY?). Also, be careful not to accidentally include both u and v in the vertex cover when these two nodes form a **duopole**, that is, when each is the other's only neighbour.

edges	Expected result
[[4, 3, 2], [2, 3], [1, 0], [1, 0], [0]]	2
[[5], [2, 5, 3], [1, 3, 5], [1, 2, 5], [], [1, 0, 3, 2]]	3
[[2, 5, 6], [6, 2, 4], [5, 4, 1, 0, 3, 6], [4, 2], [3, 2, 5, 1], [4, 0, 2], [2, 1, 0]]	4
[[2, 1, 7, 3, 9], [0, 8, 4], [4, 0, 9], [4, 0, 6], [3, 2, 1, 7], [7, 6], [5, 7, 3, 9], [4, 0, 9, 6, 5, 8], [1, 7], [2, 7, 0, 6]]	6

Since the nodes that were not chosen in the vertex cover form an independent set, finding a minimal vertex cover is equivalent to finding a maximal independent set. *To-may-to, to-mah-to*. All NP-complete search problems are fundamentally the same problem anyway, under different disguises.

80. Shotgun sequence

```
def shotgun(n):
```

Usually, programmers like to number sequence positions starting from zero, but starting from one works better for the wacky Shotgun sequence problem. Consider the infinite sequence of positive integers 1, 2, 3, 4, 5, 6, 7, 8, 9, ... Pair up numbers in positions **divisible by two** with each other and swap each number with its pair for the new sequence 1, 4, 3, 2, 5, 8, 7, 6, 9, ... Next, look at the numbers in positions that are **divisible by three**, and similarly pair them up and swap each number with its pair. In general, in the k :th round, look at the numbers currently in positions divisible by $k - 1$, and pair up and swap those numbers. After $n - 1$ rounds, the number that ended up in position n will never move again, and this function should return that number.

Denoting the number in position n after k rounds of this operation by $s(n, k)$, the base case of the recursive formula is simply $s(n, 0) = n$ for any n . If n is not divisible by k , then $s(n, k) = s(n, k - 1)$, again simple enough. If n is divisible by k , find the position m with which the position n is swapped in that round (this can be easily determined from the parity of n / k), and solve $s(m, k - 1)$ to find out what number was previously in position m . However, the automated tester will give you values of n large enough to make the resulting deep recursion hit a **stack overflow**. Hence, you need to convert your **tail recursion** into a for-loop that will be executed entirely in the same function activation.

n	Expected result
1	1
4	6
1912	2385
3348444	7610786

This “deranged” permutation of positive integers is a bizarre combination of the famous Hilbert Hotel and the Sieve of Eratosthenes in that these repeated swaps keep pushing all prime numbers infinitely far to the right along the number line, so that this function can never return a prime number for any finite n . Consider any prime number p . Since p is not divisible by any integer $1 < k < p$, it will move for the first time in round $p - 1$, when it is swapped with the integer currently in position $2p$. The integer in that position must be either the original $2p$ or some integer that was moved there in some previous round $k < p$, thus necessarily composite. The prime p will then remain in its new position $2p$ until round $2p - 1$, when it will be swapped with the integer currently in position $4p$, which, by the same logic, is also composite, and so on.

81. Card row game

```
def card_row_game(cards):
```

In this classic recursive programming chestnut, two players face each other in a game played with a row of `cards` turned face up. On alternating turns, each player can pick up either the card at the beginning of the row or the card at the end of the row. Each card contains an integer, and both players try to maximize the sum of the integers on the cards they picked up. The result of the game is the difference between the sums of the cards of the first and the second player; a negative result thus indicates that the second player won. Assuming both players make optimal decisions throughout the game, what is the end result for the given `cards` in this **principal variation** of perfect play?

The greedy strategy of always picking up the larger of the two cards does not generally produce the optimal play in this game. This problem should be modelled using a recursive **minimax** equation $C(b, e)$, where b and e are the positions of the `cards` where the current row begins and ends. The base case is $b = e$ when only one card remains, so the result of the game is the value of the card in that position. Otherwise, $C(b, e)$ equals the maximum of the two quantities $\text{cards}[b] - C(b + 1, e)$ and $\text{cards}[e] - C(b, e - 1)$, these quantities respectively giving the values of picking up the card from the beginning or the end of the current row. Note that in both cases, the recursive result is subtracted from the value of the card picked up, since that call evaluates the new situation from the opponent's point of view. In a zero-sum game, your opponent's gain exactly equals your loss, and vice versa, so one player's result is the negation of the other's.

cards	Expected result
[5, 5]	0
[3, 6, 2]	-1
[1, 6, 6, 10]	9
[2, 2, 23, 15, 3]	-5
[65, 17, 19, 68, 26, 90, 35, 28, 33, 46]	119

After solving this problem, you can analyze its **first-move advantage**, that is, how much of an advantage there is in getting to make the first move for a randomly constructed board.

82. The remains

```
def repeating_decimal(a, b):
```

Unless the prime factors of the denominator b of an integer fraction a / b are powers of 2 and 5, the decimal representation of that integer fraction ends up being infinitely long. However, for all rational numbers, the decimal representation is guaranteed to eventually become periodic, so after some initial prefix of decimal digits, the same pattern keeps repeating infinitely. Given an integer fraction a / b , this function should compute the initial part and the repeating part, returning the answer as a tuple of two strings containing the digits of these two parts. You may assume that $0 < a < b$, so that your function only needs to handle numbers between 0 and 1.

As explained in the section “Decimal expansion and recurrence sequence” on the Wikipedia page “Repeating decimal”, the initial and repeating parts of the decimal representation can be computed by performing the long division of a / b , just like you were taught back when you were a wee kid in math class, to generate the digits of the decimal representation. However, to avoid guessing when the decimal digits look like they might be repeating (both the initial part and the repeating part of the decimal representation can be arbitrarily long, even for seemingly simple fractions that don't seem to warrant this), your function needs to keep somehow track of the position where it has previously encountered each possible remainder. Once the long division encounters a remainder that it has already encountered during the same division, the rest of the division will be periodic, and you can stop and extract the repeating part as the decimal digits between the two occurrences of the same remainder. Whatever digits precede this repeating part will then be the initial part.

a	b	Expected result
1	7	('0.', '142857')
10	11	('0.', '90')
4	29	('0.', '1379310344827586206896551724')
3	44	('0.06', '81')
1	48	('0.0208', '3')
66	433	('0.', '152424942263279445727482678983833718244803695150115473441108 5450346420323325635103926096997690531177829099307159353348729 7921478060046189376443418013856812933025404157043879907621247 1131639722863741339491916859122401847575057736720554272517321 0161662817551963048498845265588914549653579676674364896073903 0023094688221709006928406466512702078521939953810623556581986 1431870669745958429561200923787528868360277136258660508083140 877598')

83. Count sublists with odd sums

```
def count_odd_sum_sublists(items):
```

Given a list of nonnegative integer `items`, this function should count the number of contiguous sublists whose sum of elements is an odd number. For example, the list `[0, 3, 2, 4]` contains six such sublists: `[0, 3]`, `[0, 3, 2]`, `[0, 3, 2, 4]`, `[3]`, `[3, 2]`, and `[3, 2, 4]`, read in order of starting position and breaking ties by length.

Of course, you could solve this problem with two nested for-loops. The outer loop iterates over all possible starting points for the sublist. The inner loop then iterates over the positions from that starting point to the end of the list, keeping track of the sum of the elements of that sublist so far. Whenever the current sum is odd, increment the tally by 1.

Short and simple, but we can actually do way better and solve this problem with just one for-loop! Have the inner loop of the previous algorithm loop through the positions of the entire list, keeping track of the sum of the elements up to that position. Maintain two integer variables, `odd_count` and `even_count`, to remember how many prefixes of the list have had odd and even sums of elements. Depending on the parity of the current element, increment the corresponding counter, and update the total tally of sublists with odd sums by the number of sublists of suitable parity that you have encountered so far.

As usual, the automated tester will feed your function lists long enough so that any “Shlemiel” solving this problem with two nested for-loops will run out of the time limit well before terminating.

items	Expected result
[2, 0, 4, 3]	4
[5, 0, 1, 2, 2]	8
[2, 1, 8, 0, 5, 6, 5, 0]	20
[9, 0, 5, 2, 8, 3, 3, 1, 9, 8, 3]	35
[22, 1, 12, 8, 5, 3, 5, 17, 8, 17, 19, 3, 1, 15, 14, 18, 8, 18, 0, 11, 18, 10]	132

84. Tailfins and hamburgers

```
def trip_plan(motels, daily_drive):
```

Your summer vacation plan is to drive along the nostalgic highway that crosses the entire country, staying in a motel each night to catch some Z's before the next day's drive across the fruited plain. You are given a list of `motels` along the highway in ascending order of distance, and you start your trip at milestone zero on the highway. (Note that this zero milestone is not the same thing as the first motel in position zero of the `motels` list.) Staying at any other motel is optional during the trip. Still, your trip must end at “The Gobbler”, always the last motel on the list, a historical landmark from a bygone era of happy motoring optimism.

You can drive any distance each day, and make your vacation take as many days as you like. However, you would prefer to have your daily drive to be as close to `daily_drive` as possible. We define the **misery** of each day to be the square of the difference between the actual distance and the `daily_drive`. For example, if you prefer to drive 300 miles each day, the misery of a day when you drove only 250 miles equals $(300 - 250)^2 = 2500$. This function should generate the trip plan that minimizes your total misery over the entire trip.

This trip plan should be returned as a list of the motels you choose to stay at each night, in ascending order. If there exist multiple travel plans with equal total misery, this function must return the lexicographically lowest such plan.

<code>motels</code>	<code>daily_drive</code>	Expected result
[7]	5	[7]
[3, 16, 32, 43]	22	[16, 43]
[16, 27, 49, 59, 72]	32	[27, 72]
[13, 23, 28, 55, 71, 82]	38	[28, 55, 82]
[13, 18, 32, 44, 81, 88, 114, 144]	49	[44, 88, 144]
[21, 41, 74, 98, 130, 171, 183, 198, 216, 267]	93	[98, 183, 267]

85. Split the digits, maximize the product

```
def max_product(digits, n, m):
```

A string contains exactly $n + m$ decimal `digits` listed in descending order. Your task is to use these `digits` to build two positive integers of exactly n and m digits, respectively, to maximize the product of your two integers. Obviously, the digits chosen for both integers will appear in descending order, but in which integer will you place each digit to maximize the resulting product?

This problem is based on the adorable classroom story given in the post “Mrs. Nguyen’s prestidigitation” by Brian Hayes. Once you work it out, the algorithm is short and sweet... provided there are no duplicate digits or zeros! However, your function must work correctly even without these restrictions.

This problem can be split into recursive two-way decisions by realizing that the first digit must go to either the first or the second integer being constructed. After this two-way choice, continue by placing the second digit in either number, and so on. The base case of this recursion is when one of the numbers reaches its required length, at which point you can stuff the remaining digits into the other number and compare the resulting product to the current maximum. Make sure to also handle each block of k consecutive identical digits as a branch of $k + 1$ possibilities of how many of these digits you place into the first number, as opposed to a branch of 2^k redundant possibilities to divvy up these identical digits between the two numbers.

digits	n	m	Expected result
'6540'	3	1	3240
'87452'	3	2	63168
'942200'	5	1	379800
'9877100'	4	3	8448700
'8866666555554222211110'	12	10	7510047809441458785210
'9998888876544332000000'	11	11	9877523839516576000000

86. Tower of cubes

```
def cube_tower(cubes):
```

The six faces of a cube are numbered from 0 to 5 so that the face opposite to the face number i is given by the expression $(i+3)\%6$. Your function is given a list of `cubes` where the sides of each cube have been coloured with colours encoded as natural numbers. These `cubes` are all different sizes and are listed in increasing order of size. Your task is to build a tower with the largest possible number of cubes under the rules that a larger cube may not be placed on top of a smaller one, and that the touching faces of adjacent cubes must always be of the same colour. You can freely rotate each cube to any orientation during the build. Return the largest number of cubes that the tower can contain under these constraints.

This problem is probably easiest to solve with the classic “take it or leave it” recursive search. Either you take the largest remaining cube as the tower's base, or you don't. Either way, continue building the best tower with the remaining smaller cubes, making sure the colour of the face of your next chosen cube always matches the face of the previously chosen cube it rests on. Since the same subproblems can and will repeat exponentially during this recursive search, you should again use the `@lru_cache` decorator to prune the repeated branches of this search tree.

cubes	Expected result
[[0, 2, 2, 1, 2, 4], [3, 1, 2, 0, 3, 2], [4, 1, 4, 4, 1, 2]]	3
[[2, 9, 1, 4, 8, 6], [1, 7, 6, 2, 3, 2], [0, 5, 5, 2, 6, 2], [2, 9, 8, 3, 6, 8]]	3
[[7, 21, 14, 13, 19, 9], [14, 16, 22, 22, 20, 7], [2, 3, 0, 4, 2, 18], [2, 19, 5, 12, 21, 8], [0, 12, 13, 11, 4, 15], [15, 2, 5, 3, 16, 7], [15, 16, 1, 2, 3, 12], [10, 13, 22, 3, 10, 15], [1, 20, 11, 6, 2, 12], [15, 9, 1, 14, 10, 1], [4, 14, 9, 17, 7, 15], [7, 22, 6, 4, 16, 19], [7, 21, 19, 18, 10, 18], [4, 18, 4, 22, 2, 5], [13, 10, 0, 20, 17, 21], [1, 10, 15, 10, 1, 3], [15, 2, 4, 6, 0, 17], [21, 21, 11, 16, 0, 18], [5, 1, 2, 14, 17, 8], [16, 19, 3, 21, 15, 12]]	10

This problem is adapted from problem 10051 at Online Judge, “Tower of Cubes”, included in the Programming Challenges compilation by Steven Skiena of “The Algorithm Design Manual” fame.

87. Gijswijt sequence

```
def gijswijt(n):
```

The Gijswijt sequence is an interesting sequence of positive integers that starts with the value 1. After that, each element of this sequence equals the largest integer k for which there exists some suffix of the preceding sequence, no matter what length, that repeats k times consecutively. For example, the element following the initial sequence 1, 1, 2, 1, 1, 2, 2, 2 is 3, since the preceding sequence contains the suffix [2] repeated three times consecutively.

This function should return the element at position n of the Gijswijt sequence, with position counting again starting from 0 for us budding computer scientists. The automated tester is guaranteed to give your function values of n in increasing order, so your function should construct this sequence into a global list initialized to contain some prefix of the Gijswijt sequence, and each time extend this sequence one element at a time until the stored sequence list contains the requested position n .

As explained in the article “Seven Staggering Sequences” by Neil Sloane, this sequence will eventually contain any positive integer. However, this sequence grows so slowly that element 4 first appears at position 219 (zero-based), and element 5 first appears at a position of about 10 raised to the power of 10^{23} . (Read that sentence again to make sure that you understood it correctly.) Your function may therefore safely assume that each element in this sequence for the values of n requested by the tester will be either 1, 2, 3 or 4. To solve the element in position n , your function should try out all possible repeated suffix lengths k from 1, 2, 3 ... and quickly determine how many times the suffix of length k is repeated. Don't be a Shlemiel who extracts entire sublists to compare for equality. Instead, use a loop to compare pairs until you find a mismatch, which usually happens after only a couple of elements.

n	Expected result
0	1
6	2
91	3
219	4
564	2

88. Sum of consecutive squares

```
def sum_of_consecutive_squares(n):
```

The original collection of *109 Python problems* included problems that determined whether a given positive integer n can be expressed as a sum of exactly two squares of integers, or as a sum of cubes of one or more distinct integers. Continuing in this same spirit, the problem “Sum of Consecutive Squares” that appeared recently in the popular Stack Overflow Code Golf coding problem collection site asks for a function that checks whether the given positive integer n could be expressed as a sum of squares of one or more **consecutive** positive integers. For example, since $77 = 4^2 + 5^2 + 6^2$, the integer 77 can be expressed as a sum of some consecutive integer squares.

This problem is best solved with the classic **two-pointers** approach, using two indices `lo` and `hi` as the **point man** and **rear guard**, respectively, to delimit the range of integers whose sum of squares we want to make equal to n . Initialize both indices `lo` and `hi` to the largest integer whose square is less than or equal to n , and then initialize a third local variable `s` to keep track of the sum of squares of integers from `lo` to `hi`, inclusive. If `s` is equal to n , return `True`. Otherwise, depending on whether `s` is smaller or larger than n , expand or contract this range by decrementing either index `lo` or `hi` (or both) as appropriate. Then update `s` to give the sum of squares of the new range, and continue the same way.

<code>n</code>	Expected result
9	True
30	True
294	True
3043	False
4770038	True
24015042	False
736683194	False

When implemented as explained above, this function maintains a loop invariant stating that if n can be expressed as a sum of squares of consecutive integers, the largest integer in this sum cannot exceed `hi`. This invariant is initially true because of the initial choice of `hi`. Maintenance of this invariant over a single round of the loop body can then be proven for both branches of the loop body: $s > n$ and $s < n$. Since at least one of the positive indices `hi` and `lo` will decrease each round, this loop must necessarily terminate after at most $2 * hi$ rounds, a massive improvement over the “Shlemiel” approach of iterating through all possibilities in **quadratic** time with two nested loops... or if the coder is especially clumsy, **cubic** time for three levels of nested loops.

89. Balance of power

```
def shapley_shubik(weights, quota):
```

In a democracy, we are accustomed to one vote per person, but in corporate meetings, individual members may have different numbers of votes. In such situations, each voter's power is not necessarily proportional to their number of votes. For example, under majority voting with four voters and vote weights `[4, 1, 1, 1]`, the first voter effectively acts as a dictator, deciding how every vote will go. On the other hand, with three voters and weights `[10000, 10000, 1]`, the Lilliputian third voter wields massively disproportionate power to unilaterally decide the outcome of every issue where the big boys disagree!

The Shapley–Shubik power index can gauge each voter's relative power in such situations. Given the `weights` of each voter followed by the `quota`, the total number of votes required to win the vote (this is normally half the sum of `weights` plus one, but can also be something else, such as the entire sum in a jury trial, or just one in a **vetocracy**), the algorithm should iterate through all $n!$ possible permutations of the n voters. For each permutation of voters, the algorithm adds up the votes from the beginning to the point of the **pivot** voter of that permutation, that is, the voter who causes the sum of votes so far to become greater than or equal to the `quota`. The Shapley–Shubik power index of each voter is defined simply as the number of times that voter ends up being the pivot voter over the $n!$ possible permutations.

Your function should recursively generate the voter permutations one at a time. However, to save time, you can stop the generation of the current permutation when you reach the pivot voter of that permutation. Before returning, add $m!$ to the power index of that pivot voter, where m is the number of remaining voters, including that pivot voter. (It might come in handy to precompute the factorials into a list for quick lookup during the recursion.)

weights	quota	Expected result
[1, 1, 1]	3	[2, 2, 2]
[4, 1, 1, 1]	4	[24, 0, 0, 0]
[4, 1, 1, 1]	1	[6, 6, 6, 6]
[9, 8, 6, 1, 1]	18	[54, 54, 4, 4, 4]
[23, 22, 17, 16, 5, 1, 1]	48	[1428, 1428, 1008, 1008, 168, 0, 0]
[38, 38, 35, 34, 34, 21, 10, 1]	153	[7872, 7872, 6528, 6528, 6528, 3840, 1152, 0]

90. Longest zigzag subsequence

```
def longest_zigzag(items):
```

As we may recall, a **subsequence** of the given list of `items` consists of some of these items, not necessarily consecutive, listed in the order in which they appear in the original `items`. For example, the list `[2, -1, 3, 2]` is one possible subsequence of the list `[8, 2, -1, 2, 3, 2]`. We next define a subsequence to be **zigzag** if the signs of the differences between consecutive elements strictly alternate between `+1` and `-1`, where the first such difference can be either positive or negative. This function should find the longest zigzag subsequence of the given list of `items`. Since this sequence is not necessarily unique, you should return its length.

This classic exercise on **dynamic programming** can be solved by tabulating the values of the function $s(i, d)$, where i is a position in the `items` list and d is the direction of the next allowed step, either `+1` or `-1`. When $d = +1$, $s(i, +1)$ equals one plus the maximum of the values $s(j, -1)$ for the positions $j > i$ for which `items[j] > items[i]`. The formula for $s(i, -1)$ is defined symmetrically. For the position m at the end of the `items` list, $s(i, +1) = s(i, -1) = 1$. You can compute these values either by tabulating them from the end of the list to the beginning, or by using a recursive formula, again aided by the `@lru_cache` function decorator to eliminate the exponentially repeated subproblem computations. The final answer to the problem is then simply the larger of the two values $s(0, +1)$ and $s(0, -1)$.

items	Expected result
[7, 3]	2
[3, 4, 1]	3
[10, 6, 4, 1, 5]	3
[5, 1, 3, 0, 1]	5
[1, 2, 3, 4, 5, 6, 7, 8, 9, 10]	2
[5, 6, 3, 2, 2, 9, 9, 0, 3]	6

91. Kimberling's expulsion sequence

```
def kimberling_expulsion(start, end):
```

Kimberling's expulsion sequence is one of the numerous wacky sequences of positive integers devised by Clark Kimberling. Assume that positions of positive integers 1, 2, 3, 4, 5, 6, 7, 8 ... are indexed starting from zero. To produce the i :th element of the result sequence, remove the element from position i , and rearrange the first $2i$ elements of the remaining sequence by **riffling** (same as in “Riffle shuffle kerfuffle”, problem 3 in the original *109 Python Problems* collection) the i elements before and the i elements after position i together with an **in-shuffle**; that is, the first element of the riffle is taken from the block that follows position i .

For example, the element in position 0 of the result is 1, leaving the remaining elements in the order 2, 3, 4, 5, 6, 7, 8, ... The next element of the result is 3, the element in position 1 of the remaining elements, and riffling its surrounding elements produces the new sequence 4, 2, 5, 6, 7, 8, ... of remaining elements. The next element of the result is 5, the element in position 2 of the remaining elements, and riffling its surrounding elements produces the new sequence 6, 4, 7, 2, 8, ... of remaining elements. And so on.

This function should return the subsequence of Kimberling's exclusion sequence from the position `start` (inclusive) up to position `end` (exclusive). The automated tester is guaranteed to give you test cases so that `start` is 0 in the first test case, and from every test case onwards, `start` will always equal the end of the previous test case. Your function should therefore store the remaining sequence information constructed so far in some global variables between the calls, so that it doesn't have to start producing each new subsequence all the way from the beginning.

start	end	Expected result
0	2	[1, 3]
2	5	[5, 4, 10]
12	14	[31, 14]
710	725	[1525, 2086, 1849, 1471, 1958, 836, 1035, 351, 648, 1263, 299, 579, 1531, 1019, 1682]

It is still unknown whether this sequence contains every positive integer.

92. Kimberling's repetition-resistant sequence

```
def repetition_resistant(n):
```

Right on the heels of the previous problem, here is another self-referential sequence devised by Clark Kimberling whose chaotic behaviour emerges organically from the iteration of a simple deterministic rule. This binary sequence consists of zeros and ones, and starts with a zero.

To determine the next bit of the sequence σ constructed so far, consider its **longest repeated subsequence** that occurs at least twice. For example, the longest repeated subsequence of 01100111 is 011, of length 3. Partially overlapping repetitions count, so the longest repeated subsequence of 010101 would be 0101. Let m be the length of the longest repeated subsequence of σ . To decide whether σ will continue with 0 or 1, look at the sequences $\sigma 0$ and $\sigma 1$, obtained by extending σ with 0 and 1, respectively. The sequence is extended by 1 if and only if the sequence $\sigma 0$ contains a repeated subsequence of length $m + 1$, but the sequence $\sigma 1$ does not contain such a subsequence. The reader can verify that this infinite sequence begins with

```
01000110101110010011110110000010100001110100101101111100010000000110...
```

This function should return the character at position n of this sequence, with positions numbered from zero. To speed up your function, note that adding a new character to the end of the sequence constructed so far can create a new longest repeated subsequence of length $m + 1$ only if that subsequence lies at the very end of the new sequence. The automated tester provides your function values of n from 0 to 9999, so your function should save the sequence generated so far and the length of its longest repeated subsequence in global variables stored between function calls.

n	Expected result
0	0
1	1
1000	0
10000	0

According to Kimberling, this infinite sequence contains every possible finite bit string. (Therefore, it actually contains every bit string not just once, but infinitely many times; DUCY?) The reader can also compare this sequence to the similarly spirited Linus sequence, problem 52 in this author's collection of Java problems for CCPS 209.

93. Game with multiset

```
def game_with_multiset(queries):
```

This nifty little problem is Problem 1913C from Codeforces, “Game with Multiset”. Your function is given a list of `queries` of the form `('A', x)` or `('G', x)`, where `x` is some non-negative integer. Your function needs to process these queries in the exact order given. To do this, your function needs to somehow maintain a **monotonic multiset** of non-negative integers, initially empty. Unlike in an ordinary set, the same value can occur in the multiset any number of times.

Each query `('A', x)` adds the integer 2^{**x} to your multiset. (Since the numbers that we are going to store are powers of two, of course you will only store the raw exponent `x`, instead of 2^{**x} .) Each query `('G', v)` asks you to determine whether the integer `v` can be constructed as a sum of some values that have been added to the multiset so far. Each value can be used in the sum up to as many times as it currently occurs in the multiset at that point. Your function should return the list of truth values for the results of the `('G', v)` queries.

This problem would be easy if each `x` could occur in the set only once, since you could just express `v` in binary and see if the corresponding powers of two occur in the set. (This would be literally just one line of code using clever **bitwise arithmetic**.) Multisets give us more freedom and possibilities in constructing the sum, so that if you don't have some particular power that you need stored in the set, using two instances of the one notch lower power (or four instances of the two notches lower power, or any suitable combination of lower powers for that matter) does the job as well. The question is how you arrange all this efficiently. One linear loop from the highest bit of `v` down to the lowest should be enough, keeping track of how many lower powers it needs to make up for the bit deficit it has accumulated down the line. (You can also note the handy Python built-in `bin`.)

queries	Expected result
<code>[('A', 0), ('A', 0), ('A', 2), ('G', 6), ('G', 3), ('G', 5)]</code>	<code>[True, False, True]</code>
<code>[('A', 3), ('A', 1), ('A', 4), ('G', 26), ('G', 23), ('G', 18), ('G', 1), ('A', 2), ('A', 1), ('G', 22)]</code>	<code>[True, False, True, False, True]</code>
<code>[('A', 4), ('A', 2), ('A', 2), ('G', 24), ('A', 1), ('A', 1), ('A', 5), ('G', 1), ('A', 4), ('A', 3), ('G', 10)]</code>	<code>[True, False, True]</code>

94. Tower of Babel

```
def tower_of_babel(blocks):
```

The side lengths of a solid three-dimensional block are given as a three-tuple (a, b, c) . A block can be freely rotated and placed on a flat surface in three ways, depending on which face you make its bottom face. Your task is to reach for the heavens by building the tallest possible tower under the constraint that a block can be placed on another block only if its bottom face is **strictly smaller** in both dimensions than the top face of the block directly below it. You are allowed to place only one block directly on the block (or on the table underneath), and are not allowed to combine multiple thin pillar blocks on top of the same block to form a wide block.

This function should return the height of the tallest possible tower that can be built from the given list of `blocks`. This can once again be done with recursion: first choose the bottom block, then build the tallest possible tower on top of it using the remaining blocks. It will probably speed up the search to preprocess these blocks in some useful manner, at least initially, by rotating each block so that its three dimensions are sorted in descending order.

blocks	Expected result
[(5, 6, 4), (5, 7, 3)]	12
[(8, 10, 1), (1, 8, 6), (10, 7, 9), (4, 6, 2)]	20
[(1, 4, 6), (3, 1, 11), (11, 2, 1), (2, 10, 4), (7, 8, 3)]	29
[(5, 4, 9), (2, 8, 8), (8, 1, 7), (4, 4, 6), (12, 13, 3), (7, 10, 9)]	33
[(1, 2, 9), (4, 10, 1), (16, 8, 11), (3, 1, 1), (14, 4, 10), (15, 3, 3), (13, 5, 15), (11, 16, 5)]	65

(This problem is problem 437, “Tower of Babylon”, in the Online Judge collection.)

95. Make a list self-describing

```
def self_describe(items):
```

A list of positive integer `items` is **self-describing** if every x that appears in that list appears exactly x times. For example, `[1]`, `[2, 4, 4, 2, 4, 4]` and `[3, 2, 1, 2, 3, 3]` would be self-describing in this sense. Your function receives a list of n integer `items`, each from the range `1, ..., n`. Assuming that you are allowed to change any of these `items` to any values of your choice, your function should compute and return the smallest number of such changes required to turn the `items` list self-describing. Since you can trivially set each element to n , this problem has a well-defined solution for every possible list of `items`.

This problem can be solved in various ways. This author suggests you first define a **recursive generator** function that yields all possible integer partitions of the given integer n that don't contain any repeated elements. For example, when $n = 10$, this generator yields the lists `[10]`, `[9, 1]`, `[8, 2]`, `[7, 3]`, `[7, 2, 1]`, `[6, 4]`, `[6, 3, 1]`, `[5, 4, 1]`, `[5, 3, 2]` and `[4, 3, 2, 1]`, one at a time.

Armed with this generator, your function should then preprocess the `items` by computing the counter dictionary `c` so that `c[x]` equals the number of times that x appears in `items`. Then, loop through all the possible non-repeating partitions of n and compute the cost of converting the `items` into a list that contains precisely the numbers that appear in that partition in their correct multiplicities, and finally return the smallest cost found this way.

items	Expected result
<code>[1]</code>	<code>0</code>
<code>[1, 2, 3]</code>	<code>1</code>
<code>[2, 2, 5, 5, 5, 5]</code>	<code>2</code>
<code>[1, 2, 4, 3, 2]</code>	<code>2</code>
<code>[3, 2, 3, 4, 3, 1, 2, 2, 2]</code>	<code>3</code>
<code>[3, 6, 2, 7, 3, 5, 7, 2, 5, 2]</code>	<code>4</code>

96. Two pins, not three, Dolores

```
def kayles(pins):
```

Kayles is a combinatorial game made famous by Conway, Guy, and Berlekamp, and is often used to teach basic concepts of combinatorial game theory. The game is a hypothetical bowling variant in which some pins are arranged in a linear row. Both players are perfectly accurate, so either player can knock down any one pin or any two adjacent pins of his choice, but can never knock down three or more pins, no matter how he throws. Two players take turns knocking down pins; the player who knocks down the last pin and leaves his opponent unable to move wins the game.

The game has a trivial winning strategy for the first player if the pins are in a single group, so to make it more interesting, assume some pins have been randomly removed from the row before the game begins. This function receives the current state as a list of pin groups. For example, the list `[7, 3, 4, 2]` means that the first group contains seven consecutive pins, the second group contains three consecutive pins, and so on. The function should determine whether the first player has a winning strategy starting from the given list of pins. This can be done by searching all possible moves for the current player and recursively evaluating the resulting positions from the opponent's perspective. The current situation is hot (winning) if there exists a move that leads to a cold (losing) situation for the opponent, so that all his possible moves from that situation lead to hot states.

The base case is when only one non-empty group of pins remains, as above. A similar argument also shows that any two groups with the same number of pins cancel each other out, since neither player can gain more than a temporary respite by making any play in either group. (DUCY?) For example, `[7, 4, 4, 4, 2, 2, 1]` plays essentially the same as `[7, 4, 1]`.

pins	Expected result
[3, 1]	True
[4, 1]	False
[2, 7, 1, 3]	True
[2, 3, 1]	False
[2, 2, 3, 1, 1, 1, 2]	False
[2, 6, 1, 5, 3, 2, 5, 5, 1, 4, 3]	True

97. Out where the buses don't run

```
def bus_travel(schedule, goal):
```

Cities in Poldavia have been numbered starting from zero. As a happy-go-lucky backpacker, you want to get from city 0 to the `goal` city by travelling on the bus network of that pastoral country. You have a `schedule` that lists the buses leaving from each city that day. Each leg of the bus trip is given as a triple (`destination`, `leave`, `arrive`), giving the times the bus leaves the current city and arrives at its `destination`. For simplicity, assume transfers between buses in the same city take zero time. Poldavia uses a 24-hour clock (`hour`, `minute`) to measure the day.

This function should determine the earliest time you can arrive at the `goal` based on the bus schedule. To achieve this with a simplified version of **Dijkstra's algorithm**, your function should keep track of the earliest possible arrival time to each city. For the starting city 0, this time is the starting midnight. For the rest of the cities, you will update their earliest arrival times whenever you find faster routes to get to these cities. If it's not possible to reach the `goal` city at all within the same day before midnight, return the midnight hour (24, 0) to indicate this.

The function should maintain a **frontier** of cities that need processing. Initially, this frontier contains only the starting city 0. Then, repeatedly pop one city from the frontier; it; it doesn't matter which one. Loop through the buses taking off from your chosen city after the earliest possible arrival time at that city. If the arrival time of some route to its `destination` is earlier than your current earliest known arrival time to that destination, update the earliest known arrival time of that destination, and insert the destination city into the frontier. Repeat this until the frontier becomes empty. (Also, since buses can't travel backward in time, you can ignore all buses whose arrival time at their destination comes after the currently known earliest arrival time at the `goal`.)

schedule	goal	Expected result
{0: [(1, (15, 54), (17, 27)), (2, (17, 46), (19, 7)), (4, (18, 1), (19, 17))], 1: [(0, (17, 27), (17, 46)), (2, (20, 29), (21, 6)), (3, (15, 14), (16, 47)), (4, (15, 27), (16, 3)), (5, (18, 3), (19, 39))], 2: [(1, (19, 7), (20, 29)), (3, (21, 6), (21, 44))], 3: [(0, (16, 47), (18, 1)), (1, (13, 53), (15, 14))], 4: [(5, (13, 37), (14, 17)), (5, (16, 3), (17, 26)), (5, (19, 17), (20, 36))], 5: [(1, (14, 17), (15, 27)), (1, (17, 26), (18, 3)), (3, (12, 46), (13, 53))]}	2	(19, 7)

98. SMETANA interpreter

```
def smetana_interpreter(program):
```

SMETANA is a wacky esoteric programming language based on **self-modifying code**. This language is not **Turing-complete** in its basic form, but it is still powerful enough to give us an interesting coding problem. Your function should simulate the execution of the `program` given as a list of statements, each statement a string of the form `'GOTO x'` or `'SWAP x y'`, where `x` and `y` are statement numbers in the list, with position indexing again starting from zero. The function should return the number of execution steps needed to terminate the program. If the program execution enters an infinite loop, your function should return `None`.

Execution starts at statement 0 and ends when it goes past the last statement in the list. The effect of `'GOTO x'` is to make the execution jump to the statement currently in position `x`. The effect of `'SWAP x y'` is to swap the statements currently in positions `x` and `y`, and continue the program execution from the next statement after the current position.

To detect the program execution being stuck in an infinite loop, you should simulate the execution of the given program in two separate instances of execution that operate in lockstep. In the first instance, **the hare** executes the program two steps forward for every step of the second instance, the **tortoise**. If the execution states of the tortoise and hare ever become equal, you can safely conclude that the program is stuck in an infinite loop.

program	Expected result
<code>['GOTO 2', 'SWAP 1 0', 'SWAP 2 1']</code>	2
<code>['SWAP 1 2', 'GOTO 0', 'GOTO 3', 'GOTO 1']</code>	None
<code>['SWAP 2 0', 'SWAP 0 3', 'SWAP 1 2', 'GOTO 3']</code>	4
<code>['SWAP 2 3', 'SWAP 1 3', 'GOTO 4', 'SWAP 4 0', 'GOTO 0', 'SWAP 4 1', 'SWAP 0 2']</code>	None
<code>['GOTO 5', 'GOTO 2', 'GOTO 8', 'SWAP 3 6', 'SWAP 6 5', 'GOTO 3', 'SWAP 0 7', 'SWAP 1 7', 'SWAP 1 4']</code>	18
<code>['SWAP 10 8', 'SWAP 1 4', 'GOTO 1', 'GOTO 7', 'GOTO 10', 'GOTO 7', 'GOTO 8', 'SWAP 6 2', 'SWAP 9 4', 'SWAP 8 1', 'SWAP 5 6']</code>	5

99. The sharpest axes

```
def vector_add_reach(start, goal, vectors, giveup):
```

This problem was inspired by yet another excellent Quanta Magazine article, “An Easy-Sounding Problem Yields Numbers Too Big for Our Universe,” that the reader should first read for a good explanation of what this deceptively simple problem is all about. Given a list of direction `vectors` that determine the possible moves from the current point, this function should compute the number of steps needed to reach the `goal` coordinates. If it is not possible to reach the `goal` state within `giveup` moves, this function should return `None`.

This problem can be solved with a simple **breadth-first search**, as with many previous problems. Initially, the first layer consists only of the `start` point. Repeatedly compute the next layer of reachable points by looping through the points in the current layer and generating the list of all points reachable from these points with the given `vectors`, terminating your search either when you find the `goal` point or when more than `giveup` layers have been generated. However, the **reversible** nature of the moves allows bidirectional search to cut the effective search depth through this exponentially large search tree in half, yielding massive savings in both time and space!

Operate two separate instances of breadth-first search, one emanating from the `start` point using the `vectors` to the forward direction, the other emanating from the `goal` point using the negated `vectors` to the backward direction. Alternate between generating the successive layers of reachable points for these two searches, terminating when either search “meets the other one in the middle” by finding some point that was already discovered by the opposite search.

The following table contains examples from two-dimensional space only to keep it compact, but your function should work correctly for points in any number of dimensions.

start	goal	vectors	giveup	Expected result
(3, 0)	(2, 2)	[(-2, 1), (0, 2)]	6	None
(3, 2)	(4, 1)	[(2, -3), (1, 2), (-1, 2)]	6	2
(5, 7)	(0, 4)	[(-3, -5), (2, 1), (-4, 3), (2, 5), (3, 5)]	15	11
(9, 1)	(8, 3)	[(-1, 2)]	10	1

100. Vidrach Itky Leda

```
def itky_leda(board):
```

This cute little problem on state space searching is taken from Jeff Erickson's collection of old homework and exam questions. The `board` is an n -by- n square grid with a positive integer in every cell. A red token initially sits in the corner cell $(0, 0)$, and a blue token initially sits in the opposite corner cell $(n - 1, n - 1)$. The problem is to exchange the locations of these two tokens in the fewest possible moves.

Each move consists of the player choosing one of the two tokens and one of the four compass directions, then moving the token that many steps, where the number of steps is given by the value in the cell currently occupied by the other token. A token may not move outside the `board` or end up in the cell occupied by the other token. This function should return the smallest possible number of moves necessary to exchange the positions of the red and blue tokens, or return `-1` if this task is impossible on the given `board`.

This problem can be solved with a breadth-first search, but, as in the previous problem, the reversible nature of the moves lets you tighten the search exponentially with bidirectional search. Same as before, you should operate two separate instances of breadth-first search, one starting from the initial state and the other starting from the goal state. Alternate between these two searches to generate the successive layers of reachable states, terminating the search when either search “meets the other one in the middle” by finding some state already discovered by the opposite search. Chopping the depth of the search tree and therefore the exponent of the algorithm's worst-case running time in half is surely nothing to sneeze at!

board	Expected result
[[1, 0], [0, 1]]	-1
[[1, 2, 4, 3], [3, 4, 1, 2], [3, 1, 2, 3], [2, 3, 1, 2]]	5

101. How's my coding? Call 1-800-3284778

```
def keypad_words(number, words):
```

The E.161 recommendation maps letters from A to Z to touch-tone keypad digits from 2 to 9 so that phone numbers that are otherwise difficult for humans to memorize can be given catchy phonetic mnemonics such as 1-800-DOCTORB. Since each digit can denote three or four different letters, the same series of digits could represent many different words and combinations of words. We want to simplify finding a catchy mnemonic with a function that generates a list of all word combinations that map to the number randomly given to us by Ma Bell, so we can scan the list and choose the snappiest mnemonic for that number.

To keep the size of the returned lists manageable, the list of words will contain only words with lengths between 3 and 7, inclusive. There are still 97,106 such words in `words-sorted.txt`, so you need to speed up somehow the search for words that match the given sequence of digits. For improved readability, you should use minus signs to separate the individual words in the mnemonic. To make the results unique to allow automated testing, this function should return the list of mnemonics in sorted order.

number	Expected result (using the wordlist <code>words-sorted.txt</code>)
'4444444'	['ghi-gigi', 'ghi-high', 'gig-gigi', 'gig-high', 'gigi-ghi', 'gigi-gig', 'gigi-ihl', 'gigi-iii', 'high-ghi', 'high-gig', 'high-ihl', 'high-iii', 'ihl-gigi', 'ihl-high', 'iii-gigi', 'iii-high']
'2244638'	['aah-gneu', 'ach-gneu', 'bag-gneu', 'bagh-met', 'bagh-meu', 'bagh-mev', 'bagh-net', 'bagh-ofl', 'bah-gneu', 'bai-gneu', 'bagnet', 'bch-gneu', 'cag-gneu', 'cai-gneu']
'8572539'	['tlr-akey', 'tlr-alew', 'tlr-alex', 'tlr-blew', 'tlr-clew']
'8634633'	['tod-gode', 'tod-goff', 'tod-hoed', 'tod-inde', 'tod-iode', 'toe-gode', 'toe-goff', 'toe-hoed', 'toe-inde', 'toe-iode', 'ume-gode', 'ume-goff', 'ume-hoed', 'ume-inde', 'ume-iode', 'undined', 'unfined', 'vod-gode', 'vod-goff', 'vod-hoed', 'vod-inde', 'vod-iode', 'voe-gode', 'voe-goff', 'voe-hoed', 'voe-inde', 'voe-iode']

102. Scatter her enemies

```
def queen_captures_all(queen, pawns):
```

A tweet by chess grandmaster Maurice Ashley inspired this cute little problem. On a generalized n -by- n chessboard, the positions of a lone queen and the opposing pawns are given as tuples (row, col) . This function should determine whether the queen can capture all enemy pawns in one unbroken sequence of moves, where each move captures exactly one enemy pawn and immediately removes it from the board. The pawns stay put while the queen executes her sequence of moves. Note that the queen cannot teleport through pawns, but must always capture the nearest pawn to her chosen direction of that move.

This problem is best solved with recursion. The base case is when zero pawns remain, so the answer is trivially `True`. Otherwise, when m pawns remain, find the pawn nearest to the queen in each of the eight compass directions. For each direction where there exists a pawn to be captured, recursively solve the smaller version of the problem with the queen, having captured that particular pawn, and attempt to capture the $m - 1$ remaining pawns in the same regal fashion. If any one of the eight possible starting directions yields a working solution for the smaller problem with $m - 1$ pawns, that gives a working solution for the original problem for m pawns.

queen	pawns	Expected result
(4, 4)	[(0, 2), (4, 1), (1, 4)]	False
(1, 3)	[(0, 3), (2, 0), (2, 2)]	False
(2, 1)	[(1, 1), (5, 4), (2, 0), (5, 3)]	True
(0, 0)	[(3, 1), (4, 3), (2, 0), (6, 4), (0, 5)]	False
(11, 7)	[(0, 4), (10, 8), (5, 8), (6, 9), (9, 6), (11, 9), (2, 13), (11, 6), (5, 0), (9, 7), (11, 4)]	True

The bottleneck of this recursion is quickly finding the nearest pawn in each direction, and realizing as soon as possible that the current sequence of initial captures cannot be extended to capture the remaining pawns. If you preprocess the pawns to encode the adjacency information as a graph whose nodes are the positions of each pawn and the initial position of the queen, you must also note that the neighbourhood relation between these positions changes dynamically as the queen captures pawns along her route, causing pawns initially separated by those captured pawns to become each other's neighbours. As the recursive calls return without finding a solution, you need to downdate your data structures to undo the updates that were made to these data structures to reflect the new situation before commencing that recursive call.

103. Sneaking

```
def sneaking(n, start, goal, knights):
```

Separated from his beloved queen, who is wreaking havoc on a different board in the previous problem, a lone chess king finds himself at the `start` position on the n -by- n chessboard populated by sleeping enemy `knights`. The king would like to escape by reaching the goal position with the exit. Unless woken, the enemy knights will remain asleep and not move during this demented version of the game of the kings. However, these sleeping knights are essentially spiders that have set up sensitive tripwires on every square they guard within a single knight's jump. Should the king step on any guarded square, all knights that guard that square will instantly wake up and behead and devour the king.

This function should determine whether the king can reach the `goal` without losing his head. So this problem isn't merely a simple breadth-first search through unguarded squares; the king can also sneak into any square occupied by a sleeping knight not guarded by another knight and quietly slice that knight's throat to continue his journey from that square. Solving some of these instances requires the king to do precisely that to unguard a square needed for the rest of the journey.

To speed this up, it might help to precompute an n -by- n table that tracks how many knights are currently guarding each square. During the recursion, update this table whenever some knight gets his throat cut, and then update the table back to the previous values whenever the recursive search backtracks to undo the slaying and restore that knight to life.

n	start	goal	knights	Expected result
4	(3, 2)	(1, 1)	[(1, 2), (2, 1), (0, 3), (2, 3), (0, 2)]	False
4	(0, 3)	(0, 1)	[(1, 2), (1, 3)]	True
5	(0, 0)	(3, 0)	[(2, 3), (2, 0), (4, 2)]	False
6	(5, 1)	(1, 4)	[(4, 4), (1, 5), (3, 1), (1, 1), (2, 5), (1, 3)]	True
8	(0, 7)	(5, 3)	[(4, 4), (5, 1), (4, 2), (5, 6), (6, 0), (3, 2), (5, 2)]	False

104. Inverse pair sums

```
def pair_sums(n, sums):
```

Given a list of n distinct positive integers, constructing the list of all possible sums of adding two integers from that list would be a good exercise for an introductory course however, since that one is just a one-liner, `sorted(x+y for (x, y) in combinations(items, 2))` for anyone who has been solving these problems this far, we would rather turn this problem upside down and ask for a function to reconstruct the original list of n distinct integers from these pairwise sums. For example, given the pairwise sums `[3, 4, 5, 5, 6, 7]`, this function should return `[1, 2, 3, 4]` as the original list of numbers.

Your search algorithm should try out all possible values for the last number of the original list of n elements, and try to fill in the smaller elements recursively from there. Note that once you have chosen this last number, the last two items of the sums list immediately give you the second and third largest numbers of the original list without any branching. Make sure to use an efficient dictionary data structure to keep track of which your current choices have used pairwise sums up and which ones are still available.

The automated tester will give your function only sums lists computed from some list of distinct positive integers, so that there always exists at least one solution. Since this solution is not necessarily unique, your returned solution must be the highest in **colexicographic** order; that is, your solution must use the largest possible value for its last element. Given that last value, the solution must then use the largest possible value for its second-to-last element, and so on.

sums	n	Expected result
[5, 8, 9]	3	[2, 3, 6]
[8, 13, 15, 18, 20, 25]	4	[3, 5, 10, 15]
[7, 8, 11, 13, 16, 17, 17, 20, 21, 26]	5	[2, 5, 6, 11, 15]
[4, 8, 10, 12, 14, 15, 17, 18, 21, 23, 25, 25, 28, 29, 29, 30, 31, 33, 34, 35, 36, 37, 38, 39, 39, 41, 42, 43, 44, 46, 47, 49, 50, 50, 50, 50, 52, 54, 55, 56, 57, 58, 60, 62, 63, 63, 64, 64, 65, 68, 70, 71, 71, 72, 75, 76, 77, 79, 83, 85, 88, 89, 92, 97, 104, 110]	12	[1, 3, 7, 11, 14, 22, 27, 28, 36, 43, 49, 61]

105. Blocking pawns

```
def blocking_pawns(n, queens):
```

A gaggle of chess *queens* have been randomly placed on the generalized n -by- n chessboard. To maintain peace and harmony among these queen bees, we need to place neutral pawns on the board so that any two queens on the same row, column, or diagonal are separated by at least one pawn. This function should compute and return the minimum number of pawns required to achieve this blissful separation.

As is usual in combinatorial problems of this nature, we need the serenity to accept what we cannot change, the courage to change what we can, and the wisdom to accept the difference. Your function should start by finding all attacking pairs of queens. Since you have no choice but to place one pawn between them, you can try out each such pawn position to see what happens when you try to solve the remaining problem with the same approach recursively. Ordering the attacking pairs and the squares between each pair appropriately ought to speed up this search by an order of magnitude.

n	queens	Expected result
8	[(1, 0), (5, 6), (0, 3), (1, 7), (5, 1)]	2
9	[(5, 0), (3, 8), (0, 5), (7, 5), (3, 5), (0, 0), (8, 2)]	7
10	[(7, 7), (9, 1), (3, 1), (1, 5), (4, 4), (8, 4), (2, 8)]	3
13	[(9, 8), (5, 6), (7, 11), (3, 10), (11, 0), (9, 4), (0, 8), (6, 0), (5, 3), (11, 11)]	8

106. Boggles the mind

```
def word_board(board, words):
```

In the spirit of the game of Boggle, you are given an n -by- n board of letters encoded as a list of lists of characters, and a list of acceptable words listed in sorted dictionary order. This function should return the sorted list of all words that can be found on the board by starting from some square and repeatedly moving into one of the four neighbouring squares that have not yet been visited during the traversal of that word. To keep these results manageable even for the humongous wordlist used in our word problems, we are interested only in words that are at least five characters long. Each move must take place in one of the four compass directions, so diagonal moves are not allowed.

This problem is best solved with a nested recursive function whose parameters are the current board coordinates and the word constructed so far. Maintain two sets on the side: the first set to keep track of the words that have already been found during the recursive search, and the second set to keep track of which squares have already been visited for the current word. Inside the recursive search function, use the `bisect_left` function from the Python standard library to look for the current word in the list of sorted words. If the current word appears in that list, add it to the set of found words. If the current word could theoretically still be extended into a longer legal word, continue the recursive search from each unvisited neighbour square.

board	Expected result (using words_sorted.txt)
<pre>[['e', 'c', 'a', 'l'], ['d', 'd', 'i', 'p'], ['i', 's', 'c', 'o'], ['d', 'n', 'u', 'l']]</pre>	<pre>['alpid', 'caids', 'cedis', 'copia', 'decal', 'disci', 'disco', 'eddic', 'laced', 'laics', 'lucia', 'lucid', 'picul', 'place', 'placed', 'plaid', 'plaids', 'pocul', 'undid']</pre>
<pre>[['u', 'o', 'h', 'a', 'r'], ['s', 'a', 'c', 'o', 'r'], ['e', 'e', 'n', 'v', 'e'], ['b', 'k', 'n', 's', 'n'], ['r', 'o', 'e', 's', 'i']]</pre>	(a list of 26 words, the longest of which is 'housebrokenness')

107. The round number round

```
def des_chiffres(board, goal):
```

One segment of the popular French television game show “Des chiffres et des lettres” gives the contestant a board with n positive integers, not necessarily distinct, and a positive goal integer that the contestant must reach by appearing on the board with a suitably clever sequence of moves.

Each move consists of the contestant removing any two numbers a and b of his choice from the board and replacing them with any one of $a + b$, $a - b$, $a \times b$ or a / b , provided that the resulting value is also a positive integer. This function should compute the number of moves necessary to reach the given goal on the given board. If this is not possible within the limit of $n - 1$ moves, after which only one number remains on the board, return `None`.

This problem can be solved by looping through all possible pairs of elements a and b on the board, and for each pair, trying the four arithmetic operations one at a time to determine recursively whether the goal can be reached from the new board produced by that move. The dizzying number of possible choices at each level of recursion sprouts into an exponentially large tree of possibilities, so perhaps the reader can devise a pair of conceptual shears to prune down this mess.

Another good idea is to use **iterative deepening** so that the nested recursive search function receives the recursion `limit` as an additional parameter. Decrement the limit in every recursive call, and treat the case where it becomes zero as a dead-end base case. In the function proper, keep calling your nested recursive search function with increasing values of `limit` until a sufficiently loose `limit` allows the search to get to the goal.

board	goal	Expected result
[3, 6, 7]	32	None
[7, 9, 12, 12]	1	1
[2, 4, 4, 12]	88	3
[1, 3, 5, 14, 14]	77	4
[6, 7, 10, 13, 15]	282	4
[4, 4, 6, 8, 8, 9]	594	4
[3, 4, 4, 9, 12, 16, 17, 20]	208	2

108. Complete a Costas array

```
def costas_array(rows):
```

Over the past decade, this author has learned so much from John D. Cook's excellent blog "Endeavour" that he doesn't even know how to count that high. A recent blog post, "Costas arrays", examines a variation of the famous N queens problem in which rooks replace queens. Of course, placing n rooks on the chessboard so that no two rooks threaten each other is trivial; without diagonal threats, any permutation of the n column numbers works as a legal placement, with one rook in each row, as with n queens.

To make this problem with rooks worth our while, we impose an additional requirement that all **direction vectors** between pairs of rooks must be unique. A placement of n rooks in this manner is called a Costas array, useful in certain radar and sonar applications. Recognizing a legal Costas array should be pretty straightforward for anyone who has gotten this far in this material. However, no algorithms to generate all possible Costas arrays of size n have been discovered, at least not more efficiently than trivial **brute-force backtracking** through all $n!$ permutations and cherry-picking the permutations whose direction vectors between all pairs of rooks are unique.

In this problem, your function is given a list of `rows` where each element is either `None` to indicate that that row does not yet contain a rook, or is the column number of the rook that has already been nailed in that row. Your function should merely determine whether it is possible to place rooks on the unfilled rows somehow to complete the permutation into a Costas array.

rows	Expected result
[3, 1, None, None]	True
[5, None, 1, 0, 3, None]	False
[4, None, None, 0, 5, None]	True
[4, 6, 3, 0, 2, None, 7, None, None]	False
[4, None, None, 6, None, 5, None, 0, 9, None]	True
[9, 0, 7, 6, None, 15, 11, None, 1, None, None, None, 4, 10, 8, None]	False
[2, None, 9, None, 4, 1, 10, 0, None, None, 6, 12, 11, None, 5, None]	True

109. Oppenhoppenheimereimer

```
def lindenmayer(rules, n, start='A'):
```

Lindenmayer systems, or simply **L-systems** for all of us too busy to spit out long words, are a famous formalism typically used to generate fractal strings converted into visually appealing graphical forms. A Lindenmayer system consists of an alphabet of characters (conventionally, uppercase letters) and a dictionary of `rules` to rewrite each character into a string of characters.

The process starts with a string consisting of the given `start` character. In each round, each character `c` in the current string is replaced by the string `rules[c]`, with all replacements performed logically simultaneously. For example, consider an L-system with the alphabet A, B and rules `{ 'A': 'AB', 'B': 'A' }`. Starting from the string 'A', the first round of rule application yields 'AB'. The second round converts this into 'ABA', which then becomes 'ABAAB', and so on.

This function should return the character in position `n` in the string produced by the given `rules`, after applying them as many times as needed to make the string long enough to include position `n`. The automated tester will give your function such large values of `n` that you can't possibly generate the entire string and then look up the answer in it, but you need to be more clever about how you execute this computation. To begin, you should define a nested utility function `length(rules, start, k)` that computes the length of the string that the `rules` produce from the given `start` symbol after exactly `k` rounds. (Just add up the lengths of the strings these `rules` produce from the characters for the replacement of the `start` symbol in `k-1` rounds. Good recursion, with `@lru_cache`.) Armed with this function, you can zero in only on the relevant substring at each round, instead of having to build up the entire universe-sized result string.

n	rules	start	Expected result
9	{ 'A': 'CAB', 'B': 'ABC', 'C': 'ABA' }	'B'	'C'
678	{ 'A': 'DABD', 'B': 'CA', 'C': 'AD', 'D': 'BC' }	'A'	'B'
10**100	{ 'A': 'AB', 'B': 'A' }	'A'	'A'

As the last row of the table above shows, the Grand Slam problem “Infinite Fibonacci Word” in the original *109 Python problems* collection is a trivial special case of this function.

Bonus problem 110: String stretching

```
# For Charmaine: Delayed response is sometimes too early
def string_stretching(text):
```

This wonderful little problem is titled “String stretching” in the Kattis problem collection. The `text` is constructed by initially setting it equal to the secret pattern, then repeatedly inserting that original pattern into a randomly chosen position in the text built so far. For example, if the secret pattern is `'bax'`, the text could develop from the empty string to `'bax'` to `'babax'` to `'babbaxaxx'` to `'babbaxaxbaxx'`, the newly inserted copy of the pattern shown with the green background in each step, to produce the final text equal to the string `'babbaxaxbaxx'`. Many other possible result texts are equally possible, of course.

Your function must compute and return the shortest possible pattern that could have produced the given `text` in this manner after some number of insertions. If there are multiple solutions, your function should return the lexicographically lowest one. Obviously, the `text` length must be some multiple of the pattern length (if nothing shorter works, the returned pattern can simply equal the `text`), and the first and the last characters of the `text` are necessarily also the first and the last characters of the pattern.

The problem is made tricky because the pattern may contain the same character multiple times. This can cause your matching search to branch in several directions whenever you encounter the first character of the pattern in the `text`. Was that position the start of yet another insertion of the pattern, or is that character also part of the pattern after the first character?

text	Expected result
'aaaccccc'	'aac'
'mimmiimi'	'mi'
'cccc'	'c'
'aadbabadbabdbab'	'adbab'
'aagaaagaagaaaaa'	'aagaa' (but not equally possible 'agaaa')
'aiaiiiaiiiaiiiaiiiaiiiii'	'aiii'
'mmmmammmmmammmmmammmmmmmmm'	'mmam'

Bonus problem 111: Casinos hate this Toronto man!

```
# For Alain: Nimis malum quod potest etiam numerare
def optimal_blackjack(deck):
```

Legendary Las Vegas casino moguls such as Benny Binion, Steve Wynn, and Veijo Bellagio built their opulent palaces of sin on the truth that the average punter has no chance against the house in the long run, unless he gets some crooked help under the table. One of James Swain's airport potboilers, featuring Las Vegas casino security expert Tony Valentine, mentioned a device that calculates the optimal blackjack strategy for all possible distributions of remaining cards, for an ultimate edge in card counting. In this problem, you will implement this device for a simplified version of blackjack over the known deck of cards. Since suits don't matter in blackjack, cards are given simply as their blackjack numerical values, with tens and faces all given as 10, and aces given as 11. Aces can be used either as 11 ("soft") or 1 ("hard") in either hand.

Each deal gives the player the first two cards and the dealer the next two. There are no immediate blackjack bonuses, but an ace and a face card just make an ordinary 21. Knowing the deck's contents, the player should take as many cards as needed to optimize not just the immediate outcome of the current deal, but the total outcome across all deals using the known deck. (Our simplified blackjack does not allow splitting, doubling down or surrendering.) Once the player has chosen to stand and has not gone bust, the dealer must keep taking cards until he has 17 or higher, soft or hard, regardless of the player's current total. The winner pays one dollar to the loser; nothing is paid in a push. Deals continue as long as the deck has at least four cards. If the deck runs out of cards mid-deal, cards are taken from the beginning of the deck cyclically.

The function should return the best possible score for the player, assuming optimal overall play in full tilt Frank Fontaine mode over the entire deck. Since the eye in the sky is not here to stop the action after suspicious plays, it may sometimes be optimal to intentionally bust in the current hand to play the remaining deals from a more favourable position in the deck.

deck	Expected result
[8, 10, 2, 11, 6, 7, 5, 6, 10, 8, 4, 7, 3, 8, 6, 10, 10, 4, 10, 10]	0
[10, 6, 6, 10, 9, 9, 2, 6, 9, 2, 5, 7, 4, 2, 10, 10, 11, 3, 2, 3, 10, 5, 8]	3
[9, 9, 9, 10, 4, 5, 10, 11, 6, 10, 10, 8, 10, 10, 9, 7, 3, 10, 10, 8, 2, 2, 9, 7, 8, 6, 5, 6, 3, 9, 10, 10, 11]	-2

Bonus problem 112: Word bin packing

```
# For Debra: Add a point whenever necessary
def word_bin_packing(words):
```

Define two words as **compatible** if their **Hamming distance** is zero;; that is, they don't have the same letter in any position. For example, 'oscar' and 'floyd' are compatible this way (the letter o is in a different position in both words), whereas 'oscar' and 'shram' are not compatible. Your task is to place given words in the smallest possible number of separate **bins** so that all words placed in the same bin are compatible with each other, and return the number of bins needed.

This problem can be solved using recursive backtracking. The recursive search function receives the list of bins built so far as a parameter, initially an empty list at the top-level call. Each bin itself is a list of words that have been placed in that bin in that branch of the recursive search. The recursive function then takes the next word and branches through the options of placing it into an existing bin where it is compatible with the words already in it, or starting a new bin that contains only that word. For each of these options, continue to place the remaining words optimally into the new list of bins.

As always in this sort of combinatorial search, accept the hard constraints you are not allowed to violate, and use them to prune the exponentially branching tree of search possibilities. Since every word must eventually be placed somewhere, you might as well take the bull by the horns and place the most difficult words first in the recursive search to find out how many bins you'll need at minimum. Before the recursive search, you should therefore sort the words by how many conflicts they have with other words.

words	Expected result
['yarb', 'inez', 'doke', 'fray']	1
['klop', 'surv', 'aery']	2
['wagh', 'wash', 'vady', 'glob']	3
['cyan', 'cute', 'geic', 'cote', 'nori', 'boke', 'name', 'size', 'abye']	6
['fana', 'vamp', 'cars', 'scap', 'guff', 'sagy', 'juju', 'orbs', 'uang', 'taha', 'piky', 'turp', 'flow', 'eaux', 'irpe']	7

Bonus problem 113: Probabilistic tic-tac-toe

```
def tic_tac(board, player, probs):
```

The trivial game of three-by-three **tic-tac-toe** can be made more interesting by assigning each cell a probability that the player trying to place his mark in that cell succeeds; failure causes the other player to place his mark in that cell! Whichever mark appears in that cell, the opponent then chooses his response under the same conditions. The two-dimensional three-by-three list `probs` gives the probability that the player successfully marks each cell on his turn. Unlike the version of the game linked above, a move cannot leave a cell empty, so exactly one of the two players' marks will appear on that cell. This ensures the game state space is a tree that can be analyzed with the expectimax algorithm, where each player chooses the move that maximizes the probability of winning.

The player who first creates a three-in-a-row with his own marks wins the game outright. If the board fills up with marks but so that neither player has three in a row, the player whose turn it is to move loses, as is customary in **combinatorial games**. Finally, if the randomly generated test board contains a three-in-a-row for both players, the player whose turn it is to play wins that game.

This function should calculate the probability that the given `player` wins the entire game from their turn to make the next move on the given board, assuming both players make optimal moves for the rest of the game. The function must return its answer as an exact `Fraction` object, and not use any floating-point computations at any point.

This function takes too many parameters to allow tabular representation of test cases, so we shall show just one example. Given the board as `[['.', 'O', 'O'], ['.', '.', 'X'], ['O', 'X', 'X']]`, and the `probs` as `[[Fraction(17, 19), Fraction(58, 67), Fraction(40, 41)], [Fraction(43, 74), Fraction(11, 13), Fraction(1, 1)], [Fraction(4, 5), Fraction(17, 24), Fraction(5, 9)]]`, the player using the marker `'O'` has the exact probability `Fraction(225, 247)` of winning the entire game.

Bonus problem 114: Bandwidth minimization

```
def bandwidth(edges):
```

An **undirected graph** consists of n nodes numbered from 0 to $n - 1$. Each node is connected to some other nodes by undirected edges, so if node u is a neighbour of node v , then node v is also a neighbour of node u . These edges are given as a list to your function so that each element `edges[u]` lists the neighbours of node u .

This function should renumber the nodes of the graph so that each node gets a unique number from 0 to $n - 1$, minimizing the maximum difference between neighbouring nodes across the entire graph. The minimum of the maximum difference between numbers on neighbouring nodes in the given graph, over all possible numberings, is called the bandwidth of that graph, which is what this function should return.

Of course, this problem can be solved with a backtracking recursive algorithm that tries all $n!$ possible ways to number the given nodes, and returns the bandwidth of the best such numbering, but this approach would be horrendously inefficient. A better technique is iterative deepening, in which you first try to construct a numbering with bandwidth 1. If no such numbering is possible, try to construct a numbering with bandwidth one larger than the current one, and so on. As soon as your recursion finds it impossible to number some node and stay within this limit, it can backtrack to try out the next possible way to number the previous node.

It might also be useful to preprocess the graph to determine the best order to fill in the numbers in these nodes. You can also use **symmetry breaking** to cut the search space in half by noticing that you only need to try numbers up to $n/2$ for the first node.

edges	Expected result
[[4, 3, 2], [2, 3], [1, 0], [1, 0], [0]]	2
[[5], [2, 5, 3], [1, 3, 5], [1, 2, 5], [], [1, 0, 3, 2]]	3
[[2, 1, 7, 3], [0, 6, 3, 4], [0], [1, 0, 6, 5], [6, 1, 7], [7, 6, 3], [5, 1, 4, 7, 3], [4, 0, 6, 5]]	4
[[2, 8, 3], [12, 11, 4, 10], [11, 0], [4, 0, 6, 8], [3, 9, 1], [13, 7, 10, 8], [12, 8, 10, 3, 9], [10, 5], [0, 3, 6, 12, 5, 10], [4, 10, 6, 12], [9, 6, 5, 8, 1, 7], [2, 1], [1, 9, 6, 8], [5]]	5

Bonus problem 115: Set splitting

```
# For Christine: Wheat and chaff
def set_splitting(n, subsets):
```

To solve the famous NP-complete set splitting problem, this function should determine whether it is possible to partition the integers from 0 to $n - 1$ into two disjoint groups (let's call these groups the “black” and the “white” groups) so that no subset in the given list of `subsets` is **unicolour**, so that every subset contains at least one element from both groups. This problem formulation therefore fits well with the spirit of diversity in our current society. Each subset is given as a list of integers.

This problem could be solved with a simple recursive algorithm that tries both colouring choices for each integer, backtracking whenever a subset contains only black or white elements. However, a faster recursive search branches on subsets rather than on integers. Initially assign all elements the same colour, let's say black, then keep track of the unicolour subsets that don't have any white elements. Each level of recursion then chooses one such all-black subset (any such subset will do, but the subset with the fewest numbers in it would be the best choice to curb the horizontal branching of the search tree), and colours each number in that subset white on its turn, provided that doing so does not make some other subset all white. Whenever a number is coloured white, every subset that contains it ceases to be unicoloured. Recursion will continue colouring the remaining unicoloured subsets until no unicoloured subsets remain.

It would be good to preprocess the `subsets` by precomputing a list `in_subset` so that the element `in_subset[e]` is the list of indices of `subsets` that contain the integer `e`. This allows the code to quickly iterate over all subsets that (de)colouring the integer `e` affects. Similarly, maintain a list `black_count` so that for each subset, `black_count[i]` gives the count of black elements in `subsets[i]`. These two lists can speed up colouring and backtracking operations by a factor of 10 or more, which is sorely needed to process the 2,000 pseudorandom instances that the automated tester will throw at this function. Other significant optimizations are also possible.

n	subsets	Expected result
6	[[1, 4], [1, 3], [2, 4], [0, 3], [0, 1], [1, 5]]	False
7	[[5, 6], [0, 3], [0, 1, 5], [2, 6], [1, 5], [2, 4], [0, 3, 4]]	True
7	[[2, 3, 6], [3, 4, 6], [0, 1, 2, 3], [3, 4, 5], [0, 2], [2, 3, 4, 6], [4, 6], [0, 1, 2, 4], [1, 5, 6], [2, 3, 5], [0, 6], [2, 5], [0, 3], [1, 5]]	False

Bonus problem 116: Domino poppers

```
def domino_pop(dominoes):
```

A row of `dominoes` has been laid down in a chain; the pips at the end of each tile do not necessarily match the pips at the start of the neighbouring tile. Each domino is a two-tuple of its pips. In a single move, you can choose any neighbouring pair of dominoes whose touching pips match. (As usual, zero pips serve as a **joker** that matches any other number of pips in the other domino.) Your chosen matching dominoes disappear, and the rest of the chain shifts left to fill the gap and keep the entire chain contiguous. This function should compute the smallest possible number of dominoes left in the chain when these removals are done in the best possible way.

This problem is adapted from the problem “Bad to the bone” from the 1992 ACM East Central Regional Programming Contest, as collected by Ed Karrels. In the original version of this problem, the leftmost pair of matching dominoes was always eliminated first. This simpler version allows an efficient **stack-based** solution that iterates through the `dominoes` in one pass from left to right: the current domino either eliminates the matching domino at the top of the stack or is pushed on top, with no branching choices along the way.

However, our version of the problem cannot be solved greedily by always removing the leftmost matching domino pair, as illustrated by the chain `[(1, 2), (2, 3), (3, 5), (2, 4)]` that allows the entire chain to be eliminated by starting with the two dominoes in the center. You must therefore adapt the stack-based solution to branch recursively to examine both possibilities for the current domino: either eliminate it or not.

dominos	Expected result
<code>[(2, 5)]</code>	1
<code>[(3, 5), (5, 6), (6, 6)]</code>	1
<code>[(6, 4), (4, 2), (2, 5), (5, 0)]</code>	0
<code>[(5, 0), (3, 6), (6, 2), (0, 3)]</code>	0
<code>[(3, 1), (6, 6), (5, 1), (4, 3)]</code>	4
<code>[(2, 3), (3, 3), (2, 2), (3, 3), (3, 2), (0, 1)]</code>	0

As an additional note, this Ed Karrels fellow is also somewhat of a pi enthusiast.

Bonus problem 117: Knight jam

```
def knight_jam(knights, player):
```

Most combinatorial games of complete information, such as tic-tac-toe, chess, and go, are partisan, meaning the moves available to each player on their turn differ from the other player's possible moves. For example, in chess, one player can move only the white pieces, whereas the other player can move only the black pieces. This problem, however, concerns a curious **impartial** game inspired by chess, in which both players have the same moves on their turn.

A set of chess `knights`, all the same colour and thus on the same side, stand on some squares of the infinite two-dimensional grid of pairs of natural numbers. On his turn, player zero can move any knight from its current square (x, y) into either $(x - 2, y + 1)$ or $(x - 2, y - 1)$. Player one, on the other hand, can move the knight from its current square (x, y) into either $(x + 1, y - 2)$ or $(x - 1, y - 2)$. A knight can jump over other knights, but cannot jump into a square already occupied by another knight, nor jump outside the game board onto coordinates whose x - or y -coordinate is negative.

Players alternate moves, and the player who gets stuck and cannot move loses. (This must eventually happen, since each move strictly decreases the knight's **Manhattan distance** from the origin.) This function should determine whether the given `player` has a winning move with the given `knights`, that is, any one move for which the opponent does not have a winning response. This can be determined by looping through the possible moves for the given `knights` for that `player`. For each possible move, determine recursively whether the opponent has a winning move from the new situation produced by that move.

knights	player	Expected result
<code>{(2, 2), (3, 0)}</code>	0	True
<code>{(0, 1), (2, 2)}</code>	0	False
<code>{(0, 1), (1, 2), (2, 1)}</code>	1	True
<code>{(4, 4), (2, 4), (4, 1), (2, 1)}</code>	0	True
<code>{(0, 4), (3, 4), (3, 3), (3, 2), (4, 1)}</code>	0	False
<code>{(0, 1), (1, 2), (1, 5), (0, 3), (2, 0), (5, 1), (5, 2)}</code>	1	True

The recursive definition of the winning move may seem a bit queer, as it lacks a base case. But the base case is there, even if you might have to squint your eyes a little bit to see it.

Bonus problem 118: Cubes on the trailer

```
# For Ron: It possibly could be that too
def cubes_on_trailer(xy, xz, yz):
```

A while back, an innocent tweet about a simple math problem of adding up cubes on a trailer sparked a debate about what reasonable assumptions are for a problem like this. Obviously, cubes cannot float in thin air, but is it reasonable to assume that all the unseen columns have been filled as high as they can to produce the given three views? Or is there some hidden catch that is not evident from the problem setup?

In this problem, the trailer bed constitutes the xy -plane as seen from above (we set the long side of the trailer to be the y -coordinate), and the back and side views constitute the xz - and yz -planes, respectively. All three planes are given as two-dimensional lists so that, for example, the expression `xy[x][y]` equals `True` whenever there is a box visible in coordinates (x, y) of that view, and `False` otherwise. This function should return the two-tuple of the minimum and maximum numbers of boxes that can be arranged on the trailer consistent with all three given views.

For example, for the views `xy`, `xz` and `yz` corresponding to the original problem Tweet, this function should return the tuple `(31, 52)`. Most replies to the tweet got the maximum number correct, but those attempting to find all possible solutions usually got the minimum value too high.

Bonus problem 119: Tom and Jerry

```
def tom_and_jerry(n, edges):
```

Engaged in a comical game of catch with alternating moves, Tom and Jerry start on nodes 0 and n , respectively, in an undirected graph with $n + 1$ nodes numbered 0 to n . Same as in all earlier graph problems in this problem collection, `edges[u]` is the list of neighbours of node u .

Both Tom and Jerry can see the entire graph and each other's current location. On his turn, the player must move from his current location into one of the neighbouring nodes. To ensure the game doesn't get stuck in a loop and inevitably moves toward a decisive end, Tom must always move forward to a neighbour with a higher number, while Jerry must always move backward to a neighbour with a lower number.

Tom wins if he can catch Jerry, whereas Jerry wins if he can sneak past Tom on the way towards node 0. This function should determine whether Jerry has a winning strategy to escape in the graph with given `edges`, with Tom making the first move. The graphs constructed by the automated tester are guaranteed to have no dead ends, and both players always have at least one move available to them at each node $0 < u < n$.

edges	Expected result
[[2, 1], [2, 3, 5, 0], [1, 0, 4], [4, 1, 6], [3, 2], [1], [3]]	False
[[2, 1], [3, 5, 0, 4], [0, 4, 7], [1, 7, 6], [7, 5, 2, 6, 1], [4, 1], [4, 7, 3], [4, 3, 6, 2]]	True
[[1], [2, 3, 5, 0, 4], [1, 7, 8], [1, 5, 6], [8, 10, 5, 1, 9], [3, 4, 1, 7], [7, 3], [6, 2, 5], [4, 10, 9, 2], [8, 4], [8, 4]]	True

This problem is adapted from an exercise in “Algorithms” by Jeff Erickson. It is an example of a pursuit-evasion game restricted to a certain type of graph.

Bonus problem 120: Pinch to a pound

```
def pinch_moves(board, player):
```

The board game of Go is famous for its complexity, which far surpasses that of chess. AlphaGo, a computer player trained through repeated self-play, recently conquered this game to beat the best human professional players. Before that, the board size had to be drastically reduced from the usual 19-by-19 for computers to have any chance against even decent local club players.

Even more curiously, Go's emergent combinatorics are deep enough that restricting the board to one dimension (the resulting game is called “pinch”) still yields a surprisingly complex game! Instead of finding the best move with minimax search, this function only needs to generate all legal next moves for the given `player` on the given `board`. Each move is given as the next board state it produces from the original board. The `board` consists of characters 'W' for a white piece, 'B' for a black piece, '.' for an empty point and 'R' for a point that is currently empty because some pieces were removed in the previous move. Return the list of resulting boards in sorted alphabetical order.

In one-dimensional Go, surrounding a solid enemy group on both sides removes it from the board, replacing that group with a sequence of 'R'. (A solid group lying at the edge of the board can be captured by surrounding it just from one side.) All the previous occurrences of 'R' turn into '.'.

A “suicide” move that causes your own group to be captured is illegal, unless that move captures an opponent group, such as white playing to the empty point of 'B.B' to capture both black pieces. One last “ko” rule says that playing to a singular 'R' point is illegal if doing so captures an opponent group. If no opponent groups get captured, playing to 'R' is legal, as is also playing to one of the points of multiple consecutive 'R' squares even if this captures an enemy group.

board	player	Expected result
'..BB'	'W'	['.WRR', 'W.BB']
'.WW'	'W'	[]
'BBB...WWW'	'B'	['BBB..BRRR', 'BBB.B.WWW', 'BBBB..WWW']
'W..W.'	'B'	['RB.W.', 'W.BW.']
'RBW.BRRB'	'W'	['RBW.B.WB', 'RBW.BW.B']

Bonus problem 121: Minimal Egyptian fractions

```
# For Dianne: The mystery of E
def minimal_egyptian(f, k):
```

The two earlier problems of Egyptian fractions had you implement the greedy algorithm and the pairing algorithm to express an integer fraction a/b as a sum of distinct unit fractions $1/n$, and you should solve at least one of those problems before attempting this problem to get an idea of how unit fractions work. Neither of the previous algorithms is guaranteed to result in the smallest possible number of distinct unit fractions. The function required here should determine whether the given `Fraction` object `f` can be expressed as a sum of at most `k` distinct unit fractions.

This problem can be solved with a recursive search that takes an additional third parameter, `n`, the smallest positive integer you are allowed to use as a unit-fraction denominator. This parameter ensures that the result will use only distinct unit fractions. The base cases are $k = 0$, where the answer is trivially `False`, and when the fraction is of the form $1/m$ with m no smaller than n , where the answer is trivially `True`. Otherwise, try out the possible values for n for the first fraction $1/n$ in the breakdown (the arithmetic used in the greedy algorithm works here just as well), and for each such n , determine recursively whether the remaining fraction $f - 1/n$ can be broken down into at most $k - 1$ distinct unit fractions whose denominators are at least $n + 1$.

f	k	Expected result
<code>Fraction(25, 51)</code>	2	<code>False</code>
<code>Fraction(4, 51)</code>	3	<code>True</code>
<code>Fraction(52, 53)</code>	4	<code>False</code>
<code>Fraction(45, 62)</code>	4	<code>True</code>
<code>Fraction(97, 145)</code>	4	<code>False</code>

The still-open Erdős–Straus conjecture asks whether there exists an integer fraction of the form $4/n$ whose breakdown into distinct unit fractions cannot be done with at most four terms.

Bonus problem 122: Unity partition

```
def unity_partition(n):
```

John D. Cook recently tweeted out the following interesting algebra fact, originally proven by none other than Ronald L. Graham himself: Every positive integer $n > 77$ can be expressed as a sum of distinct positive integers whose reciprocals add up to exactly one. For example, the integer 87 can be partitioned as $2 + 4 + 6 + 15 + 60$, and we leave it for the readers to convince themselves that the corresponding sum $1/2 + 1/4 + 1/6 + 1/15 + 1/60$ does indeed add up to unity.

This function should return the list of these integers that add up to the given n and satisfy the required property that the sum of their reciprocals equals one, in ascending order. Since the asked integer partition is not necessarily unique, to satisfy the automated testing, we force the result to be unique by requiring this function to return the **lexicographically smallest** solution list that starts with the smallest integer that can be extended to a working solution, followed by the lexicographically smallest way to complete this solution list.

For anyone who has been solving these problems up to this point, Python should make this relatively easy by using recursion to enumerate the possibilities, aided by the standard library Fraction class for exact fraction arithmetic. However, the difficulty in this problem is pruning the search space and recognizing inevitable dead ends as soon as possible, so all required results are computed within the time limit imposed by the tester.

n	Expected result
78	[2, 6, 8, 10, 12, 40]
88	[2, 4, 8, 20, 24, 30]
121	[2, 3, 8, 36, 72]
150	[2, 4, 9, 18, 26, 39, 52]
200	[2, 3, 9, 36, 60, 90]
250	[2, 3, 12, 28, 44, 77, 84]
300	[2, 3, 10, 40, 50, 75, 120]

Bonus problem 123: Ants on the rod

```
# For Rimma: Float down the river to the place of clay moons
def ants_on_the_rod(ants, width):
```

Some tiny ants have been placed on a rod with the given `width` measured in centimetres. Each ant is encoded as a two-element list `[position, direction]`, where `position` is its location on the rod and `direction` is either `+1` or `-1`, depending on whether it faces right or left. The ants have been numbered from `0` to `n - 1`, left to right, according to their initial positions on the rod. No two ants are initially in the same position.

Each ant moves at a constant speed of one centimetre per second. Whenever two ants collide, both immediately reverse direction, so the numbering of ants on the rod never changes. If two ants moving towards each other are one centimetre apart, they collide “halfway” and end up in their original positions, with their directions flipped, after one second. If the colliding ants are two centimetres away, they collide after one second and temporarily share the same location, facing in opposite directions.

An ant falls off the rod the moment its position equals either `0` or the `width`. This function should return a list of times when each ant fell off the rod. Since the ants can't pass each other along their journey, the resulting list of falloff times will necessarily have an “organ pipe” shape whose values first ascend and then at some point start descending.)

ants	width	Expected result
[[1, 1], [2, 1], [3, 1], [4, -1]]	5	[4, 4, 3, 2]
[[1, 1], [2, -1]]	5	[2, 4]
[[1, -1], [2, 1], [6, -1], [7, -1]]	9	[1, 6, 7, 7]
[[2, -1], [4, 1], [8, -1], [9, -1], [10, 1], [12, -1]]	14	[2, 8, 9, 12, 10, 4]
[[4, 1], [7, -1], [10, 1], [14, 1], [15, 1], [16, 1], [21, 1], [22, -1], [23, -1]]	26	[7, 22, 23, 22, 16, 12, 11, 10, 5]

This problem was inspired by the famous puzzle that asks you to prove that no matter how you initially place the ants and their directions, every ant must have fallen off the rod after at most `width` seconds. The flash of insight needed to prove this is that, once we ignore the individuality of these ants, this system is equivalent to one in which the ants simply quantum tunnel through each other without changing direction. Clearly, no ant can stay on the rod for more than `width` seconds.